PLAIN-ENGLISH ALPHA AND EXECUTION DOSSIER

Inside the whale's alpha: copying loses, atomic breadth wins.

A transaction-level investigation, 1,444-cell parameter atlas, and literal specification of the strongest public alpha footprint we could recover from the chain.

Public-data study covering June 19, 2026 to August 25, 2026. Analysis generated August 25, 2026.

-6.15%	+41.94%	+27.32%
return from blindly copying all 139 large signals	return when one trigger matched at least 18 maker accounts	return from those signals after the threshold was selected

The answer in thirty seconds

His observable edge is **selective, informed-looking liquidity consumption**. The important trades do not just spend more money. In one mined transaction, they take offers from an unusually broad set of accounts already waiting in the order book.

Bottom line: among otherwise eligible events, transactions that matched at least **18 distinct maker accounts** won 23 of 30 and returned +41.9% after the original conservative execution assumptions. Transactions below 18 returned -11.6%. That is the clearest public fingerprint of when this trader appears to mean the trade.

1. **The headline is misleading.** The wallet made $5,644,317, but its top five winners produced 144% of net profit. Removing those five makes the wallet negative.

2. **Blind copying is not the edge.** 139 delayed $100 copies lost $855.28. The later sample also lost -6.68%.

3. **The public trade feed hides structure.** We decoded all 149 trigger transactions. Every one was a successful V2 `matchOrders` call with the target as the BUY taker.

4. **Breadth separated signal from noise.** Broad sweeps beat the market price by +23.2 points; narrower transactions missed it by -0.6 points.

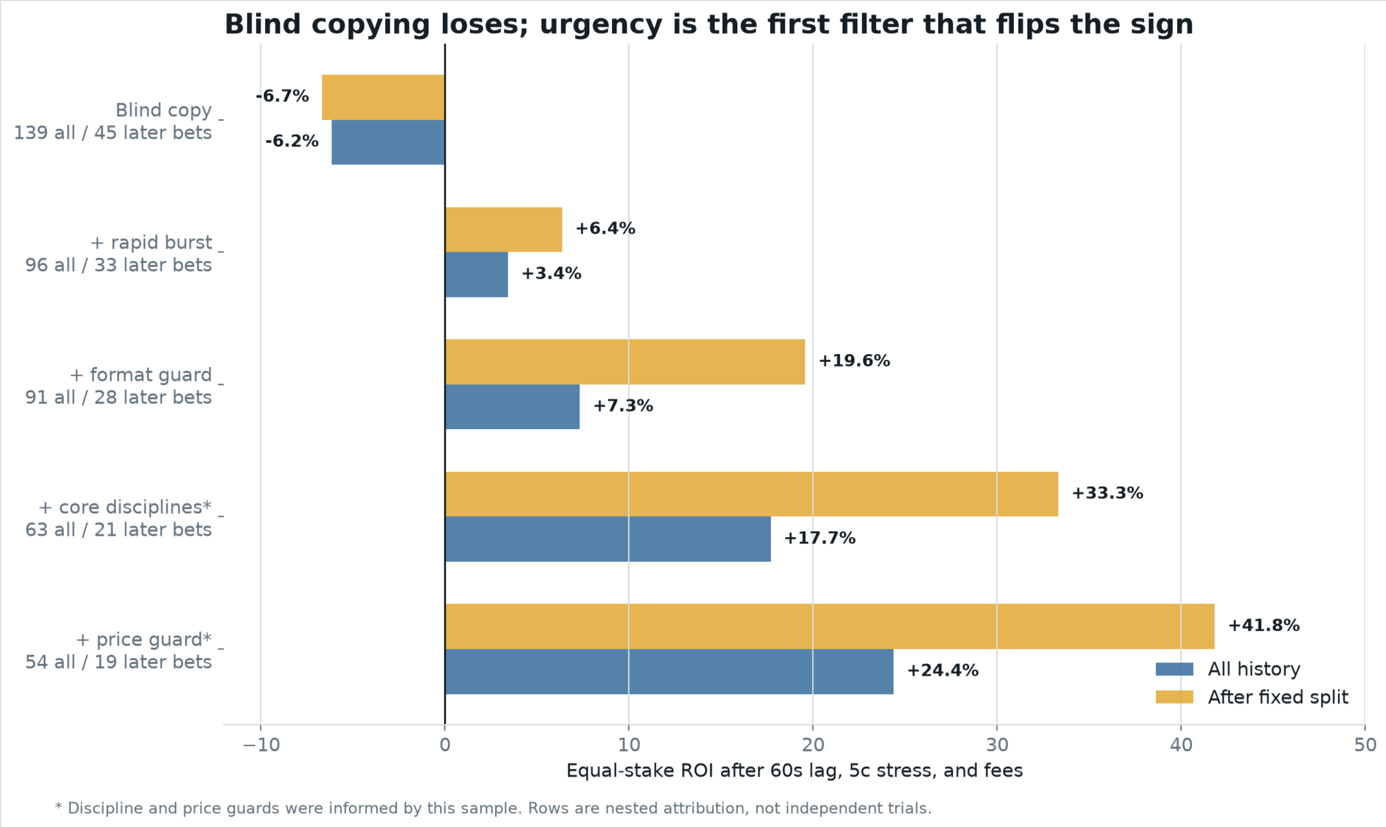

What blind copying would have done: the first row is the naive strategy. It loses over the full history and after the fixed date split. Earlier behavioral filters found useful clues, but they did not expose what was inside the triggering blockchain transaction.

First, the painful result: copying loses

The original registered test was deliberately harsh, but ordinary: it did not give the follower the whale's old price.

1. Wait until the target has crossed **$25,000** of aggressive buying and at least 70% of its net direction points to one outcome.

2. Keep only the first signal for each underlying match, so a series and its individual maps cannot be counted as independent ideas.

3. Wait **60 seconds**, then use the first unrelated public trade in the following minute as the available price. If none appears, retain the trigger price.

4. Make the price five cents worse, include the account-observed fee curve, and place the same **$100** stake every time.

The result was 80 wins and 59 losses, yet -$855.28 net profit on $13,900 of turnover. The maximum historical drawdown was $2,055.76.

> **This is what blind faith costs:** even a spectacular trader can be a bad signal to copy. The follower sees the action late, pays a worse price, does not know final position size, and cannot reproduce inventory management, maker fills, exits, or rebates.

Why winning more than half the bets was not enough

Prediction-market prices already contain a probability. Buying a contract at 60 cents means paying roughly as if it has a 60% chance to win. A strategy can win 58% of its bets and still lose if it repeatedly pays prices that require a higher hit rate.

Blind copying won 80 times. The public prices implied about 77.15 wins. The gap was only +2.0 points. In ordinary language, the extra wins were neither large enough nor valuable enough to establish a generic “copy his big buys” edge.

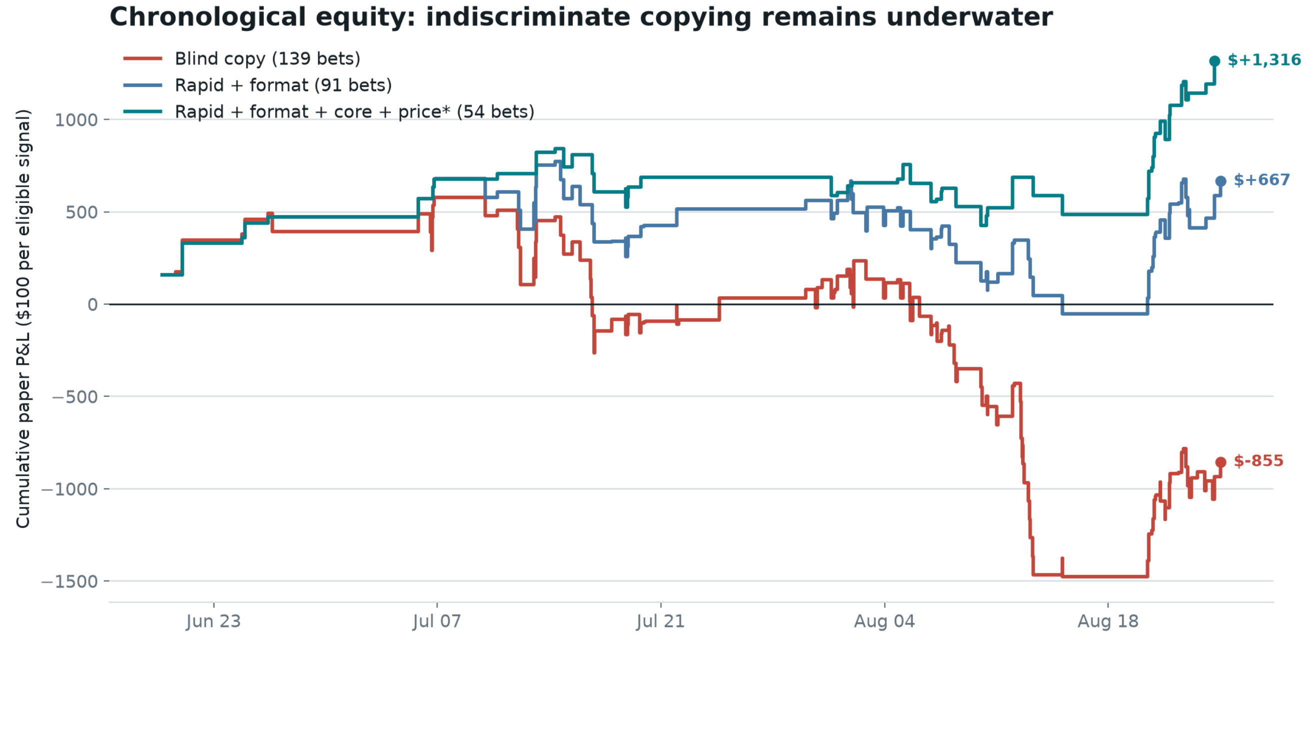

* Core-discipline and price filters are exploratory. Curves use each rule's own chronological event set.

The red blind-copy line remains underwater. More selective rules improve because they reject most trades. This was the clue that the wallet's value lives in selection, not in its address.

What if a copy bot is nearly instant?

Speed helps, but the broad replay finds no sharp 1-second-versus-5-second cliff. The dangerous variable is the price paid after the whale has consumed the available offers.

The expanded audit crosses **15 delays**, from the trigger's own second through five minutes, with **17 adverse-price assumptions** from zero through 30 cents. That is 510 latency-by-price results across blind and alpha-filtered copies before the fee and maker-threshold atlases. Every primary cell includes the observed fee curve. At the most optimistic same-second price with no extra adverse movement, blind copying returns only +3.4%. At one second plus one cent it returns +1.0%. At one second plus two cents it is already negative at -0.9%.

> **The practical threshold:** at one-second latency, blind copying historically breaks even at only about 1.53 cents of additional adverse price. Under one second plus one cent, the breadth-filtered held-out sample returned +35.6%. Being fast is not enough if the whale just removed the cheap liquidity; selecting the right trigger matters more.

The timestamps impose an important limit. Polygon block timestamps and the historical public tape are recorded to whole seconds. A 0.1-second bot and a 0.5-second bot cannot be separated honestly. The same-second scenario may also include unrelated prints whose exact within-second order is unknown, so it is an optimistic lower bound rather than a reproducible fill.

Polymarket's official lifecycle separates an off-chain `MATCHED` state from `MINED` on-chain settlement. This strategy needs the mined `matchOrders` calldata to count maker addresses, so its clock starts at the block timestamp. A bot with private or earlier CLOB-match information is testing a different signal and is not represented by this public-wallet replay.

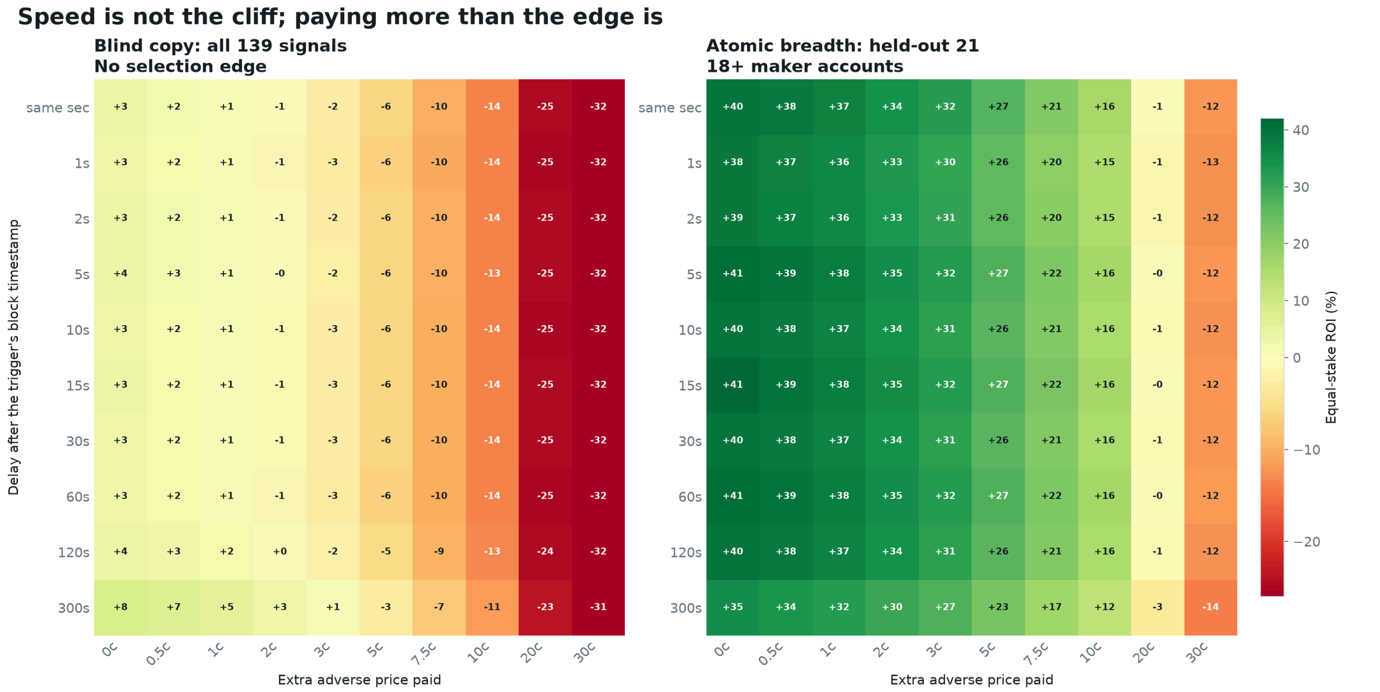

Selected anchors from the full 15-delay x 17-cost atlas. Same-second is optimistic because ordering inside a one-second timestamp is unknown. A 0.1s or 0.5s bot lies between same-second and 1s. Every cell includes the observed fee curve.

Read across for execution quality; read down for speed. Blind copying flips from pale green to red at roughly two cents in almost every sub-minute row. The held-out breadth strategy remains green much farther across. A 0.1-second or 0.5-second bot lies between the first two rows because the source timestamps are only one second precise.

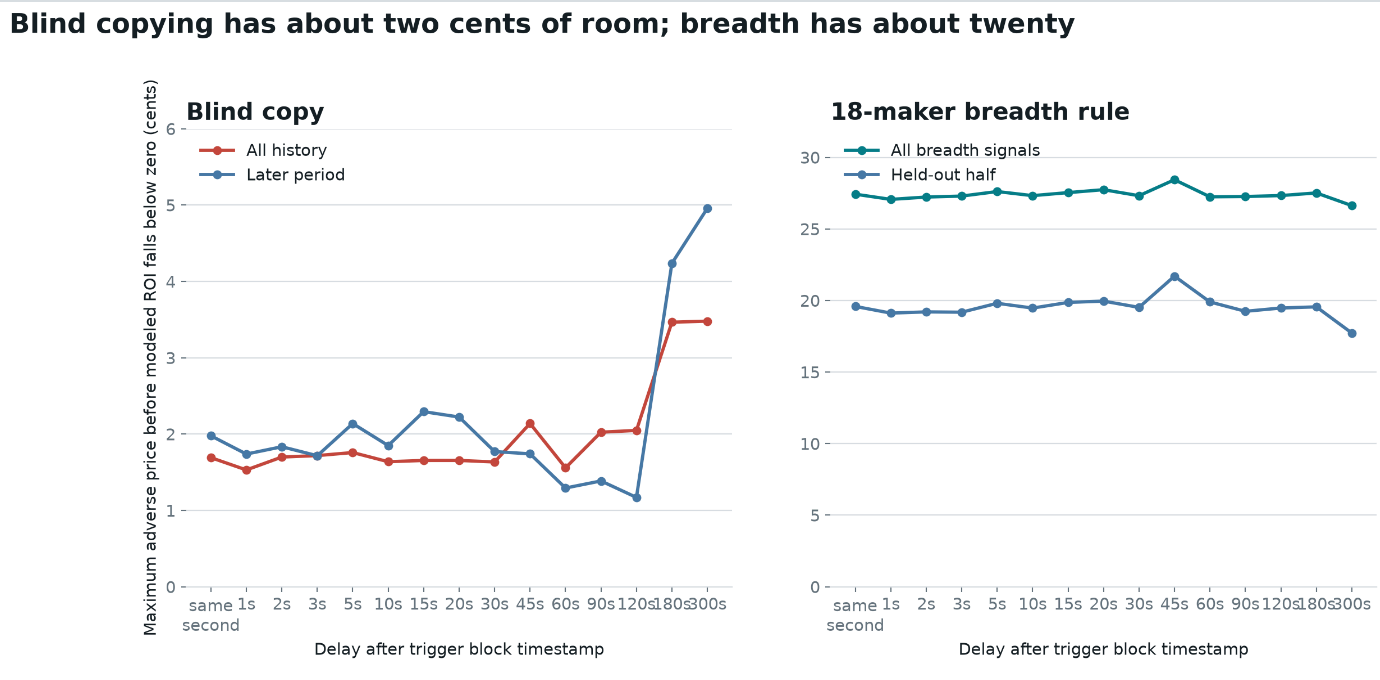

Break-even values are solved from the same equal-stake replay and include fees. The near-flat lines show no measured sub-minute latency cliff; exact historical order-book depth remains unavailable.

The blind copier has roughly 1.5 cents of room at one second. The 18-maker filter has roughly 19.1 cents in the held-out half. The near-flat lines mean this dataset does not support a claim that sub-five-second speed is the main edge.

The full copy-trading parameter atlas

A backtest should show the whole neighborhood, not one convenient assumption. This atlas varies delay, price deterioration, fees, and the maker-breadth cutoff while keeping equal stakes and event deduplication fixed.

15 delays from same-second through 300 seconds	**17** adverse-price assumptions from 0c through 30c	**6** fee-curve rates from 0% through 5%	**26** maker-breadth cutoffs from 5 through 30

ONE-SECOND SCENARIO	BLIND COPY	18-MAKER HELD OUT
No extra adverse price	+3.1%	+38.4%
One cent worse	+1.0%	+35.6%
1.5 cents worse	+0.1%	+34.3%
Two cents worse	-0.9%	+32.9%
17.5 cents worse	-22.5%	+2.3%
Twenty cents worse	-24.9%	-1.2%

What the spectrum says: blind copying is a thin-margin trade that flips near 1.5 to 2 cents. The held-out alpha filter does not flip until roughly 19 to 20 cents. The major difference comes from selecting the transaction, not shaving a few seconds from the bot.

Latency barely moves the result; execution price decides it

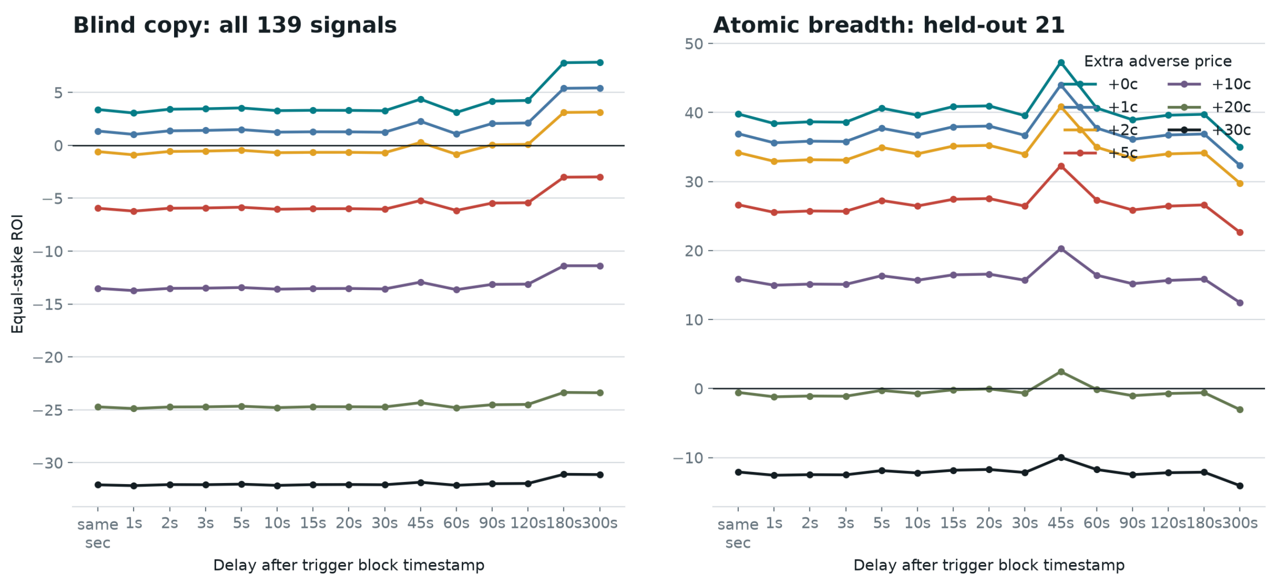

All 15 measured delays are shown. Lines are nearly horizontal below one minute. The historical data cannot rank 0.1s against 0.5s inside the same timestamped second.

All fifteen measured delays: each line holds the adverse price constant. The lines barely move through the sub-minute region, while moving between price-cost lines changes the answer immediately.

The whole cost curve: blind copying dies near 2c, breadth near 20c

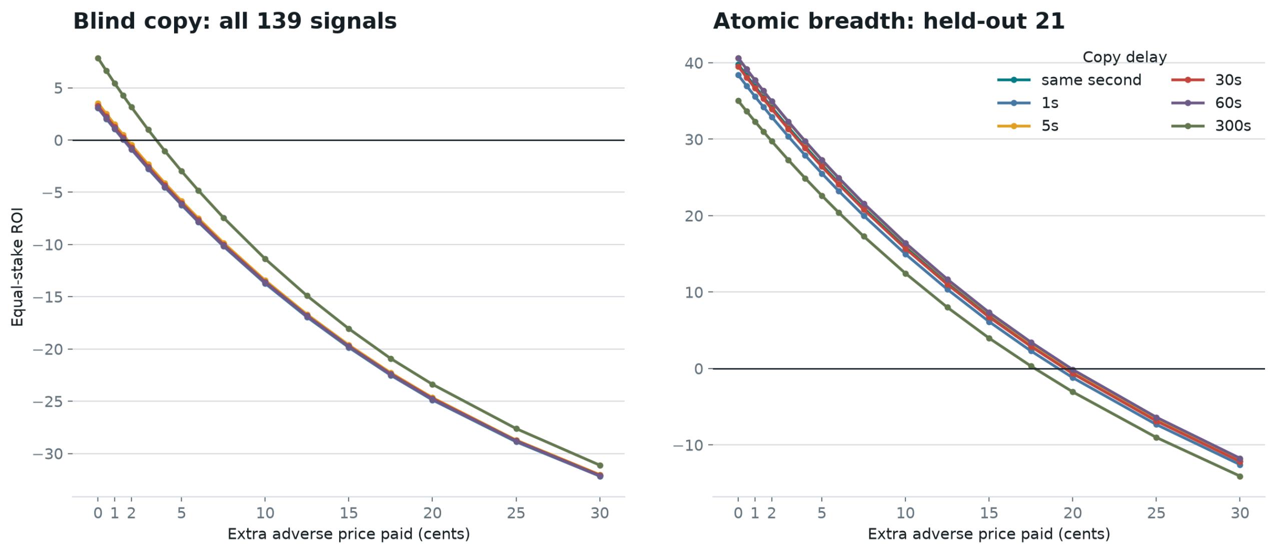

Seventeen explicit adverse-price assumptions from 0c to 30c are shown. These are stress scenarios around public prints, not reconstructed historical order-book fills.

All seventeen cost assumptions: blind-copy curves cross zero near two cents regardless of delay. The held-out breadth curves cross near twenty cents.

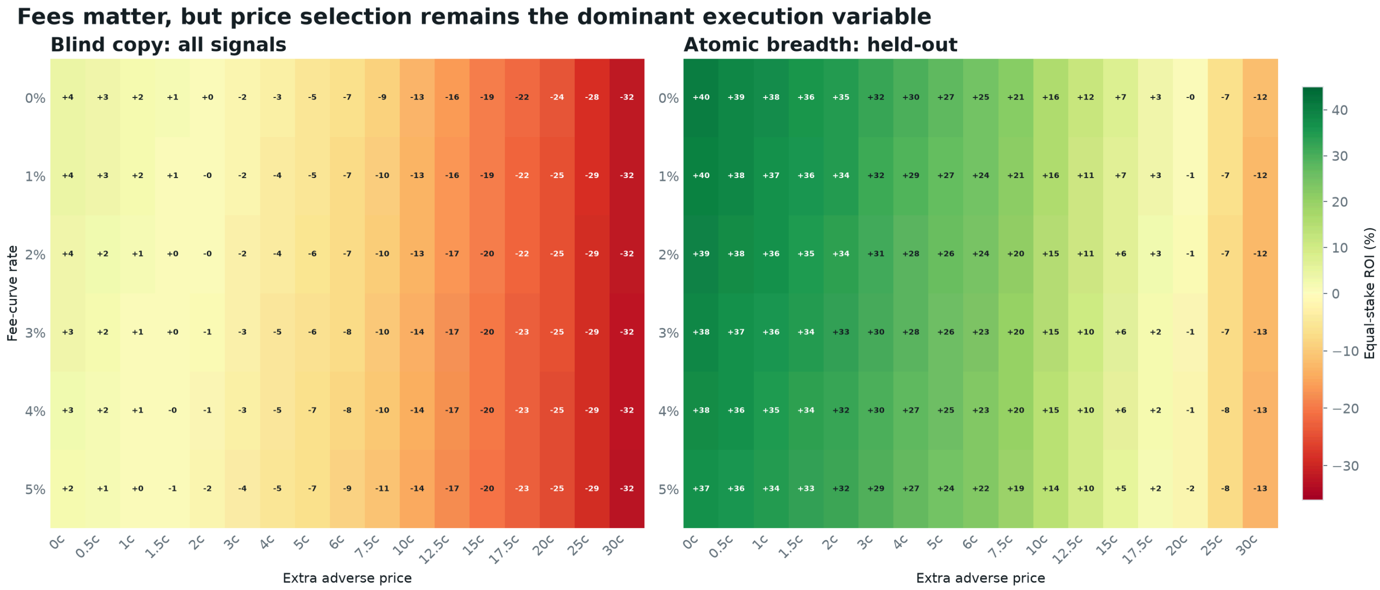

One-second replay. Each all-in price adds fee_rate x price x (1-price). The observed account curve uses 3%; other rows are explicit counterfactuals.

Fees shift the result gradually. Paying through the book shifts it much faster. The primary replay uses the account-observed 3% fee curve; every other fee row is a labeled counterfactual.

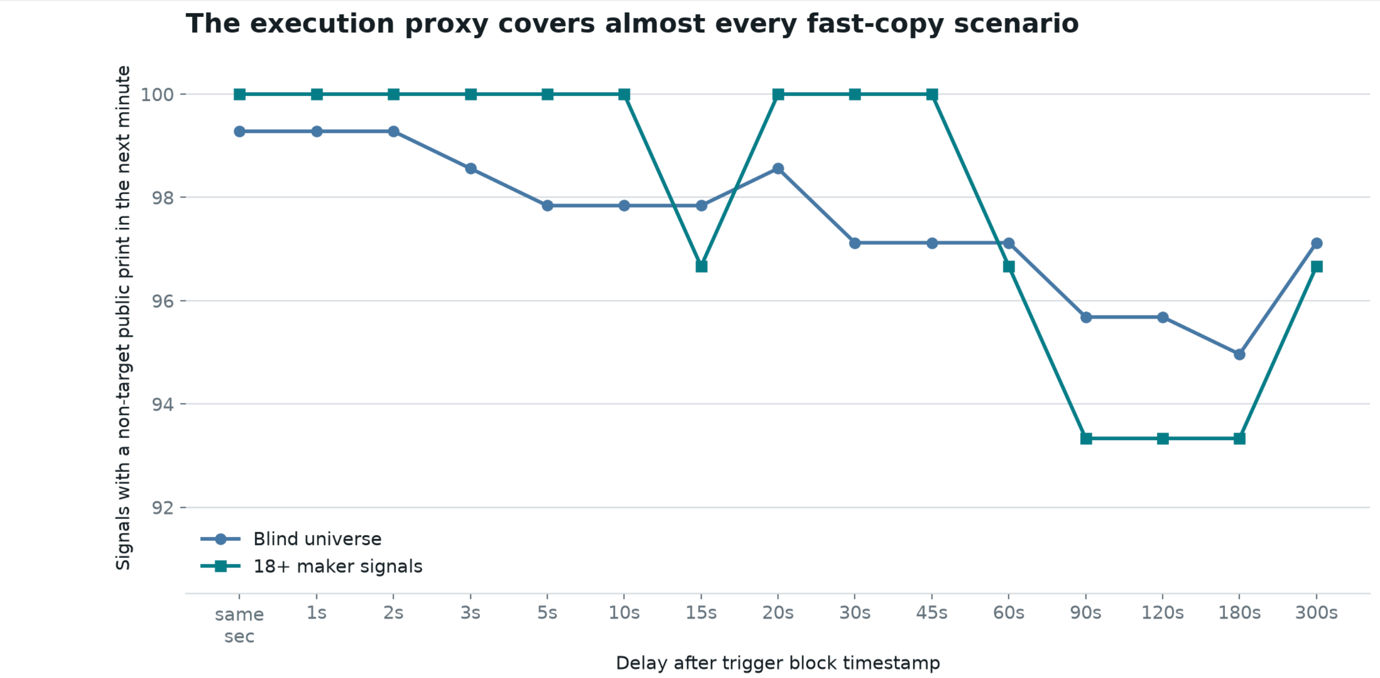

Missing-print events are retained at the trigger price rather than dropped, preventing an easy upward selection bias. Coverage is not proof that $100 of ask depth was executable.

Fast scenarios still have roughly 99% public-print coverage. Missing prints are kept with a fallback instead of being dropped, but a public print still does not prove historical ask depth or queue position.

The breakthrough came from looking inside the trade

A public feed row says the whale bought one outcome. Polygon calldata says exactly how that buy was assembled.

Polymarket's V2 exchange accepts one taker order and an array of maker orders in a call named `matchOrders`. The taker demands immediate execution. Each maker had already signed an offer. One public-looking trade can therefore be an atomic sweep through many counterparties and several prices.

One public print is not necessarily one simple trade. It can be one urgent order consuming a large section of the order book in the same mined transaction.

We fetched and decoded every one of the 149 trigger hashes. All 149 showed this wallet as the decoded taker buying the signaled token. The typical trigger matched 19 maker orders from 14 distinct maker accounts. 66 triggers crossed more than one price level. The reconstructed notional agreed with the signal to within less than one millionth of one percent.

> **Why this matters:** the final dollar amount says how much was bought. Maker breadth says how much standing liquidity the trader chose to consume at once. Those are different behaviors.

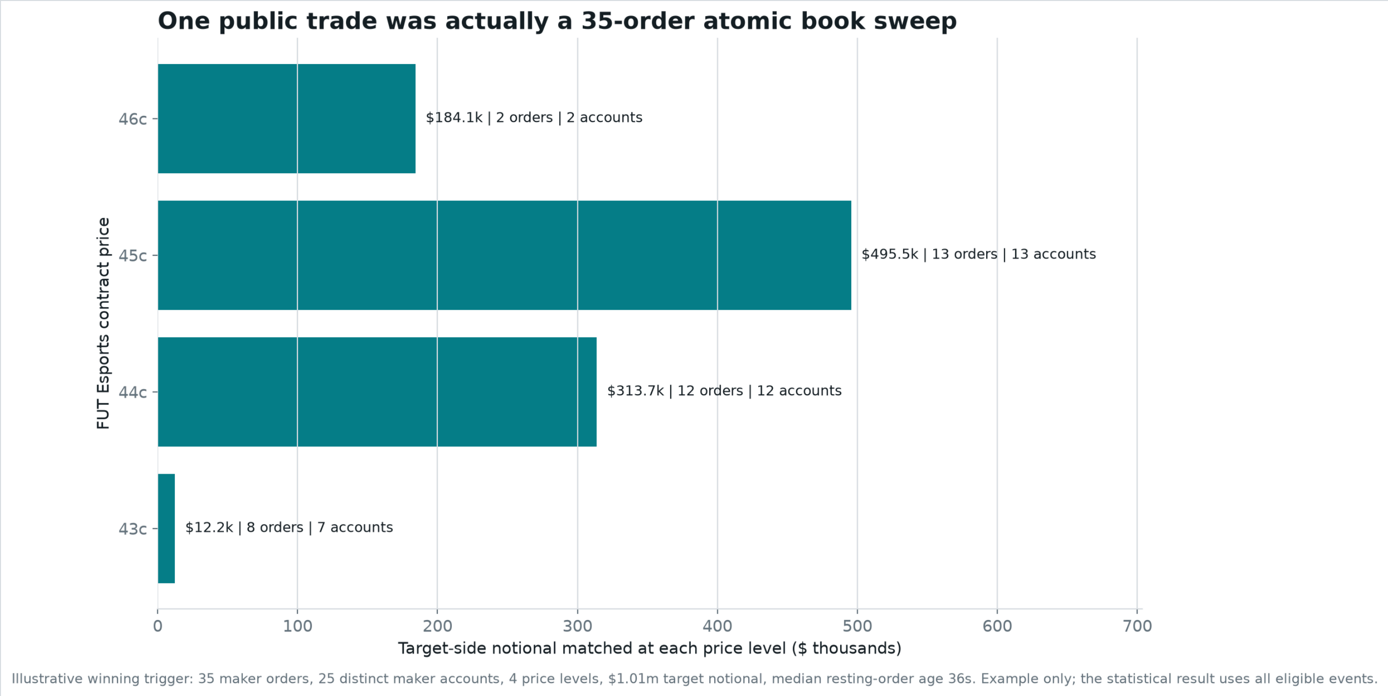

An illustrative winning FURIA trigger matched 35 maker orders from 25 distinct signed accounts across four price levels, from 43 to 46 cents. The target consumed about $1,005,518.93 of notional; the median resting order was 36 seconds old. This example explains the mechanism. The statistical result uses every eligible event.

The discovery: atomic breadth

The strongest observable fingerprint is simple: **did the triggering transaction match at least 18 distinct maker accounts?**

The comparison starts with the same conservative event universe: first canonical signal only, core tennis, soccer, and esports disciplines, full-event contracts rather than maps or short markets, target concentration of at least 70%, and trigger prices between 30 and 85 cents. That leaves 80 events. Breadth is the final split.

WHAT HAPPENED	18+ MAKERS	BELOW 18
Number of bets	30	50
Wins implied by public prices	16.04	29.30
Actual wins	23	29
Price-implied win rate	53.5%	58.6%

WHAT HAPPENED	18+ MAKERS	BELOW 18
Actual win rate	76.7%	58.0%
Gap versus price	+23.2 points	-0.6 points
Equal-stake return after costs	+41.94%	-11.56%

At a $100 equal stake, the 30 broad sweeps made $1,258.31. Even after removing the five most profitable winners, the remaining return was +20.1%. The narrower group lost $578.01.

> **The market-price test:** public prices expected about 16.04 wins; 23 occurred. The one-sided Poisson-binomial probability of at least that many wins under those market probabilities is **0.57%**. The same calculation finds no edge below 18 makers.

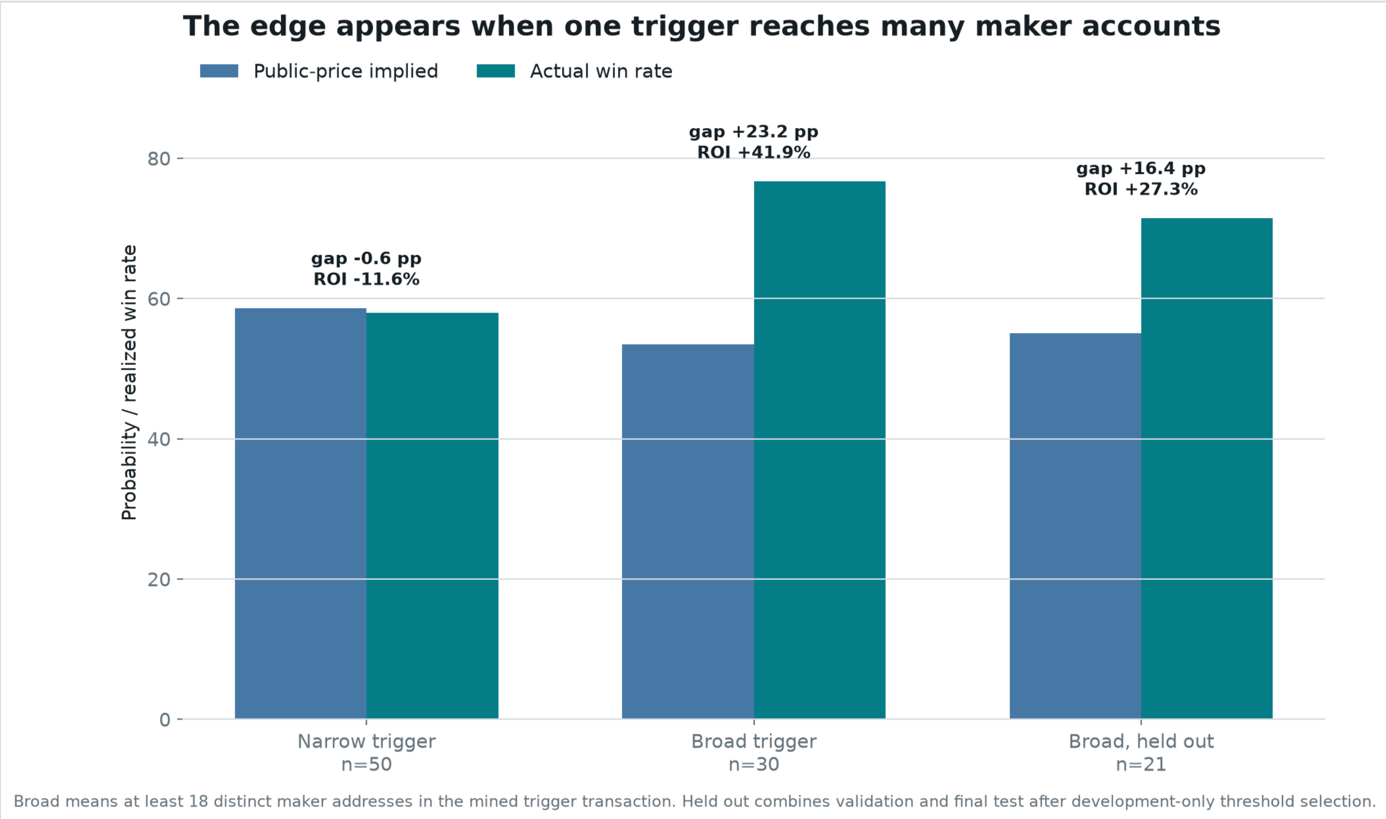

Public prices expected similar difficulty after filtering. Broad atomic sweeps won far more often than their prices implied; narrower triggers landed almost exactly where the market predicted.

Maker breadth is a graded transaction feature, not a wallet-size label

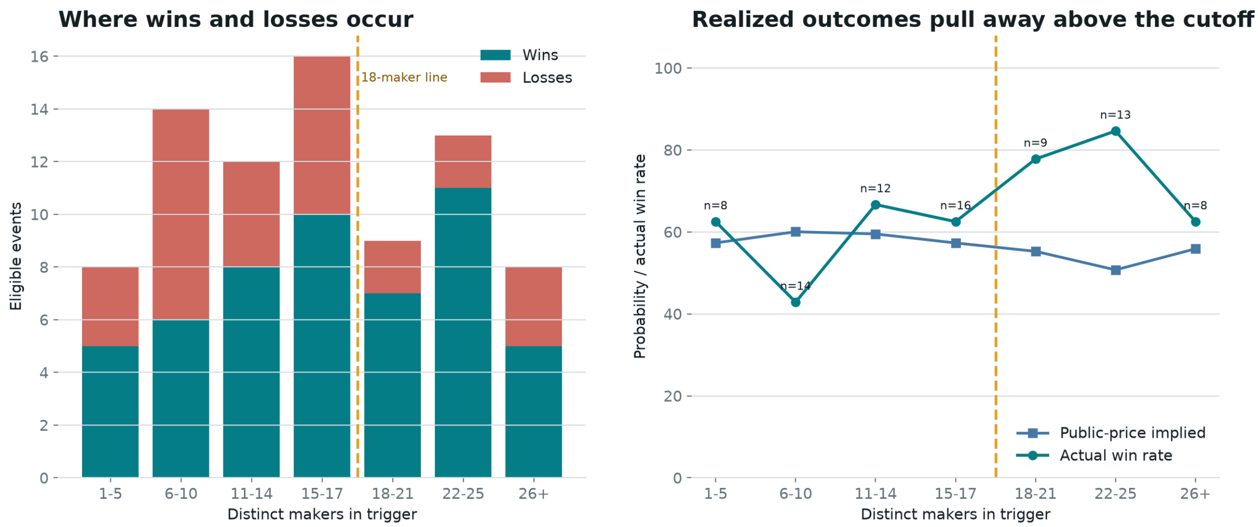

Bands are descriptive. The exact cutoff of 18 was selected on the first half only; the right panel uses the original 60s + 5c execution reference.

The 18-maker rule sits inside a graded transaction pattern. Outcomes begin separating from public probabilities around the broad-sweep region; stricter bins eventually become too small to trust.

A bigger dollar trade is not the same thing as a broader sweep

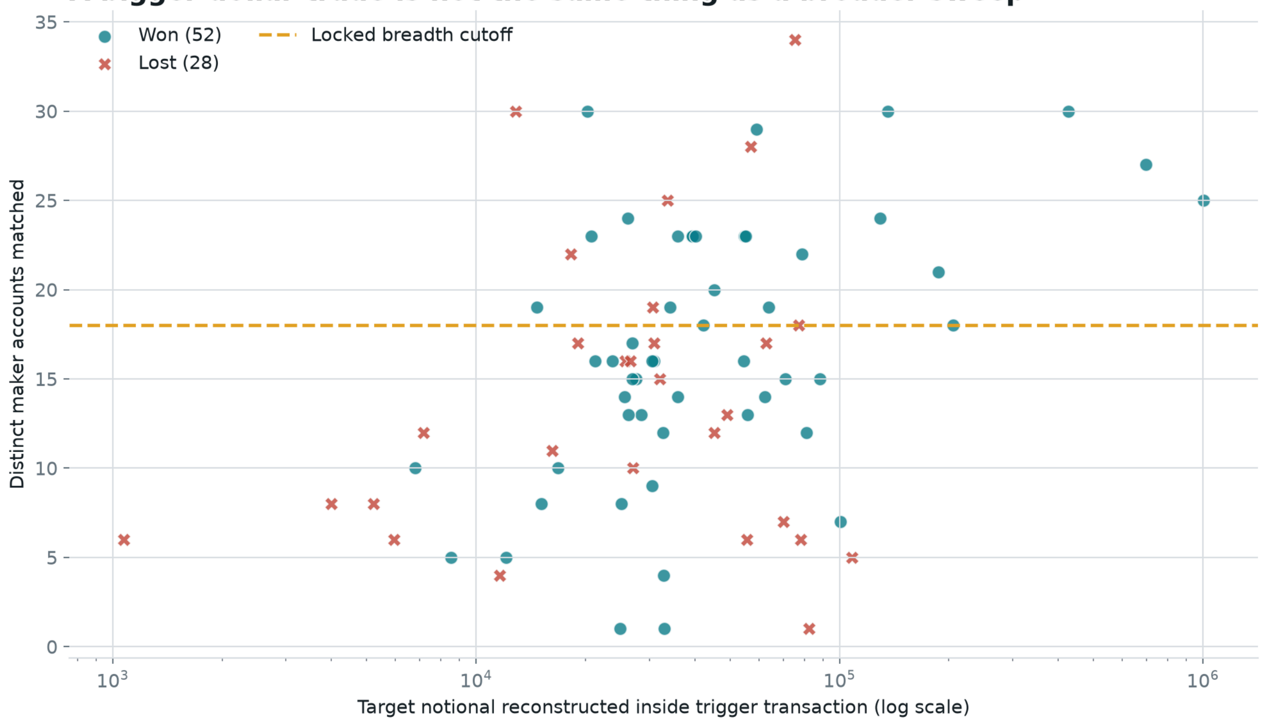

Probability-offset model after urgency and period controls: breadth OR 4.94, p=0.043; log notional OR 0.94, p=0.858. Retrospective explanatory model.

Large notional appears above and below the line and among winners and losers. The informative dimension is how broadly one order consumed standing liquidity, not simply how many dollars it represented.

Did it hold up after discovery?

The 18-maker threshold was selected using only the first half of eligible history. Integer cutoffs from 5 through 30 were compared, each requiring at least eight development bets. The next half was then scored without changing the rule.

The broad-sweep rule returned +76.1% in development, +5.1% in the next block, and +44.0% in the final block. Combined after selection, 21 signals won 15 times and returned +27.32%.

> **Held-out calibration:** public prices expected 11.55 wins; 15 occurred. Resampling whole trading days gave a held-out ROI interval from +0.7% to +58.3%.

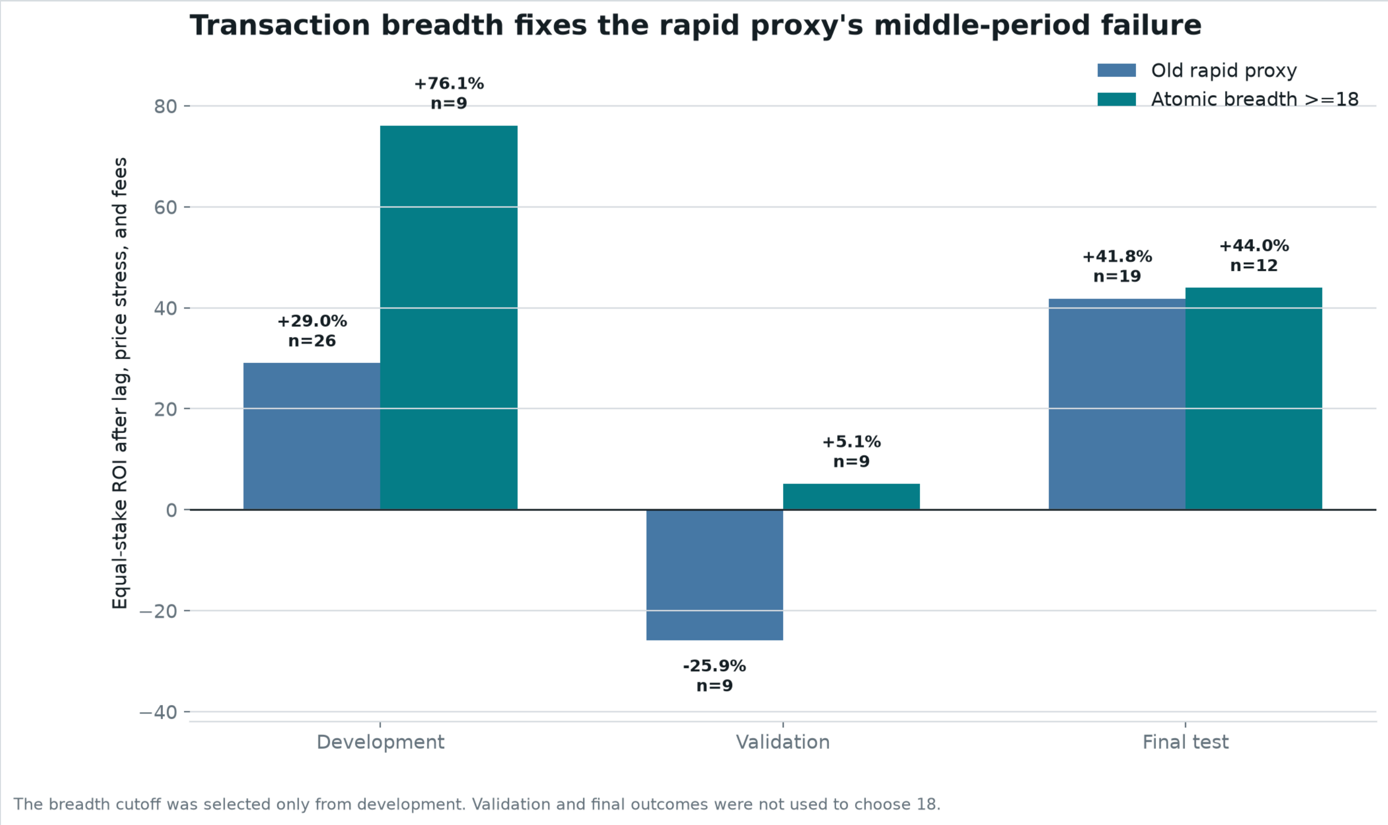

The middle period was modest, but crucially stayed positive. The final block rebounded without changing the 18-maker rule.

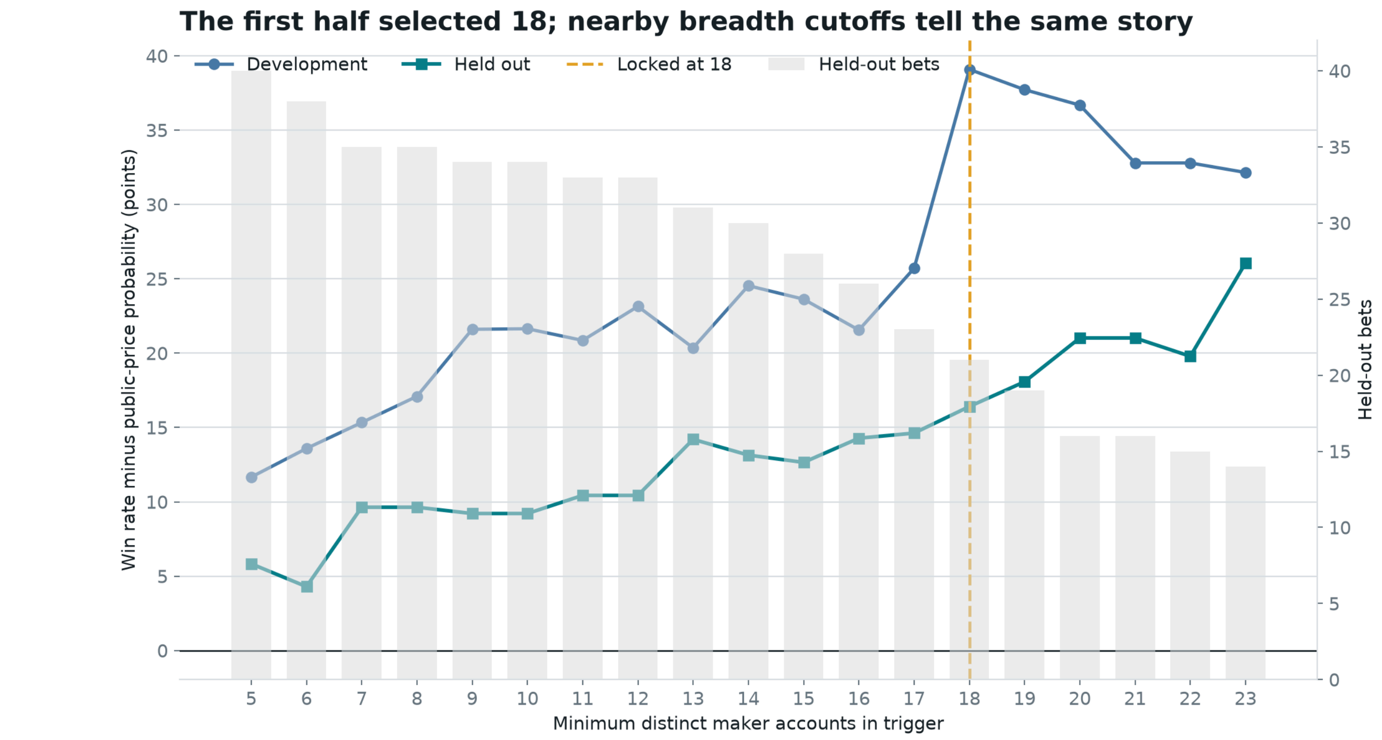

The gray bars shrink as the rule becomes stricter. The formal market-null simulation repeats the threshold search rather than treating every cutoff as an independent test.

The dashed line marks the cutoff chosen on development data. The held-out curve was not used to select 18. Nearby cutoffs tell a similar story, while very strict cutoffs leave too few bets.

His alpha, literally

The recoverable alpha is a conditional probability error: public prices understate how often the target's side wins when one eligible BUY transaction consumes liquidity from at least 18 distinct maker accounts.

Observable alpha definition

```
B(tx) = count(distinct makerOrders[].maker)
Eligible(tx) = first canonical event AND core discipline AND
full-event market
               AND concentration >= 70% AND trigger price in
[0.30, 0.85]
Signal(tx) = Eligible(tx) AND target is BUY taker AND B(tx) >=
18
Probability alpha = realized outcome - public execution-proxy
probability
```

Measured probability alpha: +23.21 points over all 30 broad sweeps and +16.41 points over the 21 signals after development selection. Below 18 makers, the corresponding gap was -0.59 points.

Measured economic alpha: +41.94% under the original 60-second plus five-cent stress; +27.32% in the post-selection half; +35.60% under the realistic one-second plus one-cent scenario.

In ordinary words: most large buys are not special. The special-looking ones are single decisions that clear offers from many signed accounts at once. That is the public footprint of conviction. The private model, information, or judgment that produced the conviction is still hidden.

This distinction matters. We did not reverse-engineer his sports model or prove private information. We recovered an **observable gating variable** that identifies when his action historically contained information beyond the market price. That is the most literal alpha statement public evidence supports.

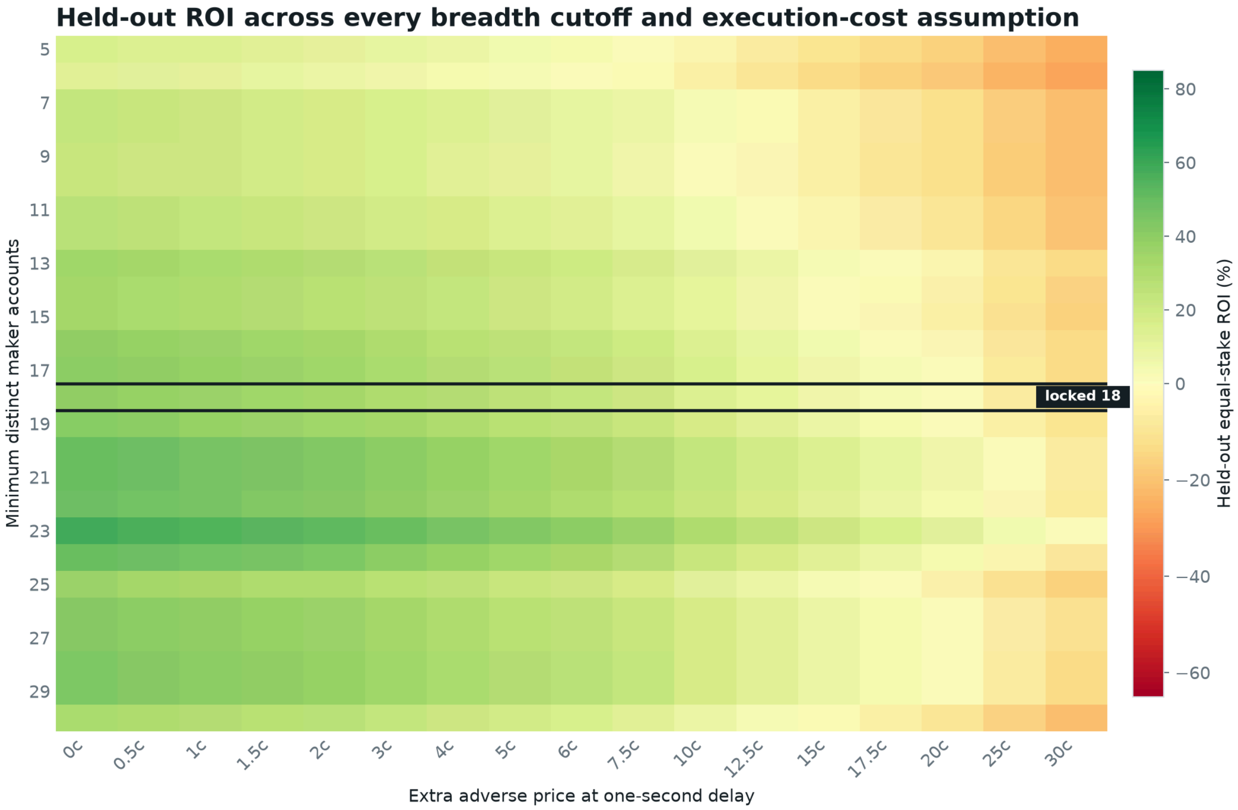

Thresholds 5-30 are descriptive sensitivity checks. High cutoffs contain very few bets and can show extreme returns. The formal rule remains the development-selected cutoff of 18.

Breadth by cost: the dark horizontal outline is the frozen 18-maker row. Neighboring cutoffs form a broad profitable region at realistic costs; very strict rows are based on tiny samples and are not alternative strategies.

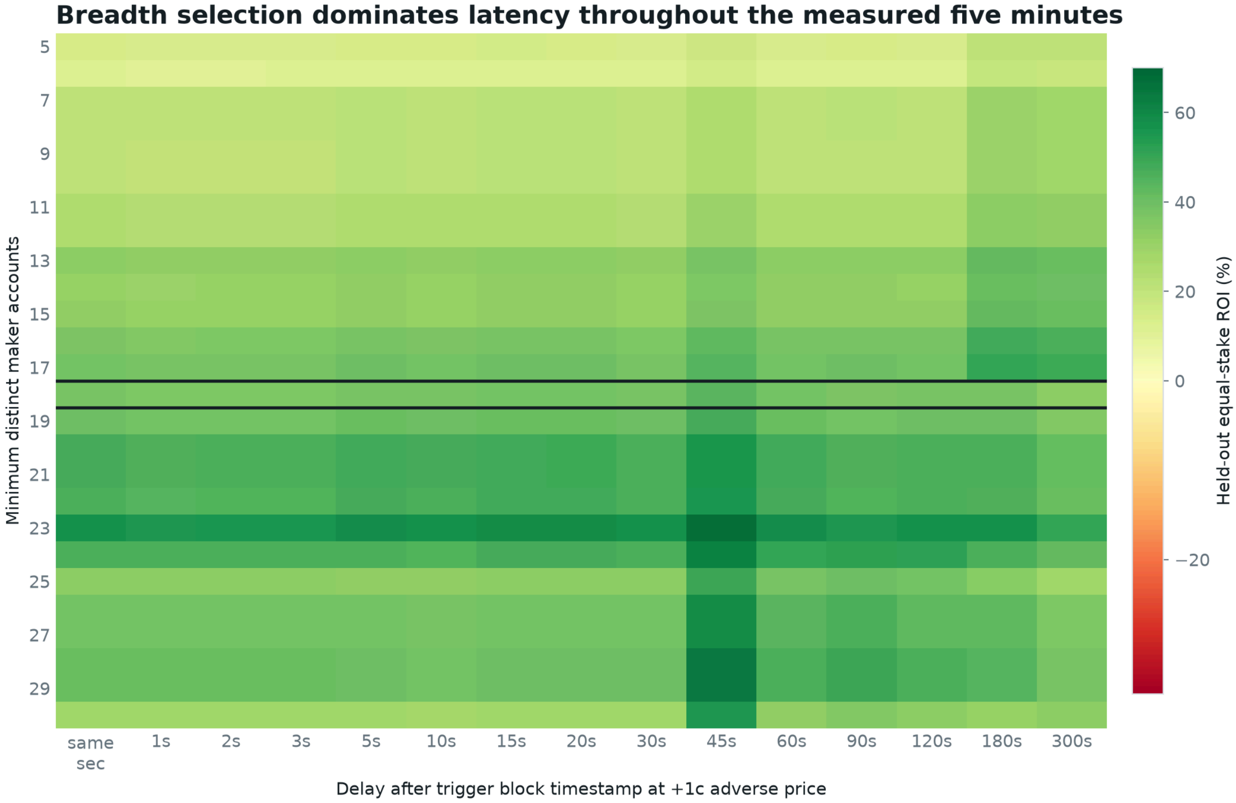

The horizontal color bands are the key result: changing the maker cutoff moves returns much more than moving from 1s to 60s. Same-second remains an optimistic bound.

Breadth by latency: horizontal bands dominate vertical changes. Choosing broad sweeps mattered much more than whether the historical proxy entered at one, five, or sixty seconds.

Speed was the clue, not the final answer

The earlier investigation found “conviction compression”: at least 80% of observed taker buying arriving in the final minute. Across the full sample, rapid signals returned +24.4% while slower signals returned -24.5%. That was useful, but incomplete.

The rapid rule lost 25.9% in its middle validation block. Atomic breadth returned +5.1% in that stage and +44.0% in the final test. Speed pointed toward urgency; calldata revealed the stronger form of urgency.

> **Flow speed asks:** did the target accumulate quickly? **Atomic breadth asks:** did one executable decision consume offers from many different signed accounts? The second measure is closer to the act of demanding liquidity now.

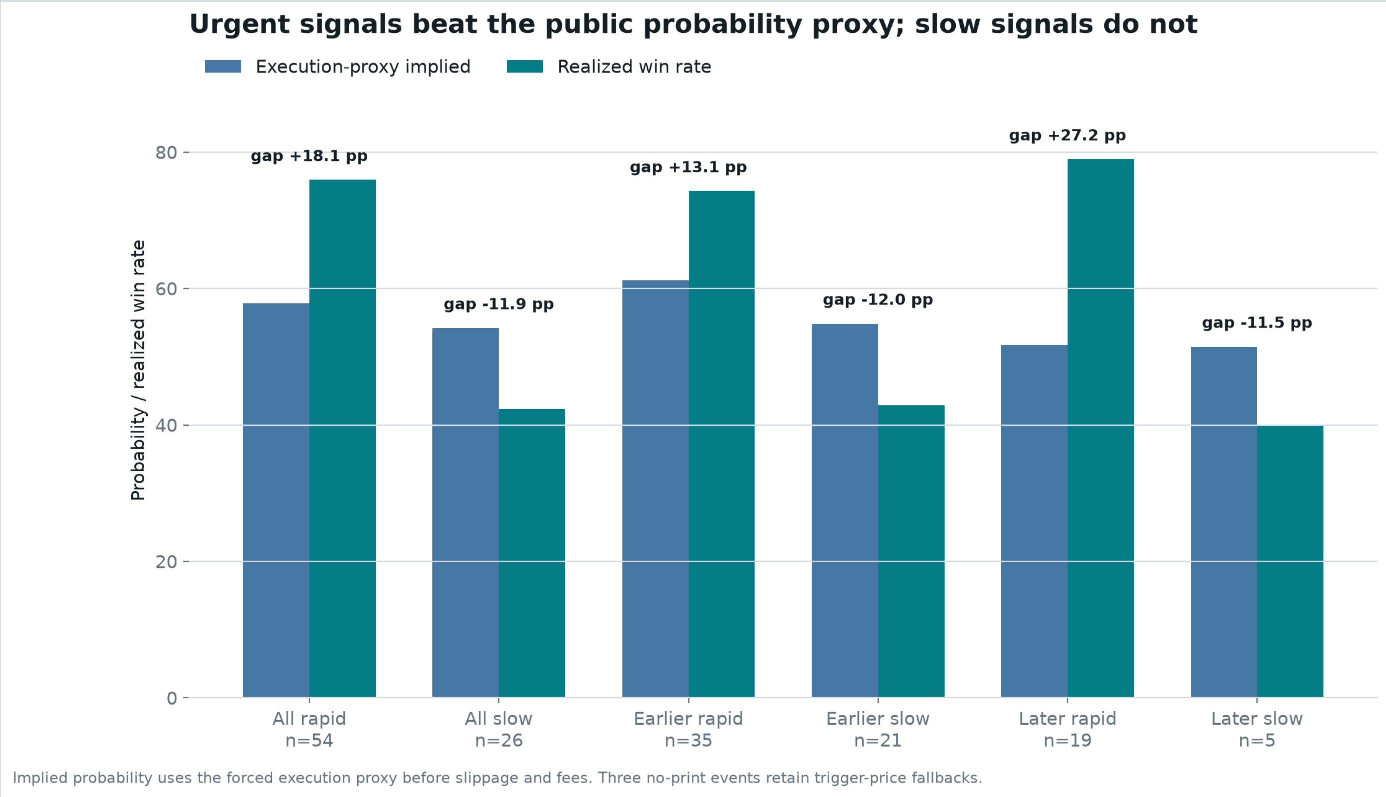

Rapid buying was the first useful behavioral clue. It remains a confidence tag, not a hard requirement, because its own middle validation period lost money.

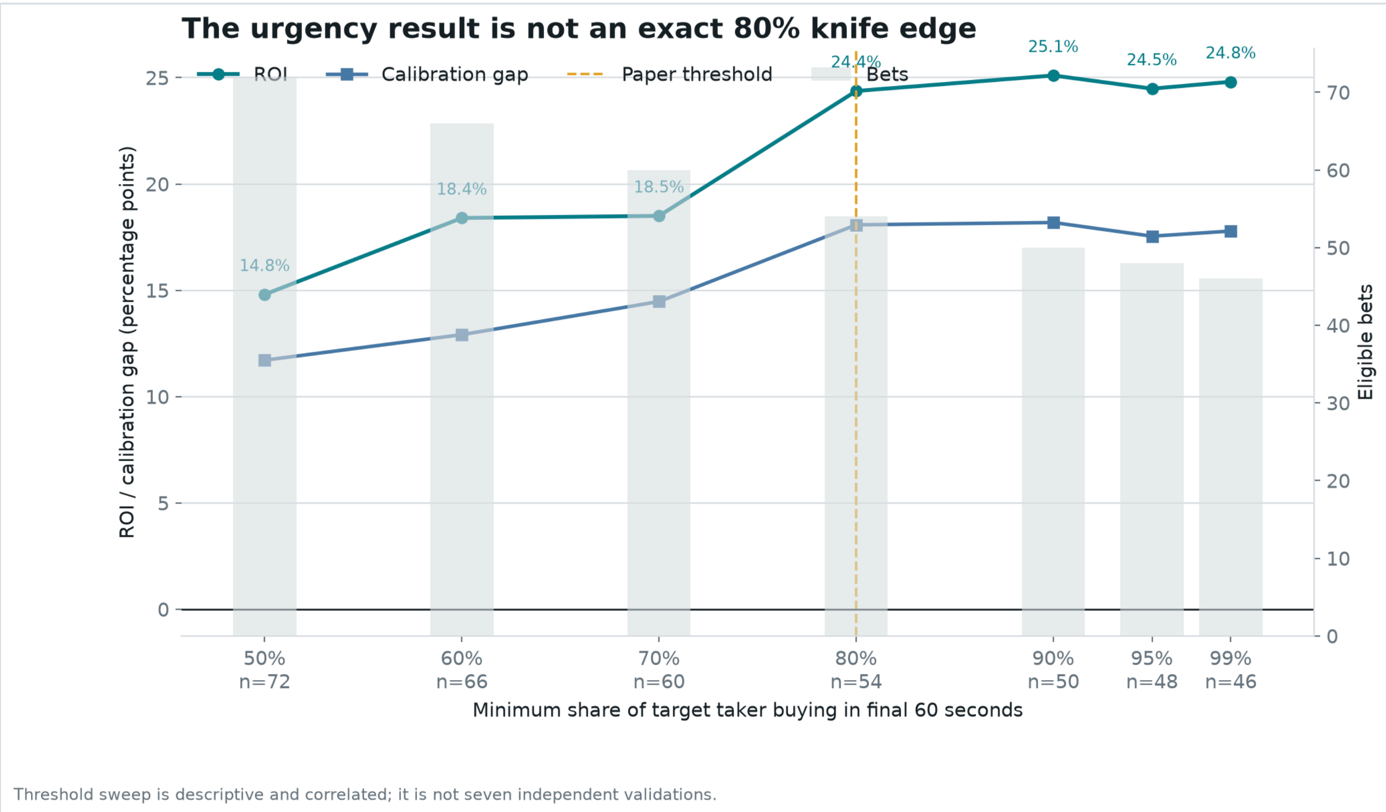

The urgency clue was not one lucky exact cutoff, but these overlapping thresholds are descriptive rather than independent confirmations.

The algorithm: atomic-breadth-18

This is the discovery translated into a rule that can be frozen and tested prospectively.

Paper-trading specification

1. Keep only the first signal for the real underlying event in core tennis, soccer, and esports.
2. Reject individual maps, single-game fragments, and short-horizon side markets.
3. Require the target to be at least 70% concentrated on one outcome and the trigger price to be 30 through 85 cents.
4. Decode the mined V2 `matchOrders` transaction and verify that the target is the BUY taker for the signaled token.
5. Count distinct addresses in `makerOrders[].maker`. Continue only when the count is at least **18**.
6. Do not add an artificial delay. Snapshot the live best ask and displayed depth as soon as the mined call is decoded. Record block-to-detection and detection-to-order latency separately.
7. Create a paper-only marketable FOK limit for a fixed $100 stake, capped at one cent above the first observed best ask and never above 0.90. Record a rejection when displayed depth is insufficient; never assume a fill.
8. Hold accepted paper fills to resolution. Mark at least 80% final-minute taker flow as high confidence, but do not make it a second gate yet.

The one-cent FOK buffer is a prospective paper convention, not a claim that historical depth existed. The replay's one-second plus one-cent cell is the closest available proxy and returned +35.6% in the held-out half. The original 60-second plus five-cent case remains the registered comparison. No live orders, martingale, inferred eventual whale size, discretionary exits, or threshold changes after a bad result.

Could this just be luck, sport mix, or bet size?

A promising pattern is easy to manufacture accidentally. The analysis attacked this one in several different ways.

CHRONOLOGY	**It stayed positive after selection.** The 21 held-out signals returned +27.3%, including +5.1% in validation and +44.0% in the final block.
COMPOSITION	**Similar markets still separated.** Relabeling broad sweeps only within discipline, three price bands, and fixed time period produced an effect this large with one-sided **p=0.0064** across 63 comparable bets.
TRADING DAYS	**One busy day did not create the result.** Resampling whole days put the broad-minus-narrow probability gap at +23.8 points, with a 95% interval from +6.2 points to +42.8 points.
THRESHOLD SEARCH	**The declared cutoff search was simulated under market probabilities.** After repeating selection and held-out scoring 20,000 times, a held-out effect this large occurred with one-sided **p=0.046**.
DOLLAR SIZE	**Bigger notional did not explain breadth.** A probability-offset model controlling for rapid flow, trigger notional, and time period estimated 4.94 times the outcome odds for a broad sweep, robust p=0.043. Trigger notional itself had odds ratio 0.94 and p=0.858.

Do not translate “18 addresses” into “18 independent humans.” They are distinct signed maker accounts. One person or market-making system can control multiple addresses. The feature measures real contract-level execution breadth, not human headcount.

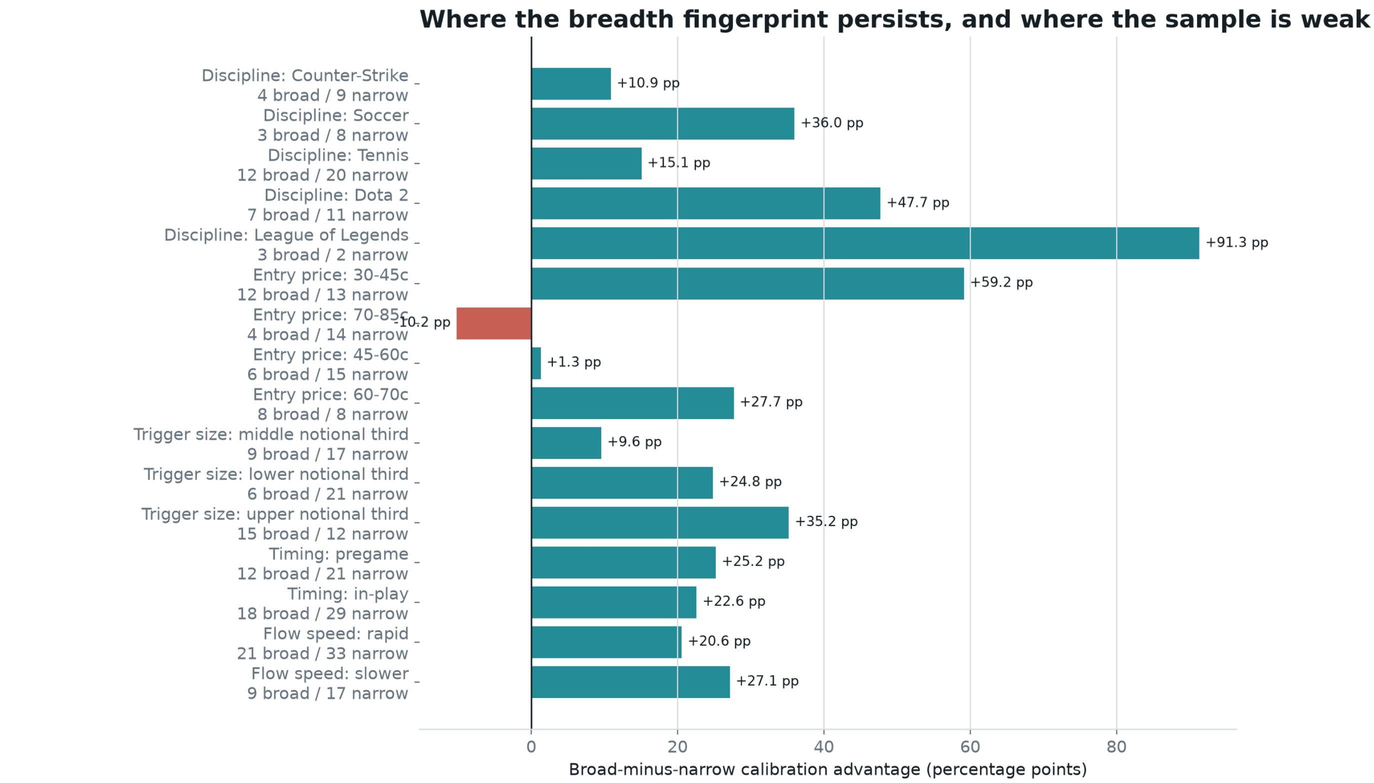

Descriptive subgroups with at least two broad and two narrow bets. These are not independent tests; small cells can swing sharply and should be treated as failure probes, not extra discoveries.

Most descriptive slices retain a positive breadth advantage. The clear weak spot is the 70-85 cent entry-price band, and several sport cells are tiny. This chart is a map of where to challenge the rule, not a menu for subgroup optimization.

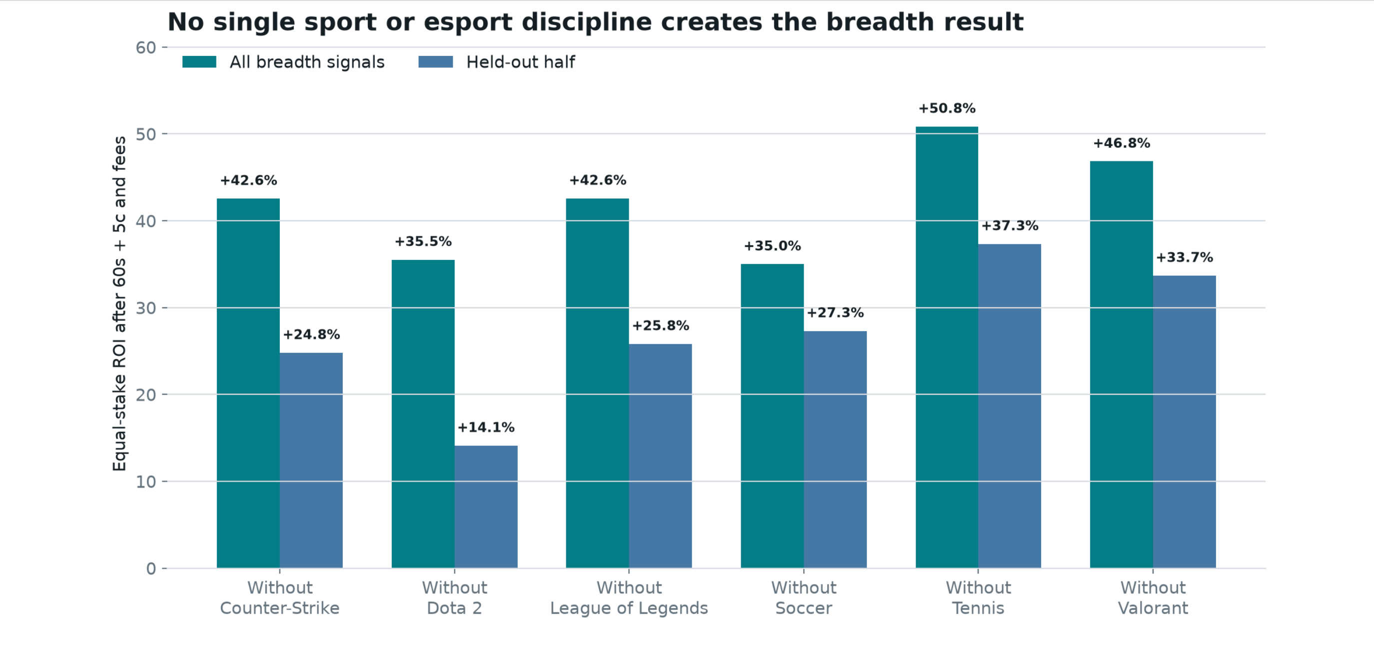

Each pair removes every breadth signal from the named discipline and reruns the same fixed rule. This guards against one category carrying the result; it does not create six independent tests.

The weakest held-out leave-one-discipline-out result was still +14.1% after excluding Dota 2. No single sport or esport category creates the aggregate sign.

What the equity curve and drawdowns actually look like

Alpha is not the same as a smooth return stream. Losing contracts still lose the entire $100 stake, and a short sample can look stable by accident.

RISK VIEW	BLIND COPY	18+ MAKERS	HELD OUT
Bets	139	30	21
ROI	-6.2%	+41.9%	+27.3%
Maximum drawdown	$2,056	$172	$172
Longest losing streak	7	1	1
Worst rolling five bets	-100.0%	-14.4%	-14.4%
Profitable trading days	20 / 42	12 / 15	7 / 10
ROI after top five winners	-13.5%	+20.1%	-3.8%

The most important risk caveat: the full 30-trade breadth sample remains positive after removing its five best winners, but the 21-trade held-out half does not. That does not erase the discovery; it shows why 200 genuinely new paper signals are required before any live-money claim.

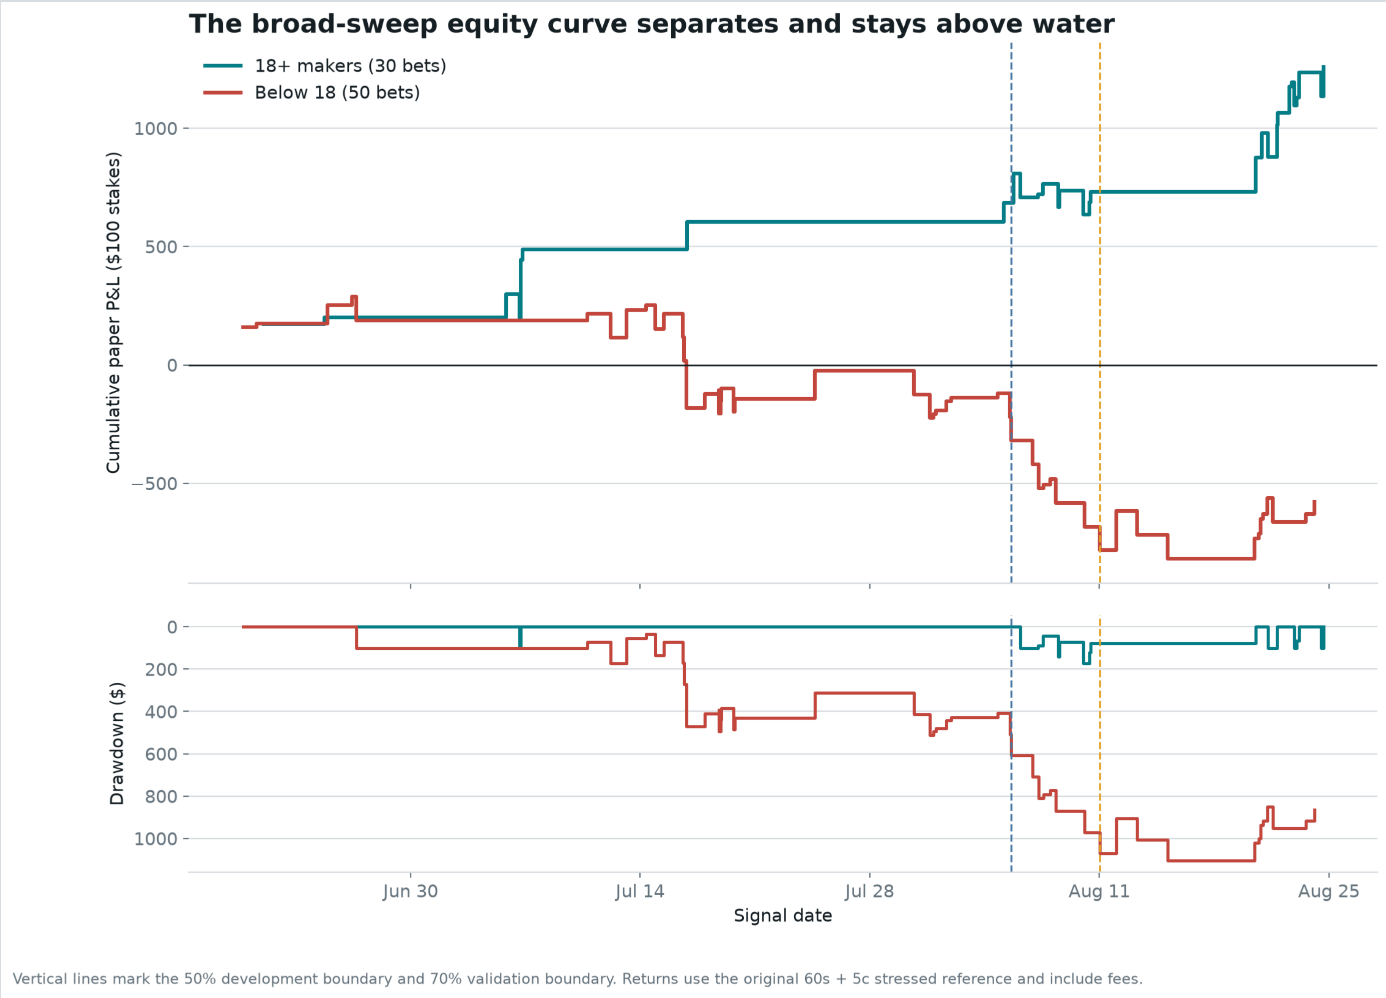

The broad curve separates early and stays positive through validation and final test. Narrow triggers deteriorate sharply. Vertical lines show the development and validation boundaries.

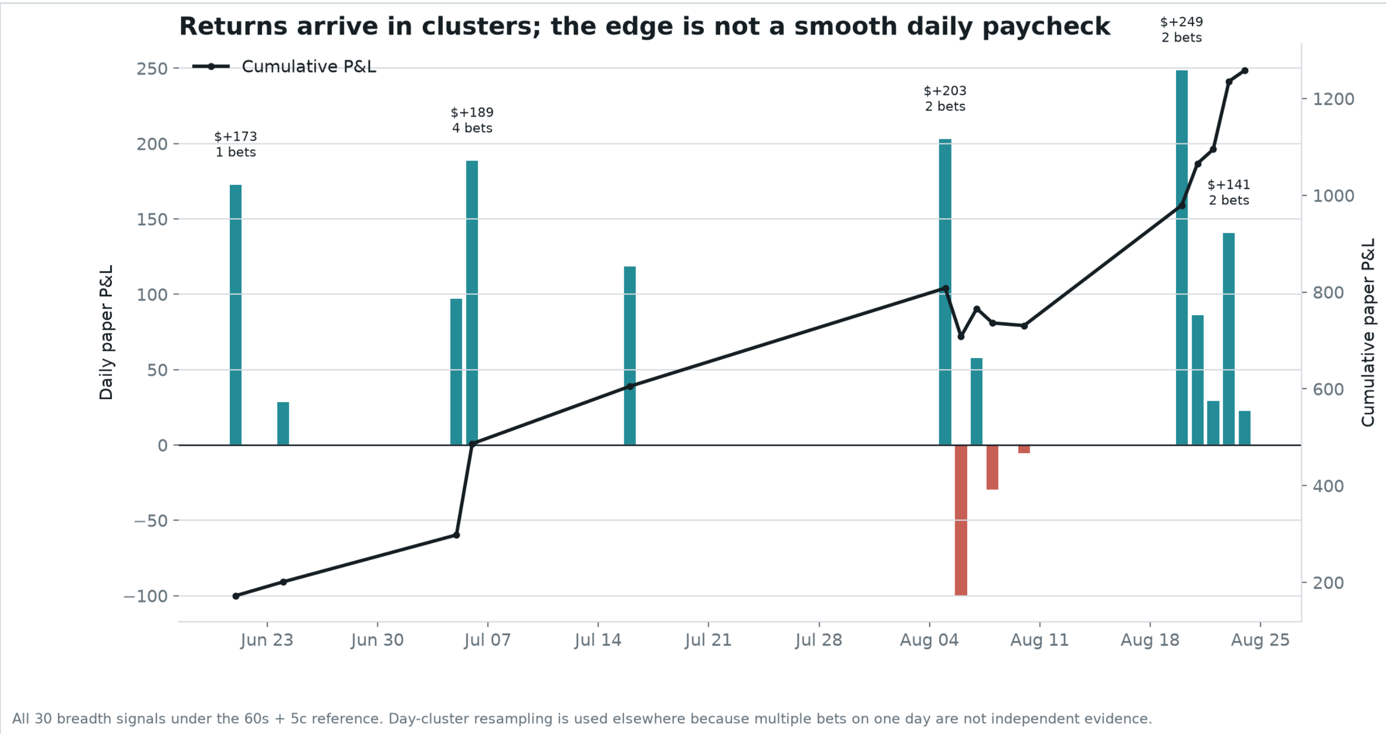

Profits arrive on clusters of signal days rather than evenly. That dependence is why the report resamples whole trading days instead of pretending every event is independent.

At the original 60-second mark, the breadth rule remains positive as the assumed price penalty rises. At ten cents, all broad signals return +29.3% and the held-out half returns +16.4%. The held-out line reaches break-even near 20 cents.

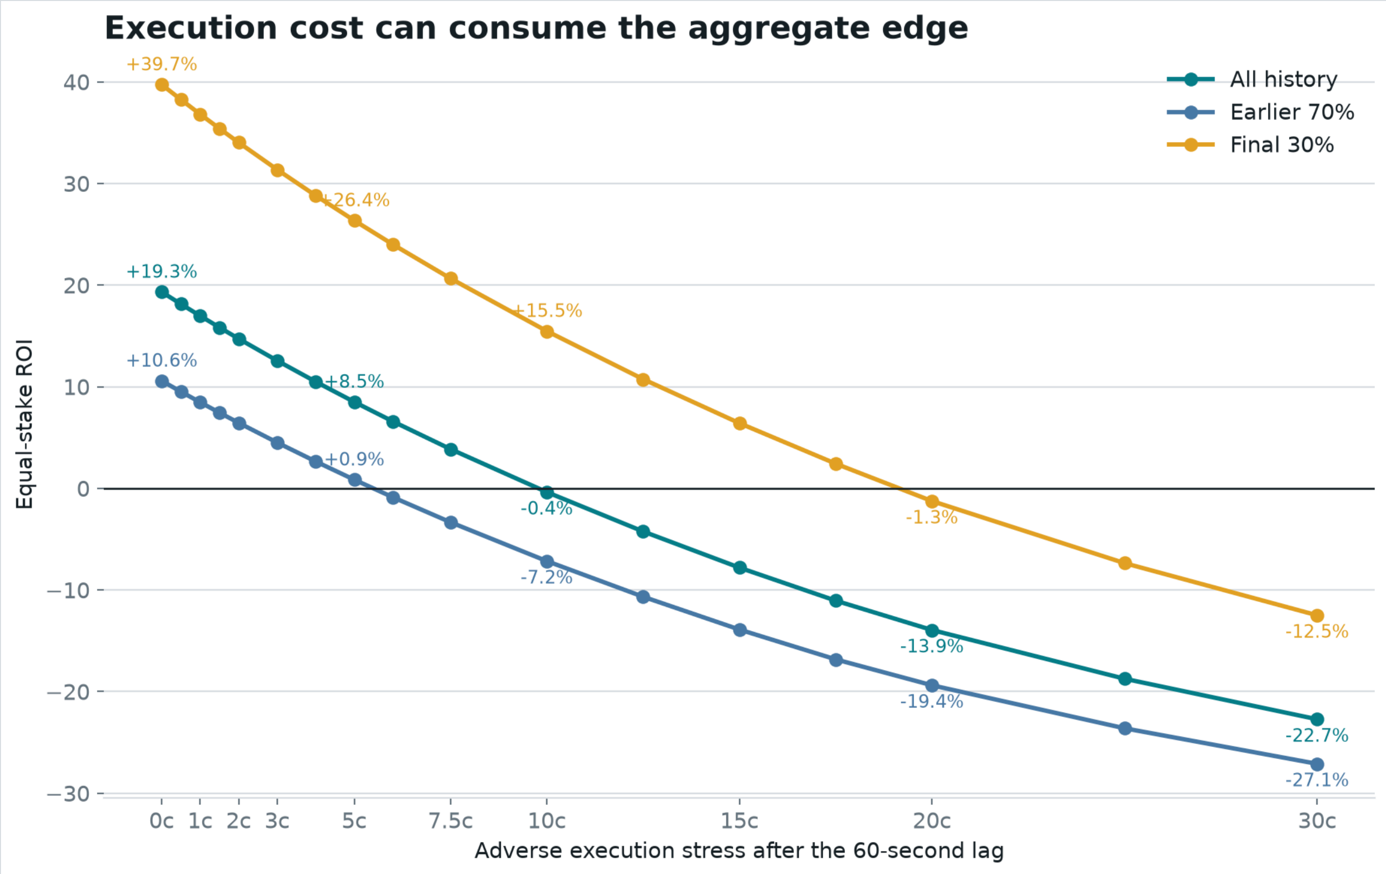

The wider eligible universe is much less forgiving. This is why a realistic copy test must report price impact rather than assuming the bot receives the whale's fill.

What the alpha reveals, and what remains hidden

CHAIN FACT	**The wallet selectively takes broad liquidity.** The triggering calldata directly names the maker orders consumed. Breadth is observable once the transaction is mined.
SAMPLE FACT	**Those broad sweeps contain the forecasting value.** They won above public-price expectations; narrower triggers did not. Dollar size does not account for the difference.
INTERPRETATION	**The behavior looks like informed liquidity demand.** The trader is willing to clear many standing offers immediately when waiting appears more dangerous than crossing the book.
SOURCE	**The information source remains invisible.** Public data cannot distinguish a superior sports model, faster public feeds, private information, coordinated

research, or disciplined human judgment.

The discovery is not the trader's secret model. It is the best externally observable signature of when that hidden process is expressing unusually strong conviction.

Why the rest of the wallet still matters

- Maker orders were 92.8% of fills but only 27.9% of dollars. Fill count alone exaggerates routine activity.
- True multi-map series earned $7,569,468 at +31.2%. Single-game and map contracts lost $4,888,740.
- The first visible buy poorly predicted eventual position size; the chronological sizing model had out-of-sample R-squared -0.108.
- The whale can size, make markets, sell, and manage inventory. A follower can reproduce none of that by copying a public feed row.

The honest reasons for doubt

- **Thirty is small.** The full breadth result has 30 bets; only 21 occurred after threshold selection.
- **The held-out win-count test is suggestive, not decisive.** Its one-sided Poisson-binomial p-value is 0.086.
- **The threshold-search simulation is barely below 0.05.** Its p-value is 0.046, not overwhelming evidence.
- **The 1,444 atlas cells are not 1,444 confirmations.** They map sensitivity around one frozen rule. Treating the best cell as a new discovery would be parameter mining.
- **The wallet and feature family were chosen retrospectively.** The null simulation corrects the stated maker-count cutoff search, not every idea considered during the investigation.
- **The history spans roughly two months.** A market regime, one operator, or one sports calendar can change.
- **The held-out result is winner-concentrated.** Removing its five most profitable winners changes held-out ROI to -3.8%.
- **Execution is still a proxy.** Public prints show activity, not guaranteed historical ask depth, queue position, partial fills, or API publication latency.

> **Correct conclusion:** this is a real, testable discovery in the available sample, not proof of future profit. It earns a frozen prospective paper test. It does not earn live capital yet.

Facts, interpretation, and speculation

CLAIM	CLASSIFICATION	WHY
Blindly copying every large signal lost money.	**Observed fact**	139 forced bets under declared execution assumptions.
All trigger transactions made the wallet the BUY taker.	**On-chain fact**	149 of 149 decoded V2 transactions.
Sweeps across 18 or more makers carried the edge.	**Observed in sample**	23 wins from 30, +41.9% ROI, with a positive held-out half.
Blind copy bots must be under one second.	**Not supported**	No sub-minute latency cliff appeared; roughly two cents of adverse price erased the blind-copy margin.
Breadth is only a disguise for dollar size.	**Not supported**	Notional was null after breadth, urgency, and period controls: p=0.858.
The trader has private information.	**Not established**	Behavior is consistent with informed trading, but public data cannot identify the source.
The algorithm is ready for live money.	**No**	Short history, 21 post-selection bets, and proxy execution still leave material uncertainty.

The honest verdict

The discovery is not “copy this whale.” That strategy lost.

His edge, as far as public evidence can reveal it, is knowing when to demand a lot of liquidity from a lot of counterparties at once. The breadth of the atomic sweep is the footprint. The hidden model or information that causes him to do it remains private.

The frozen algorithm is **atomic-breadth-18**. Historical result: 23 wins from 30, +41.94% after the original stressed costs. Post-selection result: 15 wins from 21, +27.32%.

For an indiscriminate bot, the historical break-even execution allowance at one second was only 1.53 cents. The lesson is not “buy a faster server.” It is “reject weak signals and control the paid price.”

Decision: collect at least 200 new eligible signals in paper mode without changing the rule. Record actual order-book depth, rejected orders, partial fills, and end-to-end latency. Capital should wait until a genuinely unseen sample remains profitable after costs and after removing its biggest winners.

Small glossary

Maker

A trader who leaves an order waiting in the book and provides liquidity.

Taker

A trader who accepts an available price immediately. Takers usually reveal more urgency and pay more.

Atomic sweep

One mined transaction in which a taker order matches several already-signed maker orders together.

Maker breadth

The number of distinct maker addresses consumed by that transaction. It counts accounts, not verified people.

Adverse price

How many extra cents a follower pays above the first public execution reference because of spread, consumed depth, or market reaction.

Implied probability

A 60-cent contract roughly represents a market-estimated 60% chance before trading costs.

Calibration gap

The actual win rate minus the win rate implied by prices. Positive is good only if it survives costs and honest testing.

ROI

Profit divided by total stake. A 10% ROI means $10 profit for each $100 staked.

p-value

A narrow measure of how often a result at least this large appears under a specified random null. It is not the probability that the strategy is true.

Confidence interval

A range showing statistical uncertainty under a particular resampling method. It does not include every source of bias.

Where the facts came from

This essay is a human-readable rendering of committed machine-readable evidence. The core calculations can be audited in:

- edge_analysis.json: blind-copy and atomic-breadth backtests, chronology, null simulation, controls, 1,444 parameter-atlas cells, risk diagnostics, and break-even frontiers.
- edge_features.csv: one row per reconstructed signal, including decoded maker breadth and information available by signal time.
- trigger_transactions.json: decoded `matchOrders` calldata and derived fill anatomy for all 149 trigger transactions.
- deep_analysis.json: wallet-level execution, market format, profit, fee, and concentration facts.
- market_tape.json: 143,507 unrelated public prints used for execution proxies and market response.

External context:

- Official Polymarket CTF Exchange V2 repository, including `matchOrders` execution and the signed order structure.
- Official Polymarket order lifecycle, distinguishing off-chain `MATCHED`, on-chain `MINED`, and final `CONFIRMED` trade states.
- Official public market WebSocket, documenting millisecond-stamped book and trade events available prospectively but absent from this historical second-resolution tape.
- Dubach, The Anatomy of a Decentralized Prediction Market, supporting use of authoritative on-chain direction rather than potentially ambiguous public labels.
- Easley and O'Hara, Price, Trade Size, and Information in Securities Markets, classic context for why aggressive trade structure can contain information.
- Bailey et al., The Probability of Backtest Overfitting.

Limit: this is public-wallet research, not proof of identity, distinct human counterparties, private information, causality, historical executable depth, or future profit. It is not financial advice.