PLAIN-ENGLISH ALPHA AND EXECUTION DOSSIER

Inside the whale's alpha: what copies, what fails, and the esports lead.

A transaction-level investigation, 6,244 execution-and-capacity scenarios, live esports telemetry, wallet-aligned CS2 video evidence, and independent tests that separate the trader's mechanism from a tradable rule.

Base backtest wallet sample covering June 19, 2026 to August 25, 2026. Prospective wallet, market, and live-match evidence updated August 27, 2026.

-6.15% return from blindly copying all 139 large signals	**+41.94%** retrospective return when one trigger matched at least 18 maker accounts	**-73.46%** return from blindly mirroring large passive clusters with 18+ counterparties

The answer in thirty seconds

The wallet has two observable execution modes: **it aggressively consumes dense liquidity when conviction is high, and it sometimes posts size around a live state that it believes is mispriced.** The CS2 reconstruction identifies the second mode directly. It does not turn it into an automatic strategy.

Bottom line: during M80-NAVI, the wallet first bought at 74 cents while the broadcast showed M80 leading 9-6 with a planted bomb and a 5-v-3 advantage, then absorbed $49,990 at 78 cents from 19 counterparties after M80 made it 10-6. It earned $18,382. But the same-team G2-M80 control used almost the same passive mechanism at 12-11, 3-v-3, no bomb, and lost $83,606. Across all qualifying passive clusters, the mechanical rule returned **-73.46%**. **The best-supported mechanism is match-state valuation plus selective context and passive execution; the exact event-selection and fair-value model remain unrecovered.**

1. **The headline is misleading.** The wallet made $5,644,317, but its top five winners produced 144% of net profit. Removing those five makes the wallet negative.

2. **Blind copying is not the edge.** 139 delayed $100 copies lost $855.28. The later sample also lost -6.68%.

3. **The public trade feed hides structure.** We decoded all 149 trigger transactions. Every one was a successful V2 `matchOrders` call with the target as the BUY taker.

4. **Breadth separated signal from noise.** Broad sweeps beat the market price by +23.2 points; narrower transactions missed it by -0.6 points.

5. **Esports is important, but not uniquely proven.** It carried 58.6% of wallet cost basis and $2,808,250 of realized P&L. Frozen broad sweeps returned +40.1% in esports and +43.7% in traditional sports.

6. **A wallet-linked CS2 decision is now reconstructed to the round.** The first M80 fill preceded the broadcast's 10-6 update by about 7.5 seconds, but it did not precede the favorable state: viewers could already see the bomb planted and M80 ahead 5-v-3.

7. **Copying the passive footprint is decisively wrong.** A fixed 18-counterparty mirror found 4 wins from 9 and lost 73.46%; every tested cutoff from 10 to 30 counterparties was negative.

8. **A valid Dota model did not become a valid trading strategy.** It reached ROC-AUC 0.851 on 1,499 later professional matches, then lost 6.87% in the separate market-wide replay.

9 **Size changes copyability.** At +1 cent, a current liquid-market snapshot could fully fill 94.7% of $100 FOK orders but only 26.2% at $25,000. One second after the historical broad sweeps, even an optimistic all-print ceiling covered only 38.1% of $100 requests.

10 **The market did not confirm the story before play.** Broad pregame signals had median closing-line value of -0.67 cents, with only 4 of 12 positive.

11 **The first frozen prospective audit has no bet to score.** 875 new wallet trades produced 7 raw threshold signals, 2 in the frozen universe, and zero that reached 18 makers.

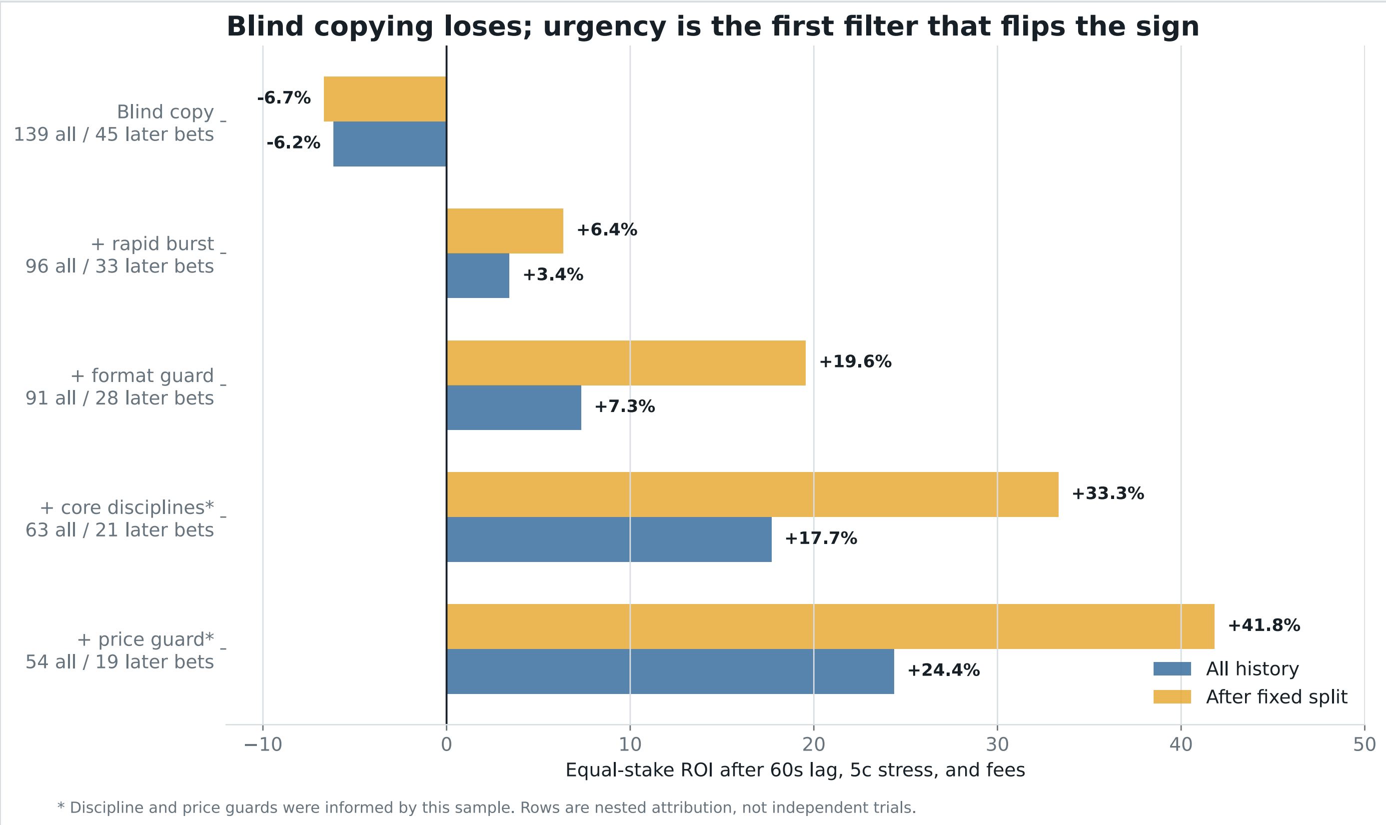

What blind copying would have done: the first row is the naive strategy. It loses over the full history and after the fixed date split. Earlier behavioral filters found useful clues, but they did not expose what was inside the triggering blockchain transaction.

First, the painful result: copying loses

The original registered test was deliberately harsh, but ordinary: it did not give the follower the whale's old price.

1 Wait until the target has crossed **$25,000** of aggressive buying and at least 70% of its net direction points to one outcome.

2. Keep only the first signal for each underlying match, so a series and its individual maps cannot be counted as independent ideas.
3. Wait **60 seconds**, then use the first unrelated public trade in the following minute as the available price. If none appears, retain the trigger price.
4. Make the price five cents worse, include the account-observed fee curve, and place the same **$100** stake every time.

The result was 80 wins and 59 losses, yet -$855.28 net profit on $13,900 of turnover. The maximum historical drawdown was $2,055.76.

> **This is what blind faith costs:** even a spectacular trader can be a bad signal to copy. The follower sees the action late, pays a worse price, does not know final position size, and cannot reproduce inventory management, maker fills, exits, or rebates.

Why winning more than half the bets was not enough

Prediction-market prices already contain a probability. Buying a contract at 60 cents means paying roughly as if it has a 60% chance to win. A strategy can win 58% of its bets and still lose if it repeatedly pays prices that require a higher hit rate.

Blind copying won 80 times. The public prices implied about 77.15 wins. The gap was only +2.0 points. In ordinary language, the extra wins were neither large enough nor valuable enough to establish a generic “copy his big buys” edge.

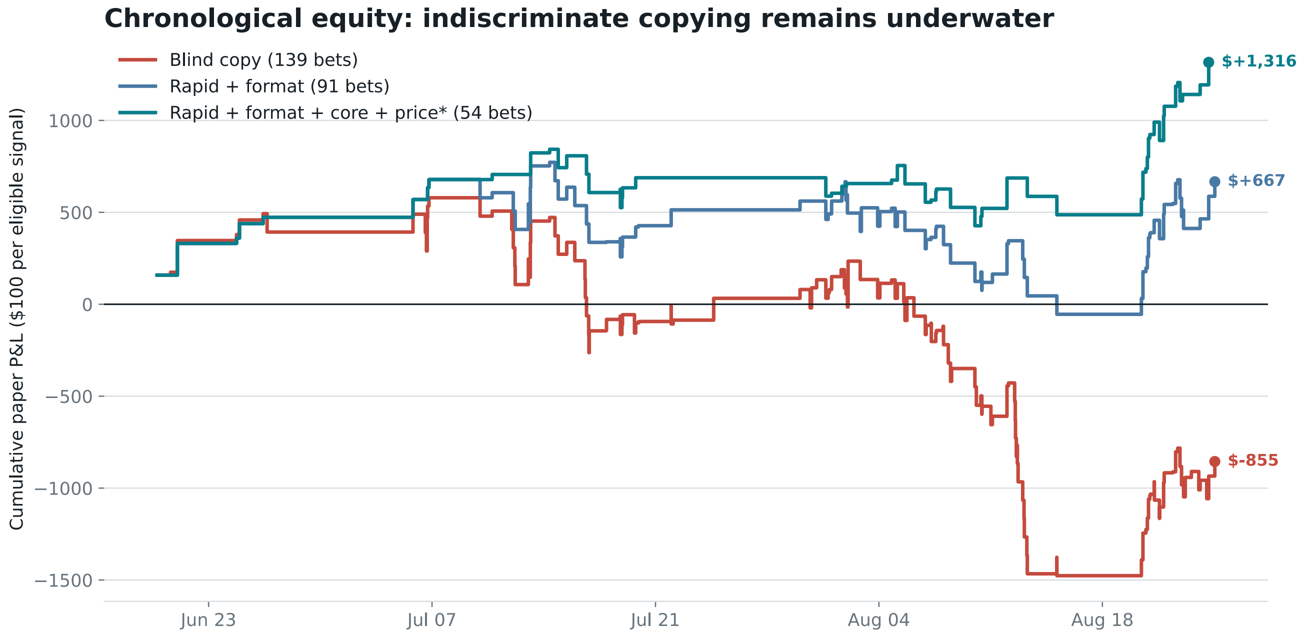

* Core-discipline and price filters are exploratory. Curves use each rule's own chronological event set.

The red blind-copy line remains underwater. More selective rules improve because they reject most trades. This was the clue that the wallet's value lives in selection, not in its address.

What if a copy bot is nearly instant?

Speed helps, but the broad replay finds no sharp 1-second-versus-5-second cliff. The dangerous variable is the price paid after the whale has consumed the available offers.

The expanded audit crosses **15 delays**, from the trigger's own second through five minutes, with **17 adverse-price assumptions** from zero through 30 cents. That is 510 latency-by-price results across blind and alpha-filtered copies before the fee and maker-threshold atlases. Every primary cell includes the observed fee curve. At the most optimistic same-second price with no extra adverse movement, blind copying returns only +3.4%. At one second plus one cent it returns +1.0%. At one second plus two cents it is already negative at -0.9%.

The practical threshold: at one-second latency, blind copying historically breaks even at only about 1.53 cents of additional adverse price. Under one second plus one cent, the breadth-filtered held-out sample returned +35.6%. Being fast is not enough if the whale just removed the cheap liquidity; selecting the right trigger matters more.

The timestamps impose an important limit. Polygon block timestamps and the historical public tape are recorded to whole seconds. A 0.1-second bot and a 0.5-second bot cannot be separated honestly. The same-second scenario may also include unrelated prints whose exact within-second order is unknown, so it is an optimistic lower bound rather than a reproducible fill.

Polymarket's official lifecycle separates an off-chain `MATCHED` state from `MINED` on-chain settlement. This strategy needs the mined `matchOrders` calldata to count maker addresses, so its clock starts at the block timestamp. A bot with private or earlier CLOB-match information is testing a different signal and is not represented by this public-wallet replay.

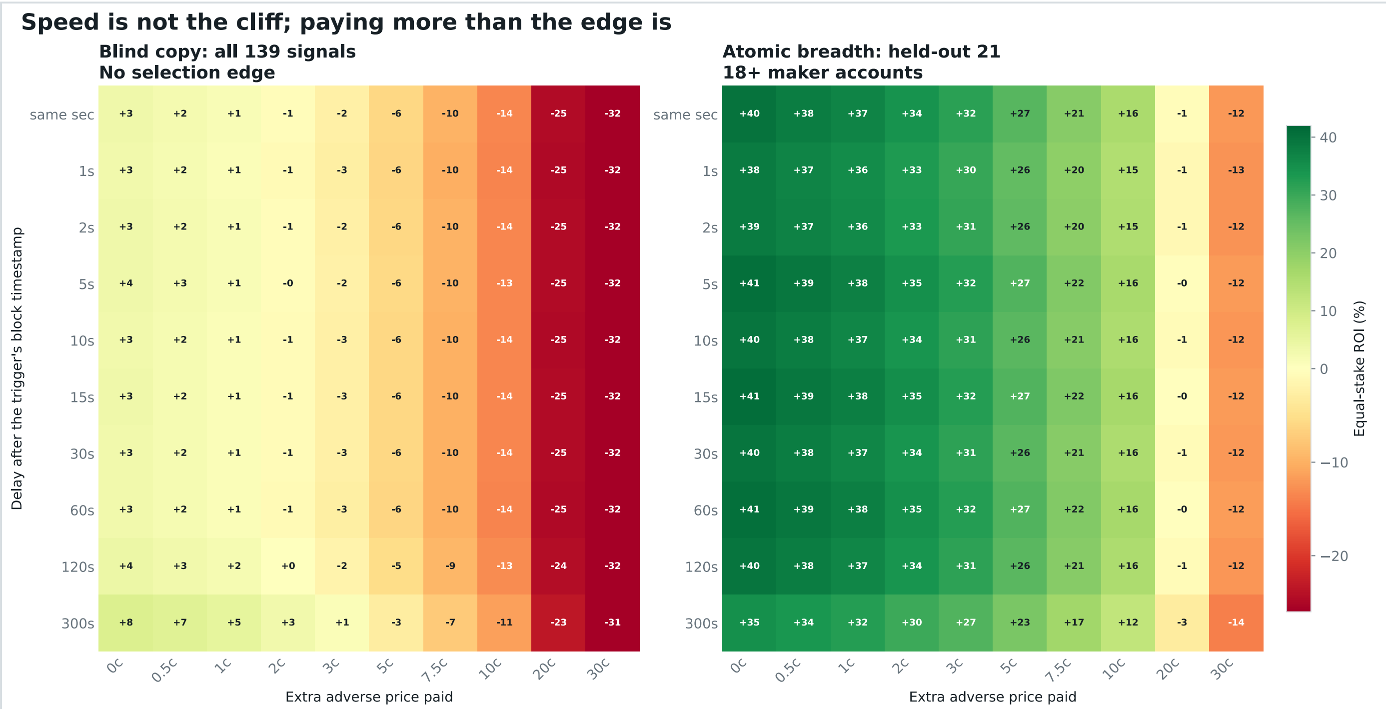

Selected anchors from the full 15-delay x 17-cost atlas. Same-second is optimistic because ordering inside a one-second timestamp is unknown. A 0.1s or 0.5s bot lies between same-second and 1s. Every cell includes the observed fee curve.

Read across for execution quality; read down for speed. Blind copying flips from pale green to red at roughly two cents in almost every sub-minute row. The held-out breadth strategy remains green much farther across. A 0.1-second or 0.5-second bot lies between the first two rows because the source timestamps are only one second precise.

Blind copying has about two cents of room; breadth has about twenty

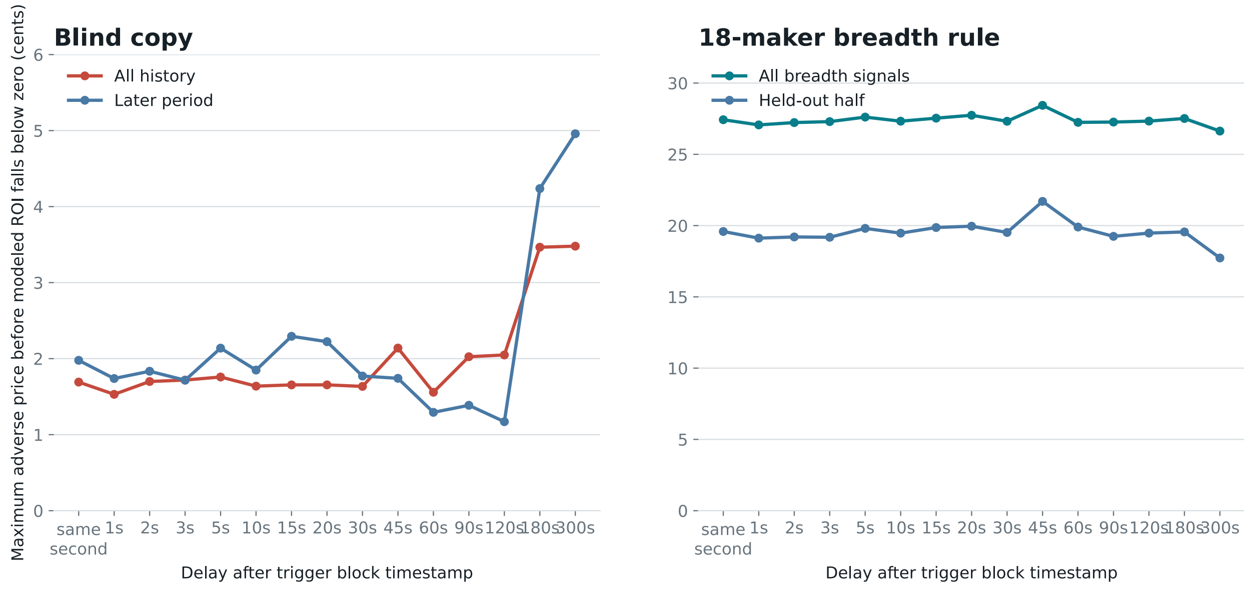

Break-even values are solved from the same equal-stake replay and include fees. The near-flat lines show no measured sub-minute latency cliff; exact historical order-book depth remains unavailable.

The blind copier has roughly 1.5 cents of room at one second. The 18-maker filter has roughly 19.1 cents in the held-out half. The near-flat lines mean this dataset does not support a claim that sub-five-second speed is the main edge.

The full copy-trading parameter atlas

A backtest should show the whole neighborhood, not one convenient assumption. This atlas varies delay, price deterioration, fees, and the maker-breadth cutoff while keeping equal stakes and event deduplication fixed.

15 delays from same-second through 300 seconds	**17** adverse-price assumptions from 0c through 30c	**6** fee-curve rates from 0% through 5%	**26** maker-breadth cutoffs from 5 through 30

ONE-SECOND SCENARIO	BLIND COPY	18-MAKER HELD OUT
No extra adverse price	+3.1%	+38.4%
One cent worse	+1.0%	+35.6%

ONE-SECOND SCENARIO	BLIND COPY	18-MAKER HELD OUT
1.5 cents worse	+0.1%	+34.3%
Two cents worse	-0.9%	+32.9%
17.5 cents worse	-22.5%	+2.3%
Twenty cents worse	-24.9%	-1.2%

> **What the spectrum says:** blind copying is a thin-margin trade that flips near 1.5 to 2 cents. The held-out alpha filter does not flip until roughly 19 to 20 cents. The major difference comes from selecting the transaction, not shaving a few seconds from the bot.

Latency barely moves the result; execution price decides it

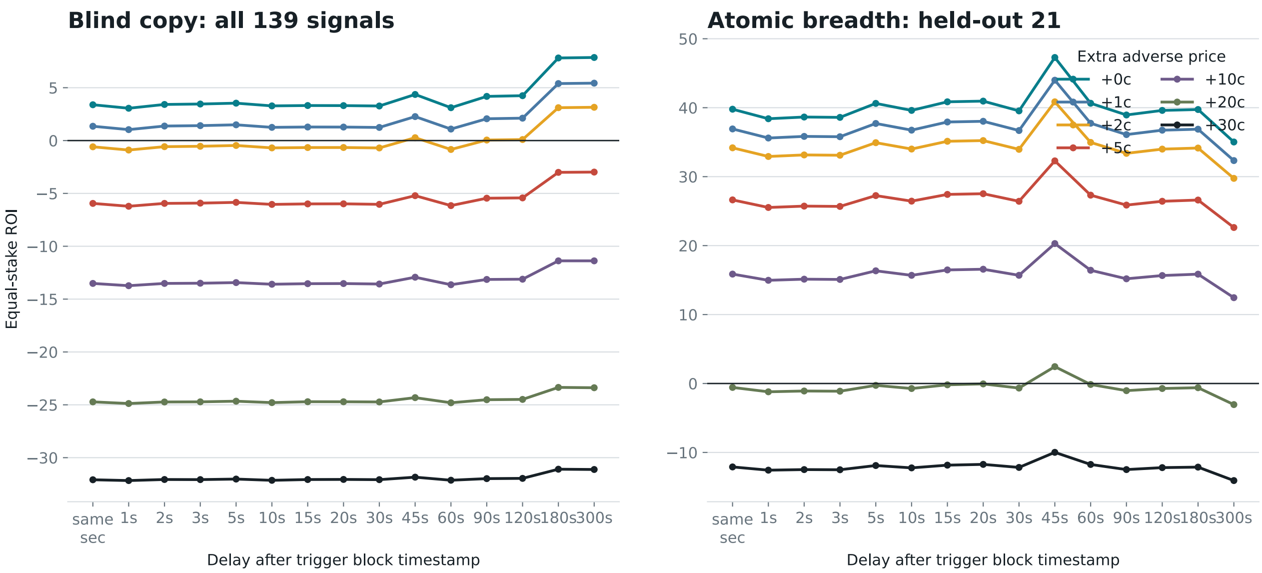

All 15 measured delays are shown. Lines are nearly horizontal below one minute. The historical data cannot rank 0.1s against 0.5s inside the same timestamped second.

All fifteen measured delays: each line holds the adverse price constant. The lines barely move through the sub-minute region, while moving between price-cost lines changes the answer immediately.

The whole cost curve: blind copying dies near 2c, breadth near 20c

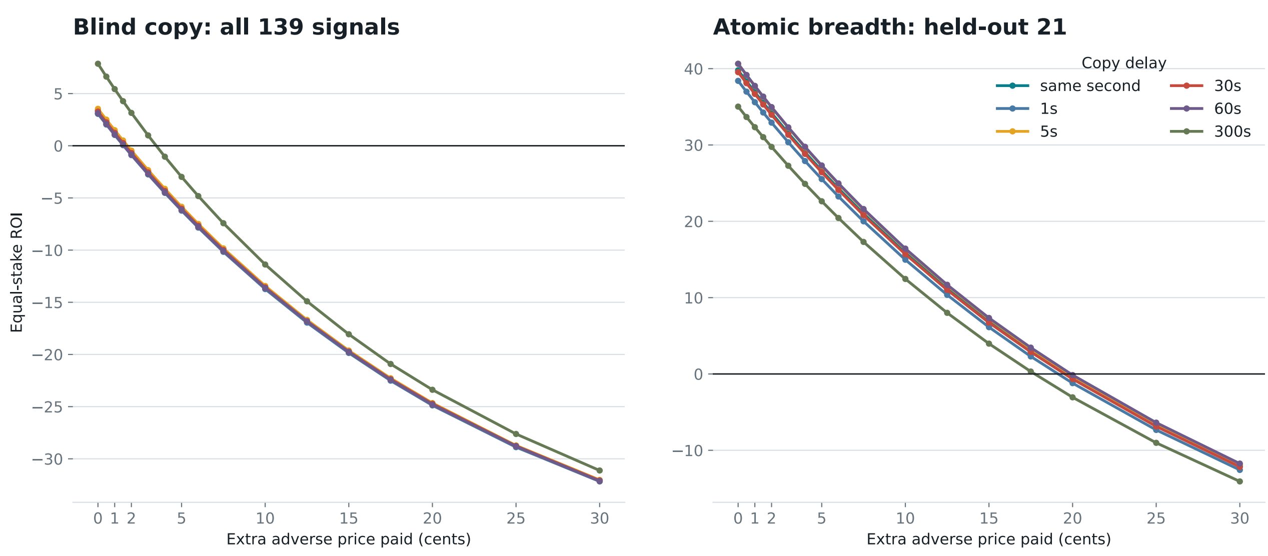

Seventeen explicit adverse-price assumptions from 0c to 30c are shown. These are stress scenarios around public prints, not reconstructed historical order-book fills.

All seventeen cost assumptions: blind-copy curves cross zero near two cents regardless of delay. The held-out breadth curves cross near twenty cents.

Fees matter, but price selection remains the dominant execution variable

One-second replay. Each all-in price adds fee_rate x price x (1-price). The observed account curve uses 3%; other rows are explicit counterfactuals.

Fees shift the result gradually. Paying through the book shifts it much faster. The primary replay uses the account-observed 3% fee curve; every other fee row is a labeled counterfactual.

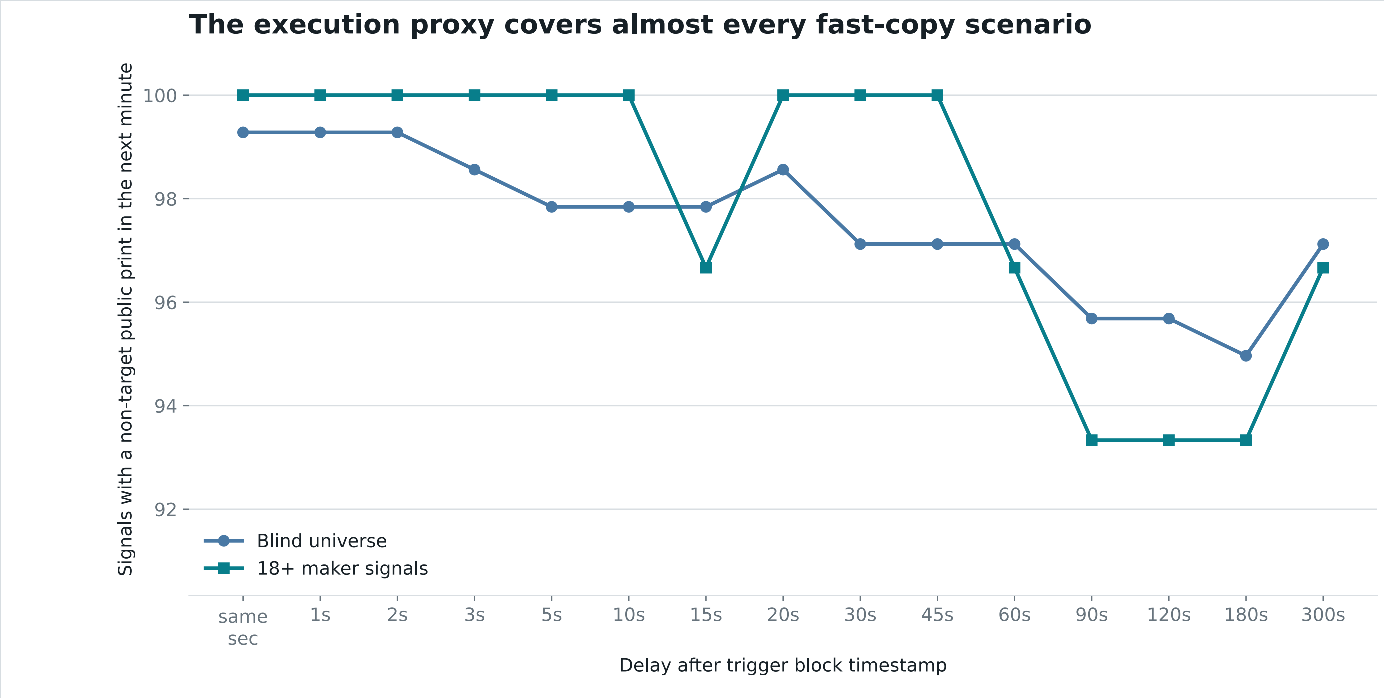

Missing-print events are retained at the trigger price rather than dropped, preventing an easy upward selection bias. Coverage is not proof that $100 of ask depth was executable.

Fast scenarios still have roughly 99% public-print coverage. Missing prints are kept with a fallback instead of being dropped, but a public print still does not prove historical ask depth or queue position.

The missing dimension: can the copy actually fit?

A backtest that changes price but assumes every requested dollar fills is incomplete. Size determines whether an FOK order trades at all, and the whale has usually consumed the most obvious liquidity before a public follower can react.

We added **4,800 capacity cells** on top of the 1,444-cell parameter atlas: five accumulation windows, four price buffers, four participation limits, ten requested stake sizes, two turnover proxies, and three strategy samples. Together the report now evaluates **6,244 execution-and-capacity scenarios**.

206 current token-side books in the liquid sports sample	**21** historical held-out breadth signals	**$8,717** median displayed depth through +1c today	**4,800** stake, speed, buffer, proxy, and participation cells

REQUESTED STAKE AT +1C	CURRENT IMMEDIATE FOK	POST-SWEEP 1S ALL PRINTS	POST-SWEEP 60S ALL PRINTS	POST-SWEEP 60S ALIGNED BUYS
$25	96.6%	52.4%	100.0%	76.2%
$100	94.7%	38.1%	85.7%	61.9%
$1,000	84.5%	9.5%	57.1%	14.3%
$10,000	46.1%	4.8%	23.8%	4.8%
$25,000	26.2%	0.0%	14.3%	0.0%

> **Do not merge these columns into one promise.** Current FOK is executable displayed depth, but it was sampled now from high-volume moneylines and not after this whale traded. Historical columns are cumulative public turnover after the signal, not simultaneous asks. The all-print column is deliberately optimistic because it treats direction-neutral prints as replaceable capacity.

Displayed depth falls away as requested copy size rises

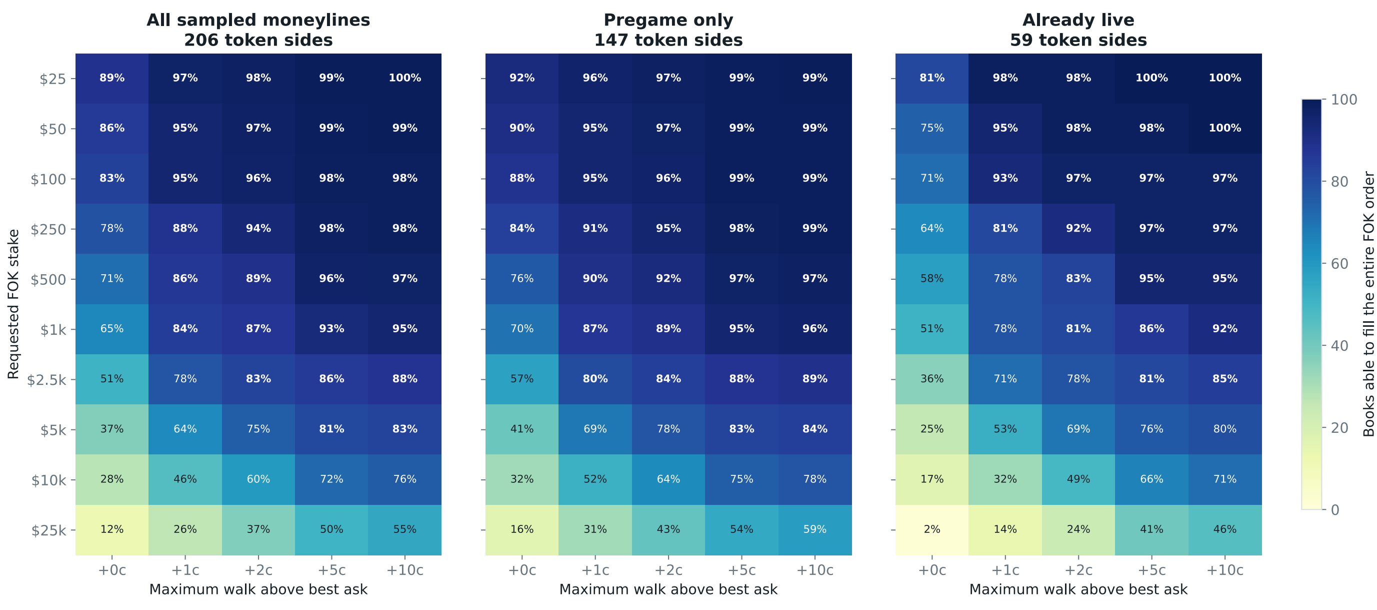

Current cross-section of high-volume sports moneyline books from the official CLOB API. A fill means displayed asks could satisfy the whole order; it is not a post-signal backtest and the sample favors liquid markets.

What size does to a real book: $100 through +1 cent fit in 94.7% of the sampled token sides; $10,000 fit in 46.1%; $25,000 fit in only 26.2%. Larger buffers raise fill coverage by authorizing worse prices.

Size creates rejection risk before it creates large measured slippage

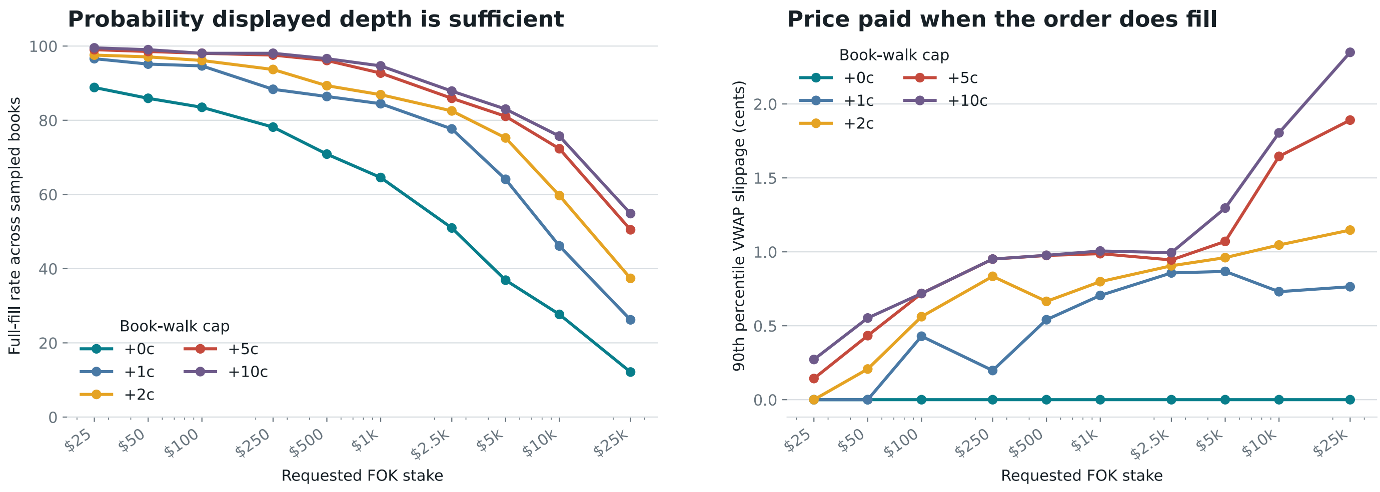

VWAP percentiles condition on complete fills, so rejected books disappear from the right panel. The apparent slippage ceiling is therefore not evidence that a large order can trade everywhere.

The left panel contains the main size penalty: rejected orders. The right panel conditions on books that could fill completely, so the modest VWAP numbers cannot be read as proof that large orders are generally easy to execute.

A liquid market before the whale trades is not liquid after the whale trades

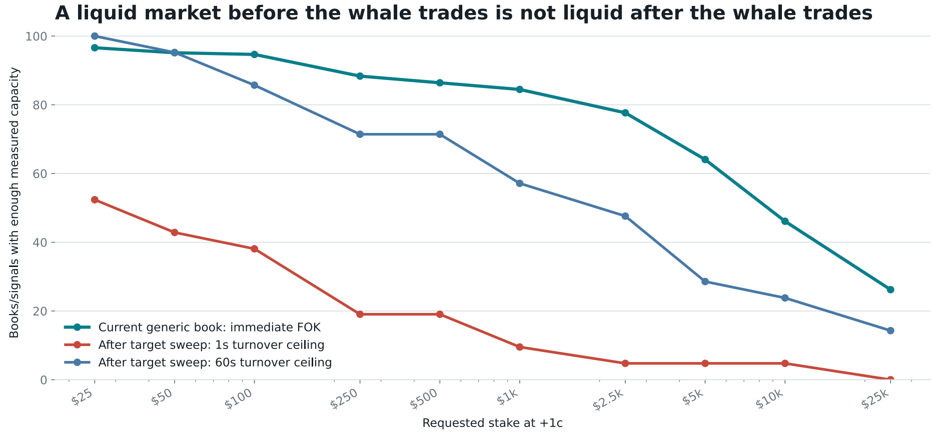

The green line is current displayed ask depth in a favorable top-volume sample. Historical lines are cumulative prints after 21 held-out broad sweeps and therefore optimistic ceilings, not FOK books. The comparison diagnoses the timing problem; it is not an apples-to-apples estimator.

The timing problem is visible. Generic liquid books often look deep before anyone trades. Immediately after a broad target sweep, the optimistic turnover ceiling is far lower because the first mover has already removed supply.

A realistic size projection is mostly a rejection projection

For the 21 held-out breadth signals, a $100 request at +1 cent found enough *total observed prints* within one second only 38.1% of the time. Waiting 60 seconds raised that optimistic ceiling to 85.7%. Restricting participation to 25% cut the 60-second figure to 71.4%. Using only reported aligned BUYs produced 61.9%.

At $1,000, the one-second all-print ceiling covered 9.5% of held-out signals. At $10,000, it covered 4.8%. Those numbers do not say the remaining trades partially filled: an FOK instruction rejects the whole order when displayed depth is short.

> **Practical paper rule:** calculate the ordinary risk cap, then reduce it to at most 10% of displayed ask notional through the one-cent limit. Reject the signal below $25. Walk every eligible ask level to estimate VWAP, submit no order in this repository, and record would-fill versus would-reject. This avoids pretending a $100 strategy scales linearly to $10,000.

After a broad sweep, historical public turnover is scarce at copy-bot speed

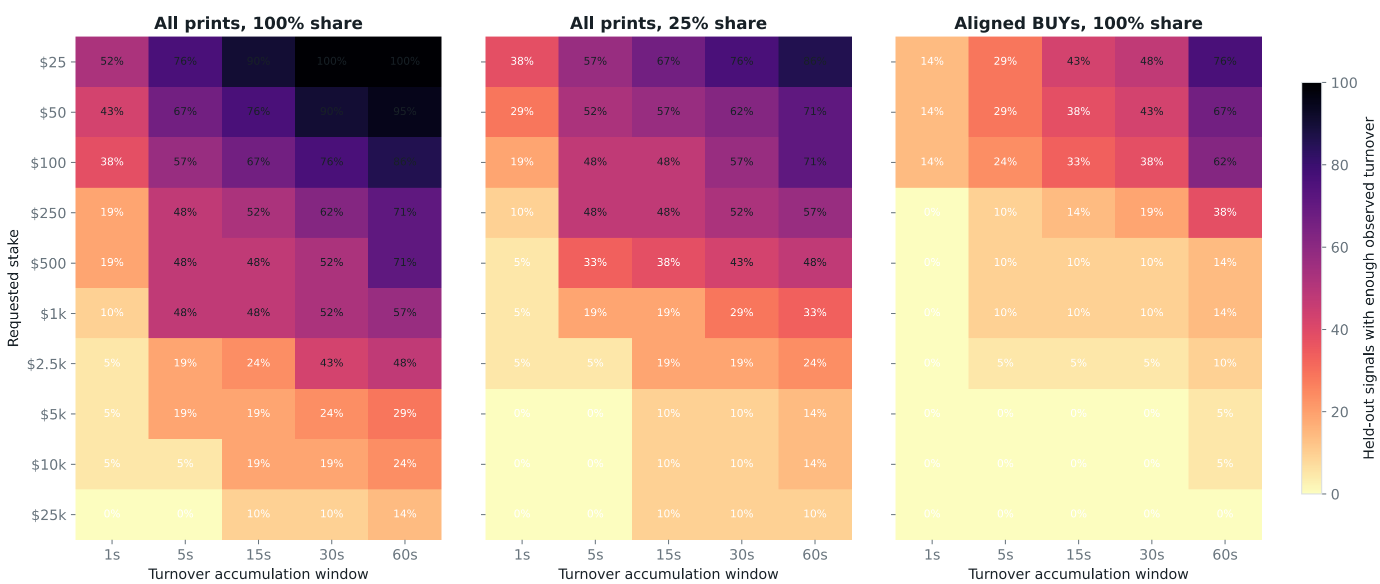

Twenty-one held-out breadth signals, +1c price buffer. Public prints accumulate over each window; they are not a simultaneous ask book. All-prints is an optimistic direction-neutral ceiling, while aligned BUYs is narrower but noisy.

Capacity disappears fastest at large stake and short delay. The middle panel limits the hypothetical follower to 25% of observed turnover. The right panel uses the narrower reported-side proxy. None is a reconstructed historical book.

Bigger requested stakes turn a strategy into a sparse fill lottery

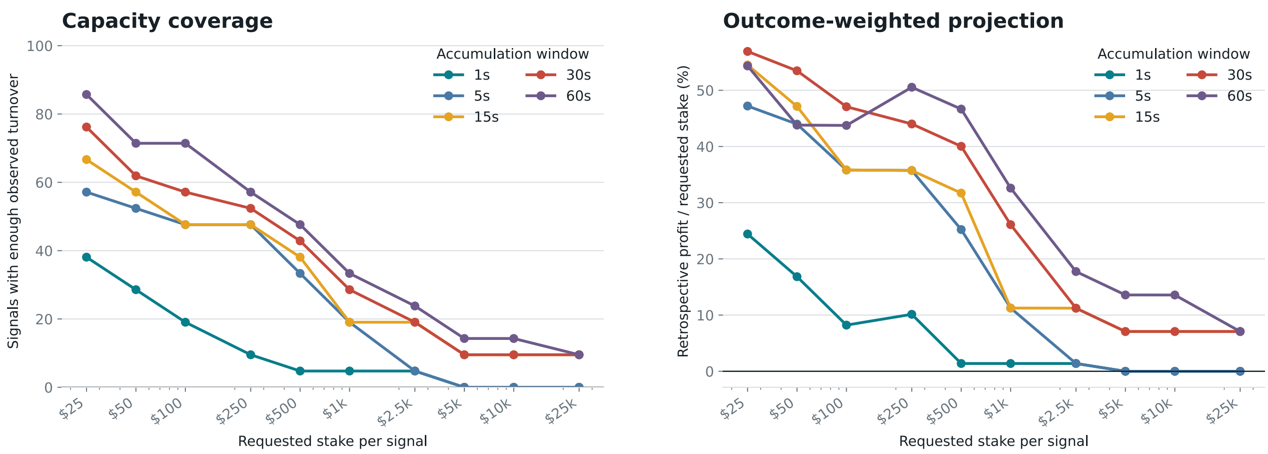

Held-out breadth signals, optimistic all-print turnover, +1c, and only 25% participation. The right panel uses known outcomes after selecting capacity-covered events; it is descriptive and selection-biased, not an investable return forecast.

The outcome-weighted panel is intentionally labeled as retrospective: capacity selects which known events remain in the sample. It is useful for diagnosing scale and selection effects, not for forecasting profit.

The breakthrough came from looking inside the trade

A public feed row says the whale bought one outcome. Polygon calldata says exactly how that buy was assembled.

Polymarket's V2 exchange accepts one taker order and an array of maker orders in a call named `matchOrders`. The taker demands immediate execution. Each maker had already signed an offer. One public-looking trade can therefore be an atomic sweep through many counterparties and several prices.

> **One public print is not necessarily one simple trade.** It can be one urgent order consuming a large section of the order book in the same mined transaction.

We fetched and decoded every one of the 149 trigger hashes. All 149 showed this wallet as the decoded taker buying the signaled token. The typical trigger matched 19 maker orders from 14 distinct maker accounts. 66 triggers crossed more than one price level. The reconstructed notional agreed with the signal to within less than one millionth of one percent.

> **Why this matters:** the final dollar amount says how much was bought. Maker breadth says how much standing liquidity the trader chose to consume at once. Those are different behaviors.

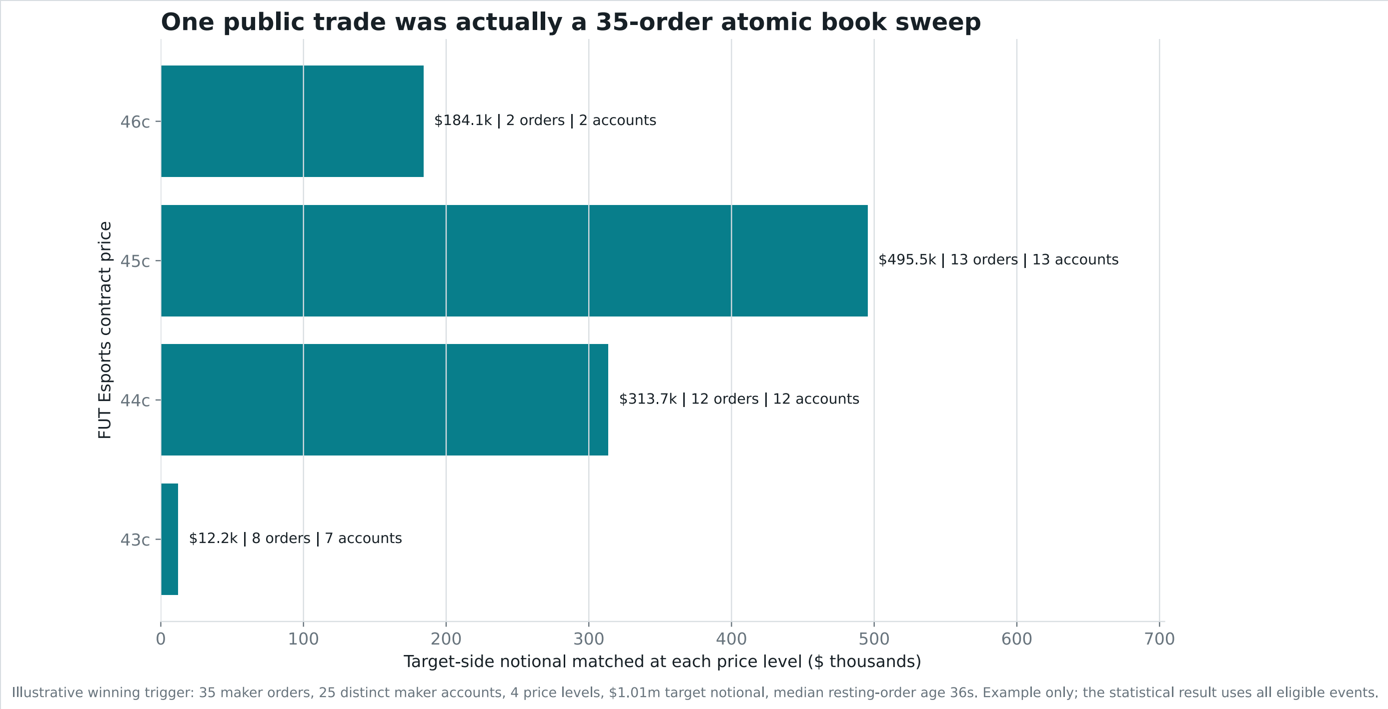

An illustrative winning FURIA trigger matched 35 maker orders from 25 distinct signed accounts across four price levels, from 43 to 46 cents. The target consumed about $1,005,518.93 of notional; the median resting order was 36 seconds old. This example explains the mechanism. The statistical result uses every eligible event.

The discovery: atomic breadth

The strongest observable fingerprint is simple: **did the triggering transaction match at least 18 distinct maker accounts?**

The comparison starts with the same conservative event universe: first canonical signal only, core tennis, soccer, and esports disciplines, full-event contracts rather than maps or short markets, target concentration of at least 70%, and trigger prices between 30 and 85 cents. That leaves 80 events. Breadth is the final split.

WHAT HAPPENED	18+ MAKERS	BELOW 18
Number of bets	30	50
Wins implied by public prices	16.04	29.30

WHAT HAPPENED	18+ MAKERS	BELOW 18
Actual wins	23	29
Price-implied win rate	53.5%	58.6%
Actual win rate	76.7%	58.0%
Gap versus price	+23.2 points	-0.6 points
Equal-stake return after costs	+41.94%	-11.56%

At a $100 equal stake, the 30 broad sweeps made $1,258.31. Even after removing the five most profitable winners, the remaining return was +20.1%. The narrower group lost $578.01.

> **The market-price test:** public prices expected about 16.04 wins; 23 occurred. The one-sided Poisson-binomial probability of at least that many wins under those market probabilities is **0.57%**. The same calculation finds no edge below 18 makers.

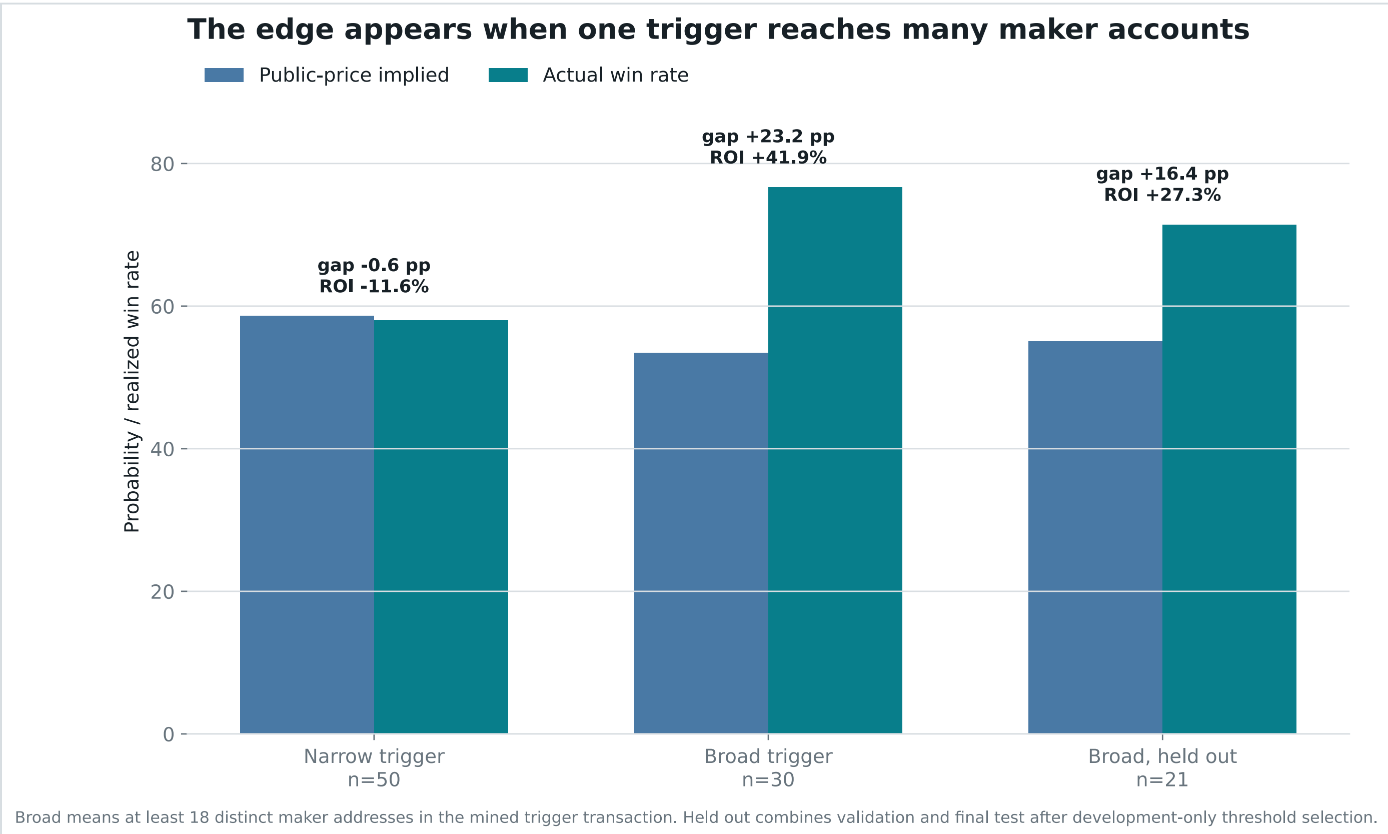

Public prices expected similar difficulty after filtering. Broad atomic sweeps won far more often than their prices implied; narrower triggers landed almost exactly where the market predicted.

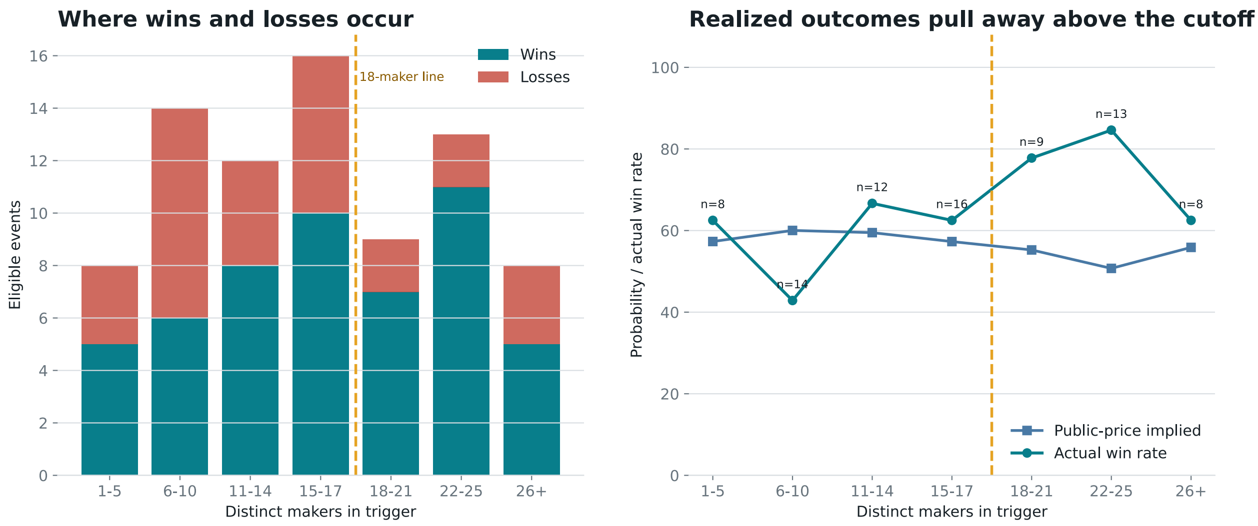

Maker breadth is a graded transaction feature, not a wallet-size label

Bands are descriptive. The exact cutoff of 18 was selected on the first half only; the right panel uses the original 60s + 5c execution reference.

The 18-maker rule sits inside a graded transaction pattern. Outcomes begin separating from public probabilities around the broad-sweep region; stricter bins eventually become too small to trust.

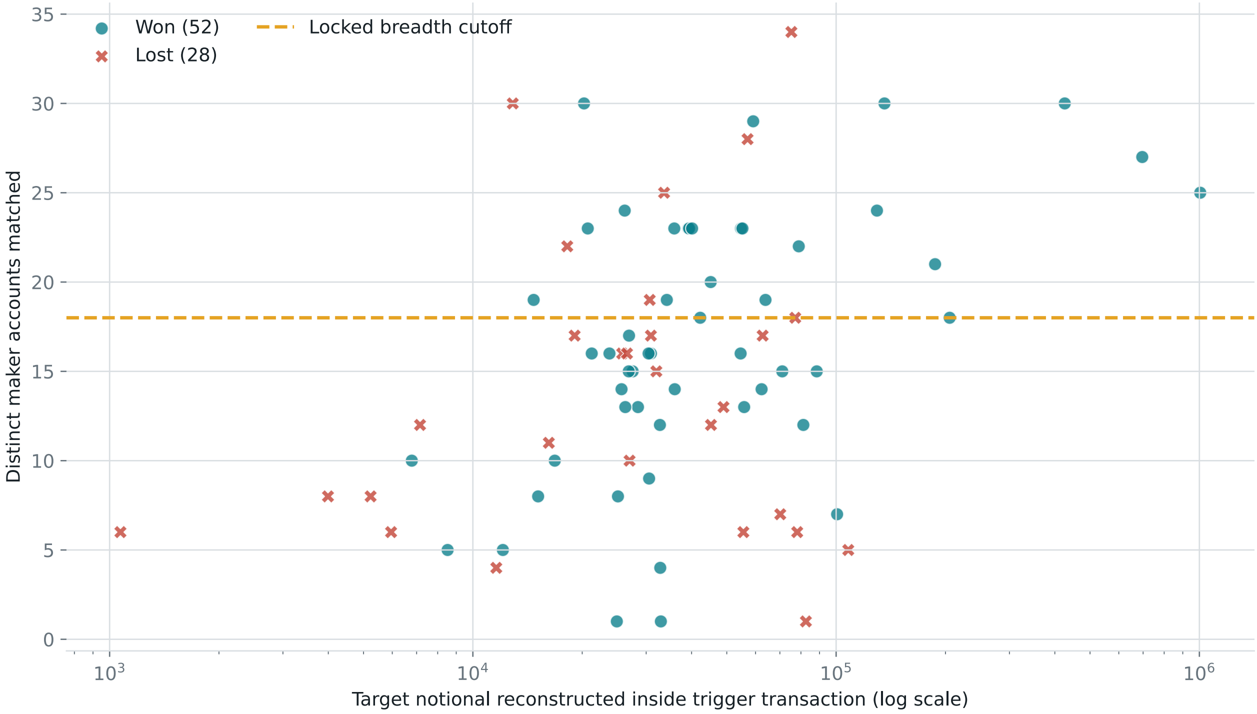

A bigger dollar trade is not the same thing as a broader sweep

Probability-offset model after urgency and period controls: breadth OR 4.94, p=0.043; log notional OR 0.94, p=0.858. Retrospective explanatory model.

Large notional appears above and below the line and among winners and losers. The informative dimension is how broadly one order consumed standing liquidity, not simply how many dollars it represented.

Did it hold up after discovery?

The 18-maker threshold was selected using only the first half of eligible history. Integer cutoffs from 5 through 30 were compared, each requiring at least eight development bets. The next half was then scored without changing the rule.

The broad-sweep rule returned +76.1% in development, +5.1% in the next block, and +44.0% in the final block. Combined after selection, 21 signals won 15 times and returned +27.32%.

> **Held-out calibration:** public prices expected 11.55 wins; 15 occurred. Resampling whole trading days gave a held-out ROI interval from +0.7% to +58.3%.

Transaction breadth fixes the rapid proxy's middle-period failure

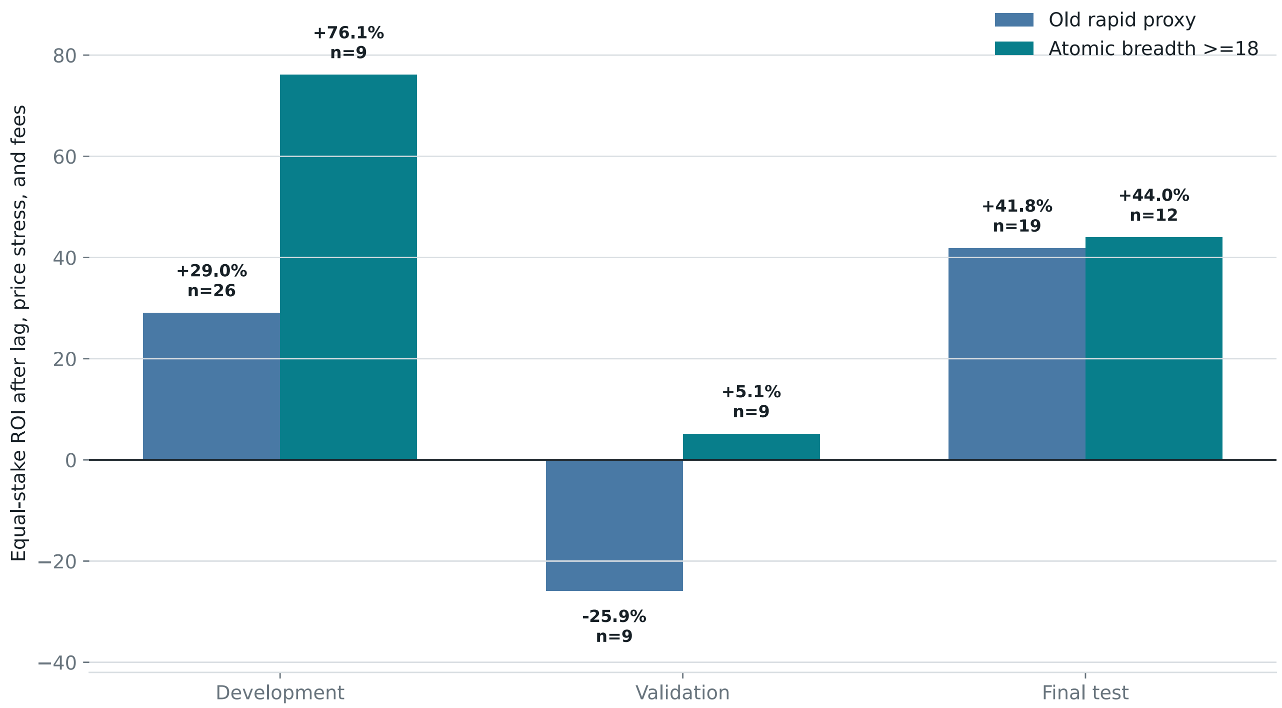

The breadth cutoff was selected only from development. Validation and final outcomes were not used to choose 18.

The middle period was modest, but crucially stayed positive. The final block rebounded without changing the 18-maker rule.

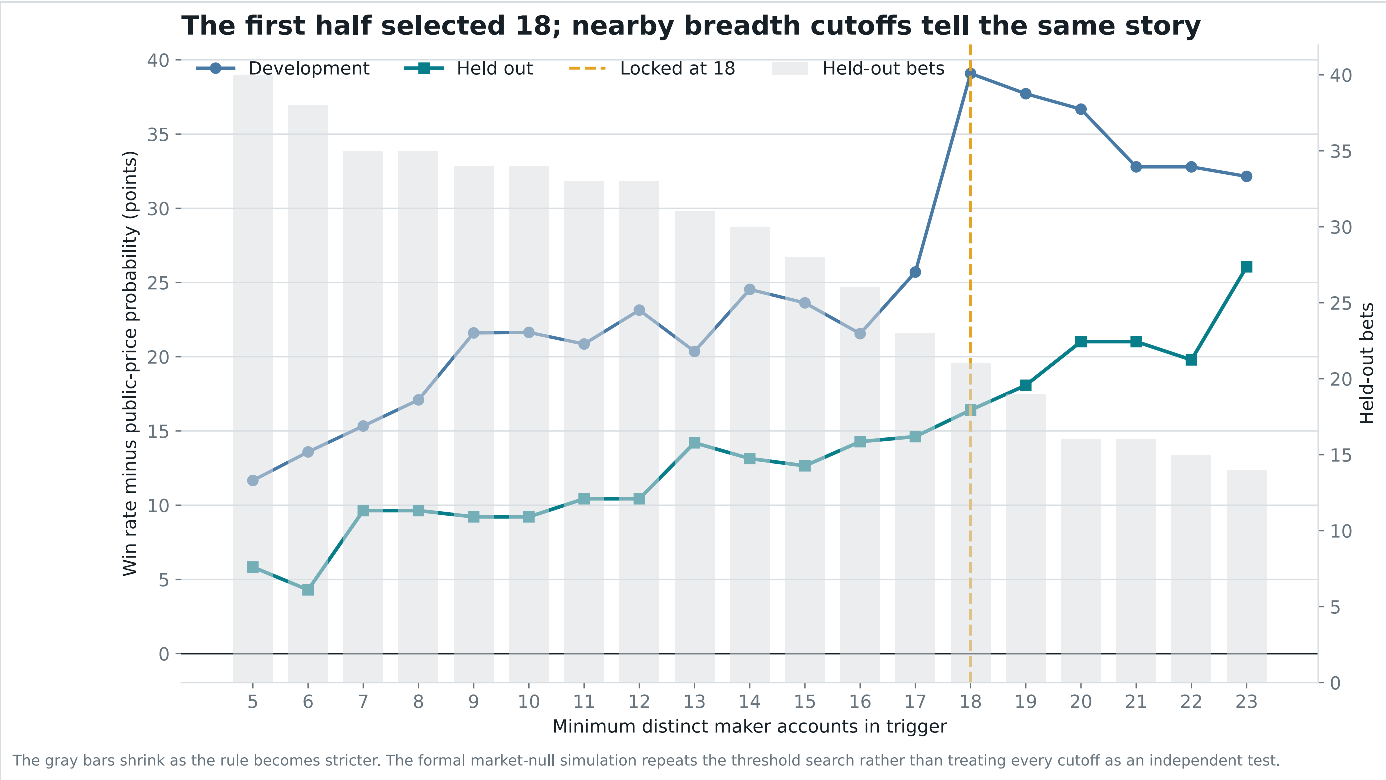

The dashed line marks the cutoff chosen on development data. The held-out curve was not used to select 18. Nearby cutoffs tell a similar story, while very strict cutoffs leave too few bets.

His alpha, literally

The strongest historical target-taker footprint is a conditional probability residual: public execution prices understated how often the target's side eventually won when one eligible BUY consumed liquidity from at least 18 distinct maker accounts. That is a fingerprint of conviction after the directional decision, not the forecasting model that made the decision.

Observable alpha definition

```
B(tx) = count(distinct makerOrders[].maker)
Eligible(tx) = first canonical event AND core discipline AND
full-event market
               AND concentration >= 70% AND trigger price in
[0.30, 0.85]
Signal(tx) = Eligible(tx) AND target is BUY taker AND B(tx) >=
18
Probability alpha = realized outcome - public execution-proxy
probability
```

Measured historical probability residual: +23.21 points over all 30 broad sweeps and +16.41 points over the 21 signals after development selection. Below 18 makers, the corresponding gap was -0.59 points.

Measured historical economic residual: +41.94% under the original 60-second plus five-cent stress; +27.32% in the post-selection half; +35.60% under the one-second plus one-cent price proxy.

> **In ordinary words:** most large buys are not special. The special-looking ones are single decisions that clear offers from many signed accounts at once. That tells us when the trader was unusually convinced. It still does not tell us why he chose that team.

The wallet-linked CS2 cases sharpen the distinction. In the M80 win, the target acted in an already favorable but unresolved round state, then supplied most of its size passively after the round. In the G2 loss, it supplied similar M80 liquidity from a far worse state and lost. Across the complete reverse-breadth population, counterparty count had no edge. The unresolved alpha is therefore the **selection and valuation residual**: team strength, roster, map, economy, series context, event quality, better telemetry, human judgment, or another input that says what each live state is actually worth.

Is esports his moat?

It is a major specialization, but the available evidence does not prove it is unique. Esports represents 204 markets, $33,006,312 of cost basis, $2,808,250 of realized P&L, and +8.51% ROI. That is 58.6% of all deployed cost basis.

The frozen broad-sweep result is strong in esports: 12/15 wins and +40.14% ROI. Dota is the strongest esports subgroup at 6/7 and +63.05%. But traditional sports independently produced 11/15 and +43.75%. The compact-fresh subset contains only 2 esports observations, versus 10 in tennis and soccer.

> **Verified answer:** esports is probably an operational specialty and Dota is the best mechanism laboratory because detailed game state exists. It is not verified as the wallet's exclusive source of alpha. Calling it the moat would ignore equally strong traditional-sports results and tiny category samples.

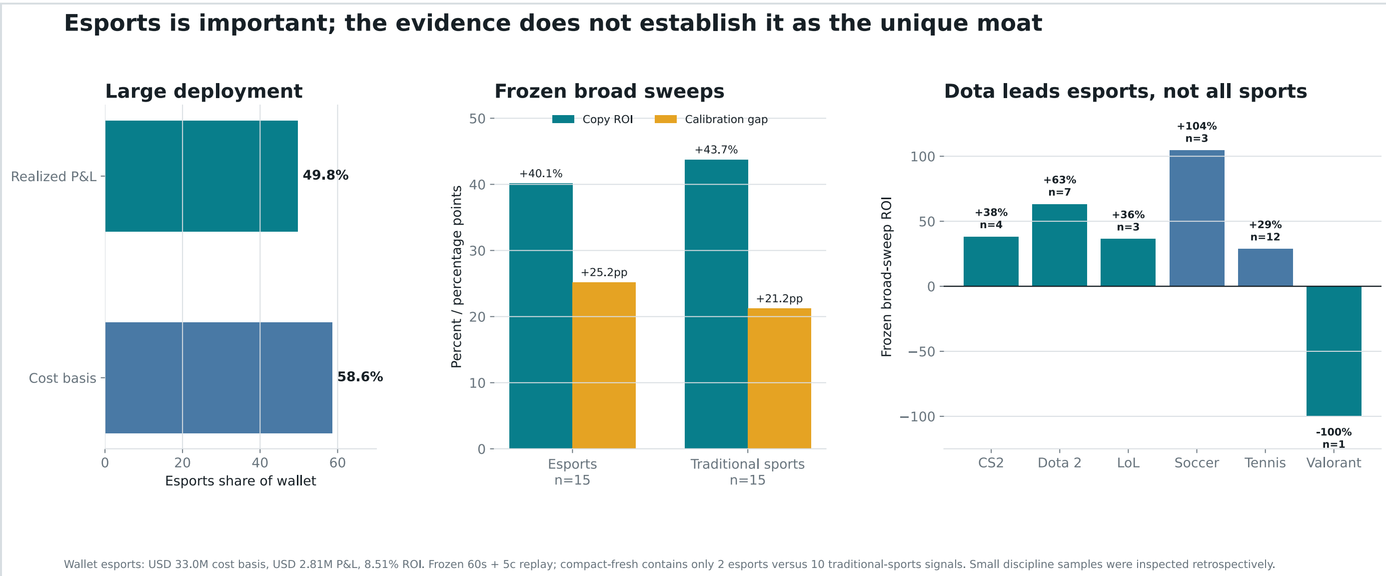

Esports dominates deployed capital, but not the frozen signal's return. Soccer and tennis contribute independently, while Dota's 7-bet result is too small to establish a unique category moat.

A live CS2 test found where the public feed loses the race

We recorded the public Polymarket sports and market WebSockets in the same process and reconstructed series-winner best bids, asks, and midpoints around one-round CS2 score changes. The capture contained 19 directional round updates. 18 had a usable book at both the preceding poll and one second after the changed score, across 6 beneficiary outcome tokens.

The public score observations were roughly 20.4 seconds apart. Of the 18 analyzable changes, 11 had already moved at least half a cent toward the round winner one second before the changed score arrived. The same 11 had moved by receipt. Their mean beneficiary midpoint move was 0.94 cents at -1 second and 1.03 cents at +1 second. Resampling the four games as clusters put the -1-second mean between 0.444 and 1.444 cents. From receipt through +1 second, the incremental mean was only 0.028 cents, with a game-cluster interval from 0.000 to 0.079.

> **The sharper discovery:** in this small live sample, the public scoreboard was a delayed confirmation, not an alpha source. In 11 of 18 analyzable transitions, a bot reacting 0.1 seconds after that message would already have missed at least a half-cent move. A bot copying the wallet after on-chain publication starts later still: it sees the decision only after the trader selected the event, obtained the information, and executed.

> **What remains unproved:** the feed does not expose an authoritative round timestamp, the sample is only 18 analyzable updates across 4 games, and the wallet did not trade those specific games. The result verifies public-feed staleness relative to the CLOB. It does not prove who repriced it, name a private vendor, or establish that the target owns a faster feed.

Live CS2 audit: this public score feed trailed the market in the captured sample

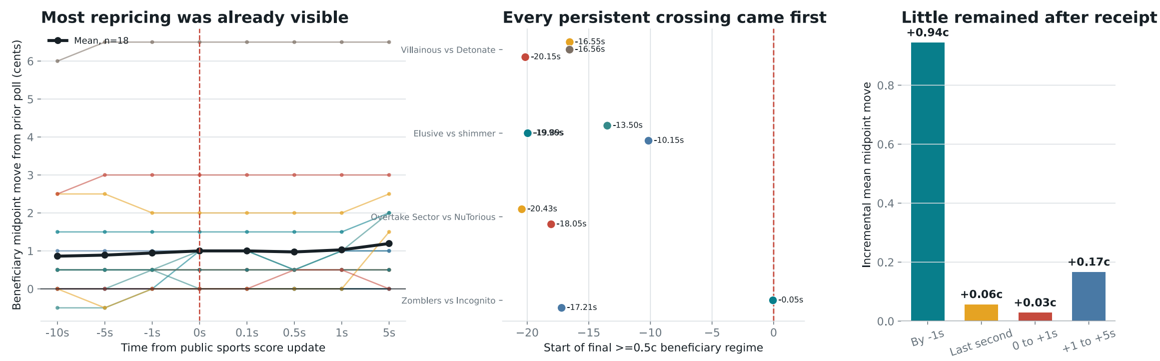

Local receive-time event study: 18 analyzable one-round updates across 4 games. The public feed exposed no authoritative round timestamp and refreshed about every 20.4s. Midpoints are not fills. This diagnoses feed staleness; it does not identify the target wallet's vendor or prove a private-data moat.

Every persistent half-cent beneficiary move in the analyzable sample started before the changed public score arrived; the median start was 17.21 seconds earlier. Most of the apparent lead is bounded by the feed's polling interval, so the chart diagnoses stale public telemetry rather than measuring the true event-to-market latency.

The wallet-linked CS2 case reveals the mechanism

This is the closest verified view of what the target was doing inside a live esports market. It is a mechanism reconstruction, not a profitable strategy claim.

The winning M80-NAVI trade

The target bought M80 in 29 fills over 23 seconds, deploying $63,374. At the first fill, the timestamp-aligned broadcast showed **M80 9-6, round 16, bomb planted, five M80 players alive against three NAVI players**. The target paid 0.74 before the broadcast first displayed M80's 10-6 round win about 7.5 seconds later. This was an early read of a visible high-probability state, not evidence of advance knowledge.

After the round, the execution mode changed. 78.9% of total quote notional was passive. From 17:04:39 to 17:04:42 UTC, 19 public taker wallets bought NAVI against 27 target maker fills. That gave the target $49,990 of M80 exposure at a 0.78 VWAP in three seconds. The full market earned $18,382 before maker rebates.

> **What that means:** the target was not chasing its own published transaction. It first crossed the book while the round was strongly favorable, then waited at a price and let other traders supply the opposite side. A copier cannot recreate those passive fills merely by being 0.1 seconds fast; it would need the same fair value, a timely quote, queue position, and incoming opposite flow.

The controls that stop a false breakthrough

Same team, same passive shape, opposite result: against G2, the target bought $83,179 of M80 and 84.5% was passive. Its key $50,310 cluster filled against 18 counterparties. The aligned broadcast showed G2 12-11 M80 with 19 seconds left, no bomb planted, and a 3-v-3. M80 needed the round for overtime, lost it, and the target's full market result was -$83,606.

A rule-sensitive loss: the target also bought $18,804 of 1WIN Map 2 exposure, mostly aggressively, before technical problems ended the map at Nemesis 7-12 1WIN. Polymarket's committed rule resolved an unfinished Map 2 at 50-50. The target later sold near 50 cents and realized -$1,248. The wallet is neither always right nor protected from market-specific resolution rules.

The broad test fails: across 405 merged wallet markets, the fixed passive rule found 9 resolved clusters with at least $25,000, 18 maker fills, and 18 unique public taker counterparties. It won 4, lost $813,357, and returned -73.46%. Counter-Strike alone was 1/2 at -36.10%. Every tested counterparty cutoff from 10 through 30 was negative.

> **Best current alpha hypothesis:** a low-latency match-state probability model, filtered by event and team context, with selective aggressive entry and passive quoting when the book is wrong. **Verified:** the state, timestamps, execution roles, counterparties, and outcomes in these cases. **Not recovered:** the target's exact fair-value calculation, upstream data source, queue logic, or a prospective rule that makes money.

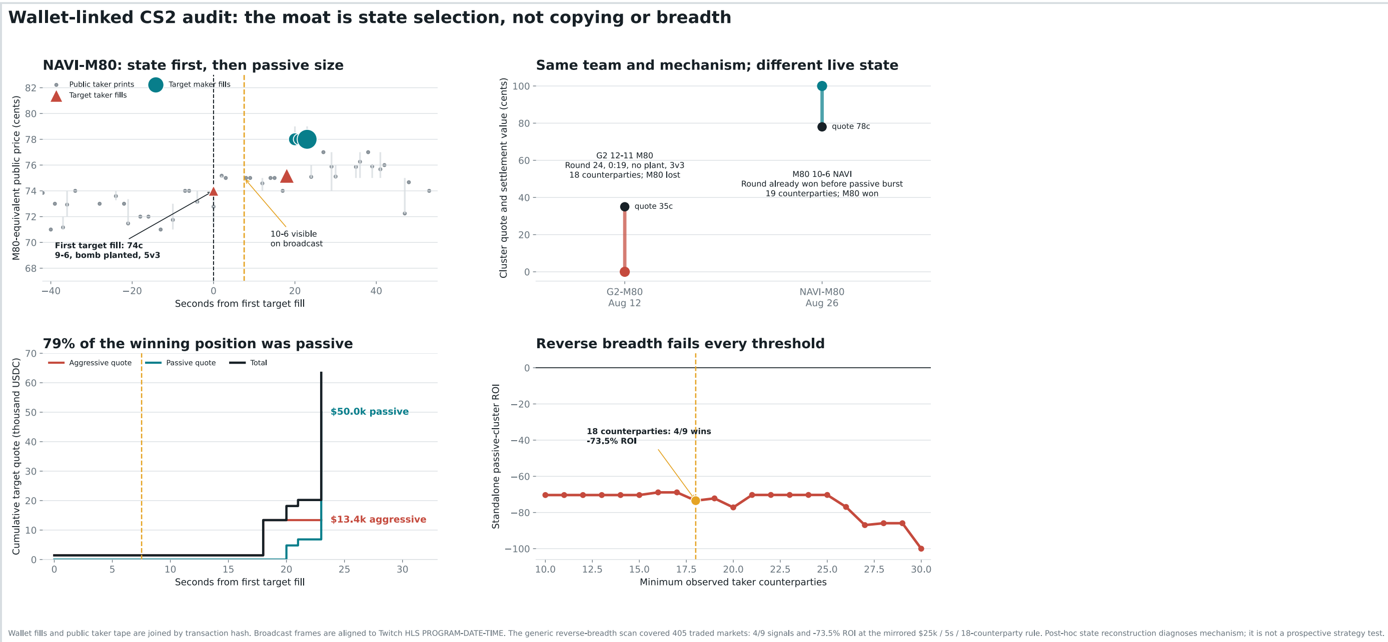

The M80-NAVI winner is real, but the G2-M80 control and the full passive-cluster population reject blind imitation. The useful discovery is the state-dependent execution mechanism; the state-selection model still has to be identified prospectively.

One Dota trade is now tied to the game itself

We matched 29 of 33 Dota threshold signals to professional matches in OpenDota. Of those, 12 occurred during a map, 13 before the first map, 4 between maps, and 4 could not be matched. Among the seven frozen broad Dota series signals, only one aligned to an in-game parsed state; five matched signals were before map one and one was unmatched.

The in-game broad case is specific. On August 5, 2026, Team Liquid was 13,433 gold and 10,921 XP ahead of Team Falcons, with a 38–32 kill lead. Liquid destroyed the top tier-two tower seven seconds before the target bought at 0.414 across 30 makers; the top tier-three fell twelve seconds after. The independent state model estimated 66.1% for Liquid against a modeled all-in follower price of 43.7%.

> **What this proves and does not prove:** the transaction was consistent with live Dota state and followed a concrete objective. It proves one state-aware decision. OpenDota's parsed replay is historical verification, not a low-latency live feed, and one beautiful case cannot explain six other broad Dota series signals.

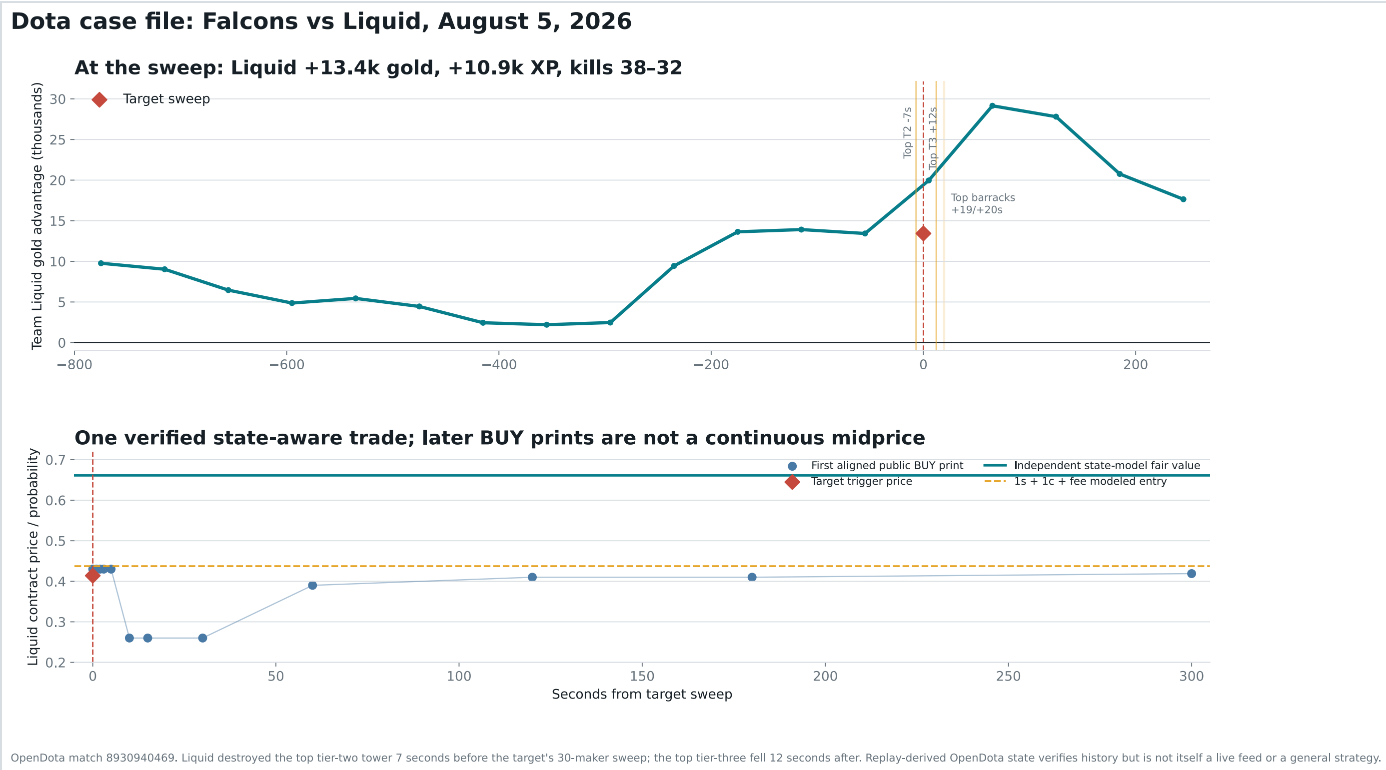

The strongest case-level link: a 30-maker Liquid sweep landed seven seconds after a tier-two tower fell while Liquid held a large state advantage. The lower panel deliberately labels later BUY prints as sparse execution observations rather than a continuous price.

We built the obvious state model, then tried to break it

A logistic model was trained on 10,000 professional Dota matches strictly before the earliest wallet Dota signal. It uses gold advantage, XP advantage, game minute, and time interactions, with one deterministic sampled minute per match and side-symmetry augmentation. Wallet outcomes were not used to fit it.

On a later chronological test of 1,499 matches and 2,998 side observations, it achieved Brier score 0.157 versus 0.25 for a coin flip, ROC-AUC 0.851, and well-tracked calibration bins. This establishes that the model reads ordinary Dota game state. It does not establish that Polymarket misprices that state.

Conditioned on the wallet's twelve mapped in-game signals, the natural nonnegative model-edge split contained all five winners; the five model-negative bets all lost. A five-point paper gate found 6 bets, 5 winners, and +61.3% ROI after a one-second public-print proxy, one adverse cent, and a 3% fee. But this was inspected after wallet outcomes, had only four day clusters, and its cluster interval ran from -42.8% to +109.0%.

A real Dota state model was built without wallet outcomes

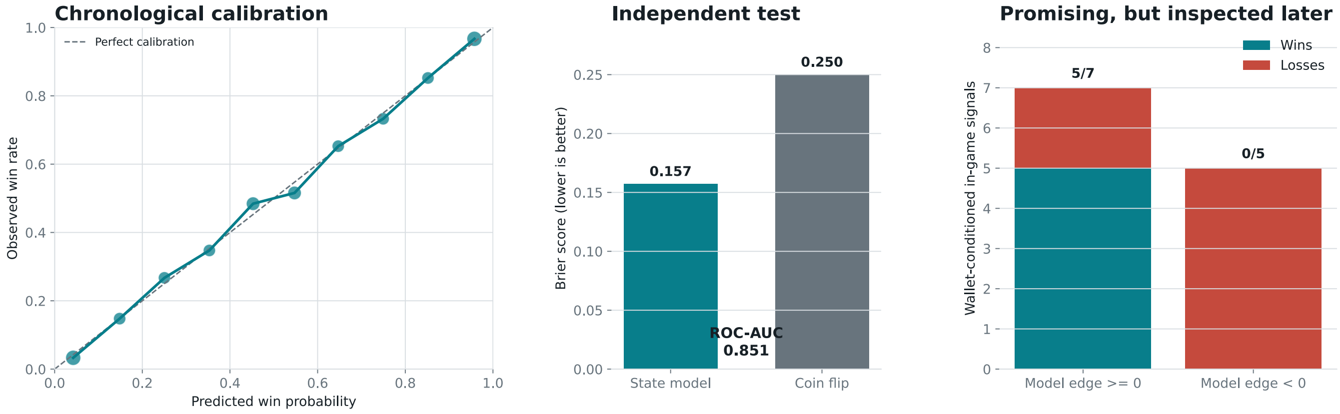

Training: 10,000 professional matches strictly before wallet signals. Chronological test: 1,499 later matches / 2,998 side observations. The 7-versus-5 wallet split is only 12 correlated, post-outcome-inspected observations (one-sided Fisher p=0.027).

The game-state model is genuinely predictive out of sample. The wallet-conditioned split is striking but tiny and retrospective, so it was treated as a hypothesis generator rather than a result.

The independent Dota test failed

The five-point gate was then detached from the target wallet and run across every mapped Dota child-moneyline market in a fixed later window, August 25, 2026 through August 26, 2026. It required the first model divergence from minute five onward, then the first unrelated public print at least one second later, one adverse cent, the fee curve, and no fallback if no print arrived within sixty seconds.

The scan covered 18 child markets with parsed state and public tape. It produced 17 signals and 9 conservative print-proxy fills. Six won. Net result: -$61.87 at -6.87% ROI, with a median actual print delay of 17 seconds.

> **This is the decisive falsification:** a calibrated state model plus apparent market disagreement is not enough. Faster 0–15 second assumptions stayed negative for the same nine-fill cohort. A 60-second row turned positive only because three additional markets became fillable, so it is a different cohort, not evidence that waiting helps.

All three losses first signaled in minutes five or six. Requiring minute eight would have produced 5/5 in this one-day sample, but that pattern was noticed after inspecting the losses and is heavily correlated by series. It is recorded as a future hypothesis, not promoted into the strategy.

Independent result: scoreboard state alone is not the target's alpha

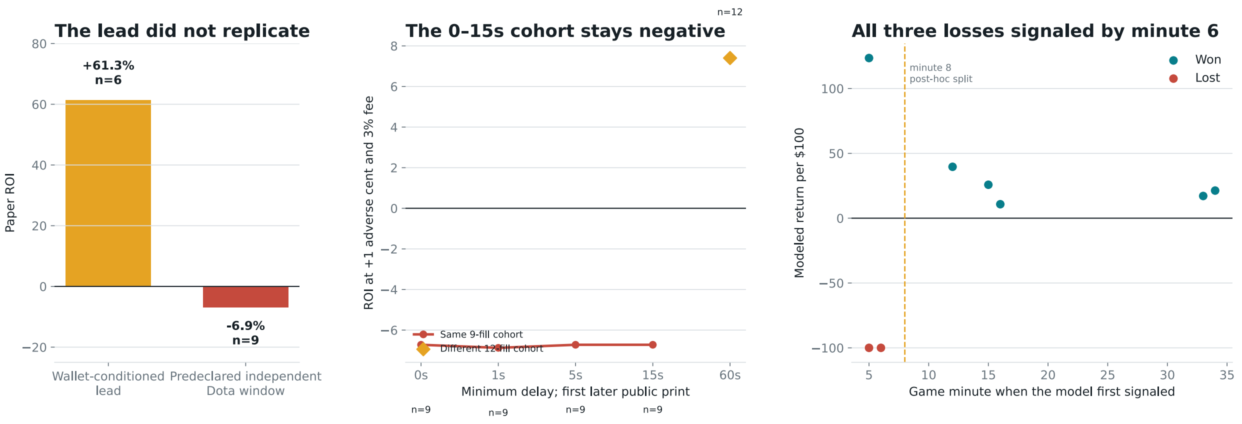

Window declared before outcome review, August 25–26: 18 parsed markets, 17 model signals, 9 conservative print-proxy fills, 6 wins, -6.87% ROI. All fills occurred on one UTC day; no historical ask-depth proof. The minute-8 line was noticed after losses and is not a strategy.

The chart puts the attractive retrospective lead beside the independent loss. The post-hoc minute-eight line is shown so future tests can freeze it, not so this report can claim it already worked.

What the failure reveals about the hidden edge

Evidence-based alpha stack

```
Observed wallet result
= upstream event / side selection that remains unrecovered
+ live-state information that can partly be modeled
+ decisive atomic execution visible as maker breadth
- price reaction, fees, depth, and delay paid by a follower
```

Recovered: a calibrated Dota state scorer, the public sports-to-market data path, exact FOK depth walking, the wallet's broad-sweep execution fingerprint, direct evidence that the sampled public CS2 score feed trailed CLOB repricing, and one wallet-linked CS2 sequence showing aggressive entry from favorable live state followed by large passive quoting.

Rejected: generic scoreboard value, indiscriminate reverse breadth, team loyalty, and blind speed. The target itself lost a similarly shaped M80 trade and a rule-sensitive 1WIN map trade.

Not recovered: a profitable rule for choosing which team, match, and live state deserves the bet. Plausible missing inputs include team and player priors, economy, map, roster and lineup information, series context, a faster licensed telemetry source, or disciplined rejection of low-quality events. None is asserted as the wallet's secret without evidence.

This is a real narrowing of the search space. The wallet is not simply reading gold advantage and buying any contract that looks cheap. It applies another selection layer before expressing conviction. That layer, rather than reaction speed alone, is now the highest-value target for prospective research.

The target appears before part of the reaction, but that reaction is a copier's cost

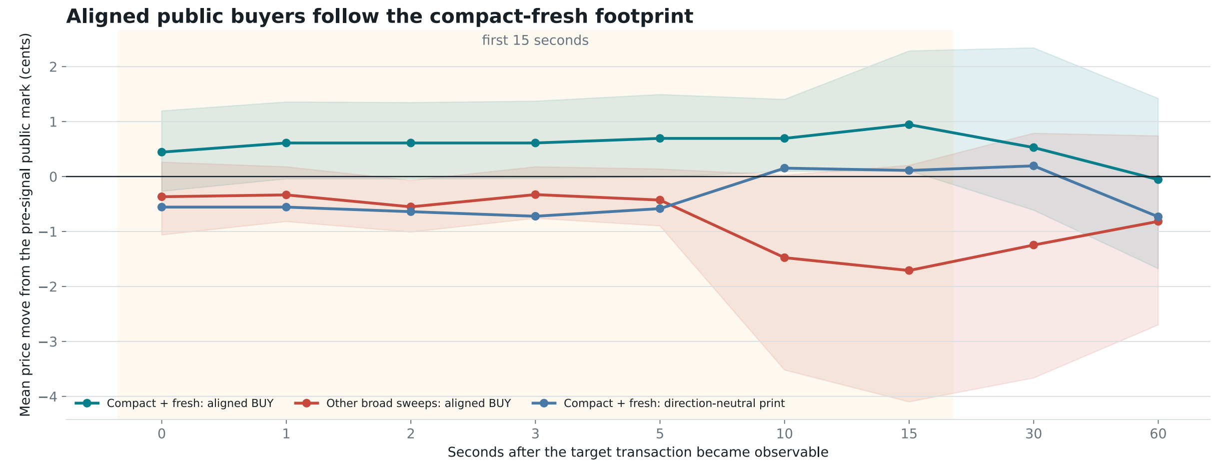

Compact + fresh n=12 across 9 UTC-day clusters; other broad n=18. Shading is the day-cluster 95% bootstrap interval for aligned BUY prints. Same-second ordering is unavailable, and this is mechanism evidence rather than a profitable entry rule.

Across 12 compact-fresh signals, aligned public BUY prices moved about 0.94 cents by 15 seconds; the day-cluster interval was 0.11 to 2.28 cents. The target appears before part of the reaction, which is useful mechanism evidence and an extra cost for a copier.

The closest answer to "what is his secret?"

Within broad sweeps, the best transactions are **compact and fresh**: many separate makers are consumed, but the order does not chase through a long price ladder, and the typical maker order was signed within five minutes.

Exploratory mechanism fingerprint

```
Strong(tx) = distinct makers >= 18
             AND distinct execution price levels <= 3
             AND median maker-order age <= 300 seconds
```

Development selected the three-level and 300-second bounds from the declared grid. It found 5 bets, all winners, at +115.9%. The unchanged rule then found 7 held-out bets, 6 winners, at +63.17%. Across all history it went 11/12 at +85.16%; the other broad sweeps went 12/18 at +13.14%.

Interpretation: this looks less like reckless price chasing and more like decisive acceptance of dense, recently posted liquidity. The wallet is taking a lot of independent-looking supply while the price ladder remains compact. That geometry is a closer footprint of high conviction than dollar amount, flow speed, or maker count alone.

Two attractive explanations failed

- **Not stale-quote harvesting.** Median maker age was 68.3 seconds in broad sweeps and 71.7 seconds in narrow ones. More importantly, broad winners consumed maker orders with median age 43.3 seconds, versus 378.4 seconds for broad losses. Winners were fresher, not staler.
- **Not one recurring set of weak counterparties.** The 30 broad signals touched 388 unique maker accounts; 102 appeared in multiple broad signals. Winners and losses had similar prior-seen maker shares, 46.7% versus 50.0%. Prior target-side outcomes against those makers also failed to separate the groups cleanly.

> **This is the sharpest lead, not a solved private model.** The compact-fresh family was proposed after inspecting the wallet. A 20,000-draw null that repeated the stated grid gave p=0.0314, but it cannot correct for every idea considered. The held-out sample contains seven bets, and its day-cluster 95% ROI interval runs from -8.7% to +110.4%.

The closest observed fingerprint: many fresh makers, few price levels

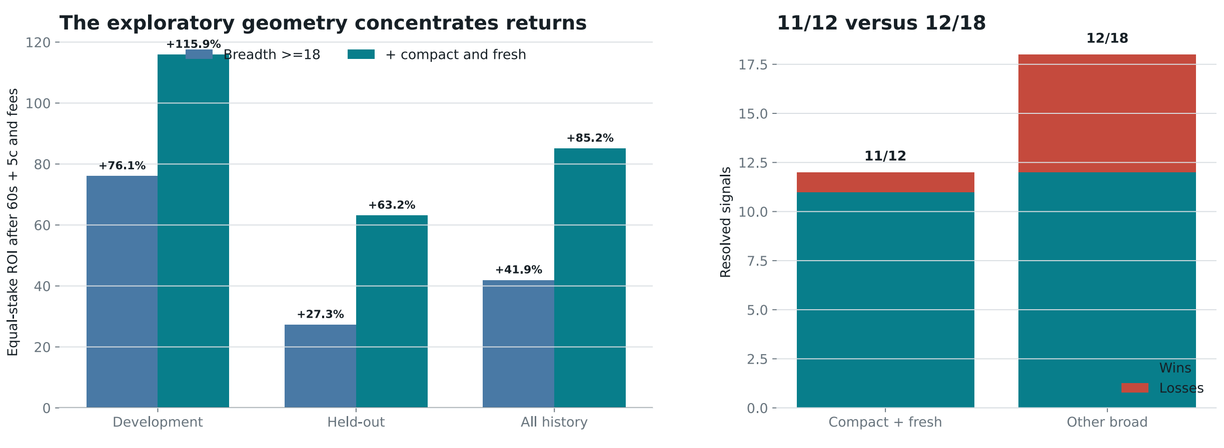

Definition selected on development data: >=18 makers, <=3 price levels, median maker age <=300s. Held-out: 6/7 and +63.17%, but day-cluster 95% CI spans -8.69% to +110.42%; this post-hoc mechanism is a lead, not a cracked private signal.

The second-stage fingerprint concentrates the result, including in the held-out block, but the sample is too small to promote it into a live strategy. The paper monitor records it as a shadow tag while the frozen 18-maker rule remains the primary gate.

The closing line refuses to confirm the secret

A genuinely informed pregame bet should often move toward the trader's side before play. That independent check did not happen here.

We found a final non-target public print before recorded start for all 33 eligible pregame events. The median print was only 7 seconds before start. Among the 12 broad pregame sweeps, median closing-line value was **-0.67 cents** and mean value was -1.06 cents. Only 4 of 12 closed higher for the target side.

A one-sided sign test for positive broad-sweep CLV gave p=0.927. Broad and narrow CLV distributions were not distinguishable either, with two-sided Mann-Whitney p=0.340. The broad group still won 9/12, leaving a +21.2 points settlement gap even at closing prices, but the market itself did not validate that information before the games began.

> **This blocks the “arm the wallet” conclusion.** Settlement outcomes look exceptional, but pregame price discovery does not confirm them, the compact-fresh pregame subset contains only 1 event, and historical fill capacity is uncertain. A live-money system would be acting on an unresolved contradiction, not a cracked code.

The market did not validate the wallet's broad sweeps before play

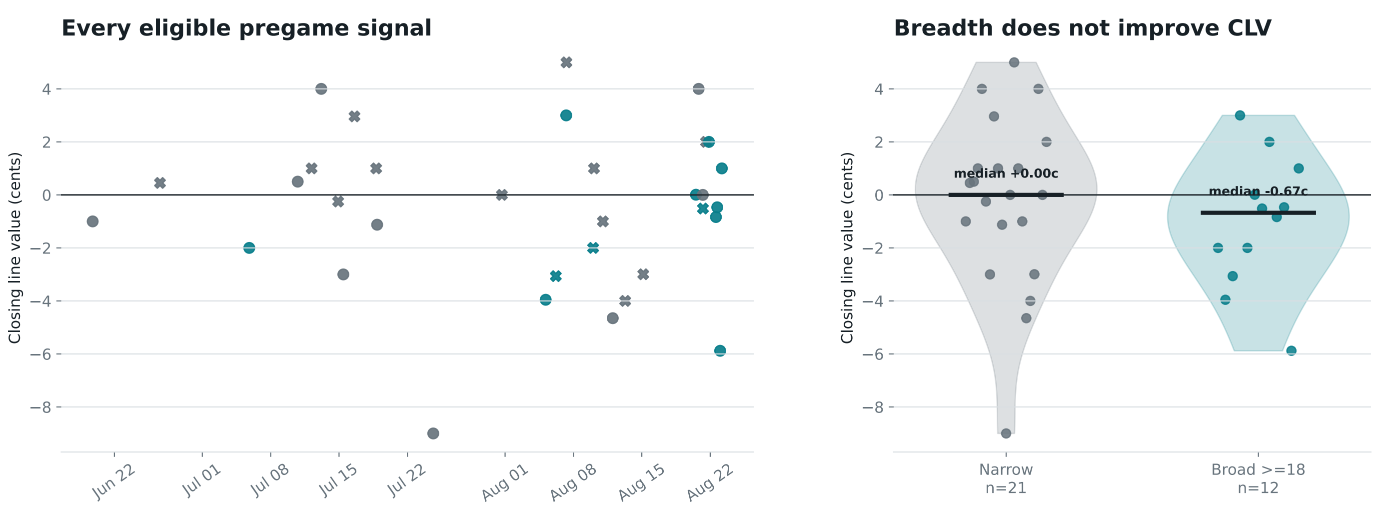

Green denotes broad sweeps; circles won and X marks lost. CLV is the final non-target public print before recorded start minus trigger price. Broad median was -0.67c; only 4/12 were positive (one-sided sign p=0.927).

Broad sweeps won more often at settlement, yet their prices generally did not strengthen before play. This is the report's strongest negative result and the main reason the mechanism remains a prospective paper hypothesis.

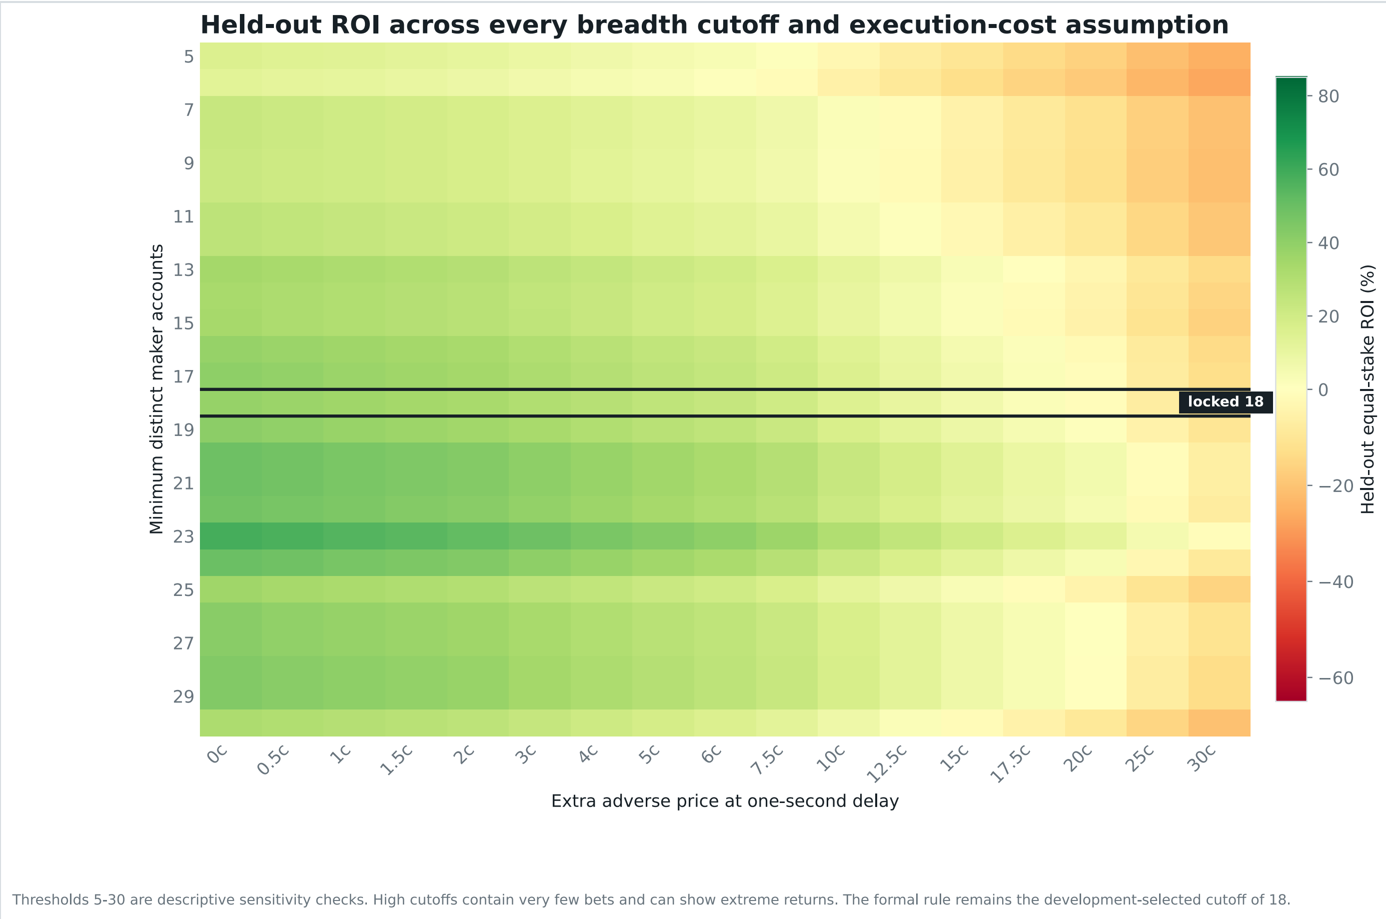

Breadth by cost: the dark horizontal outline is the frozen 18-maker row. Neighboring cutoffs form a broad profitable region at realistic costs; very strict rows are based on tiny samples and are not alternative strategies.

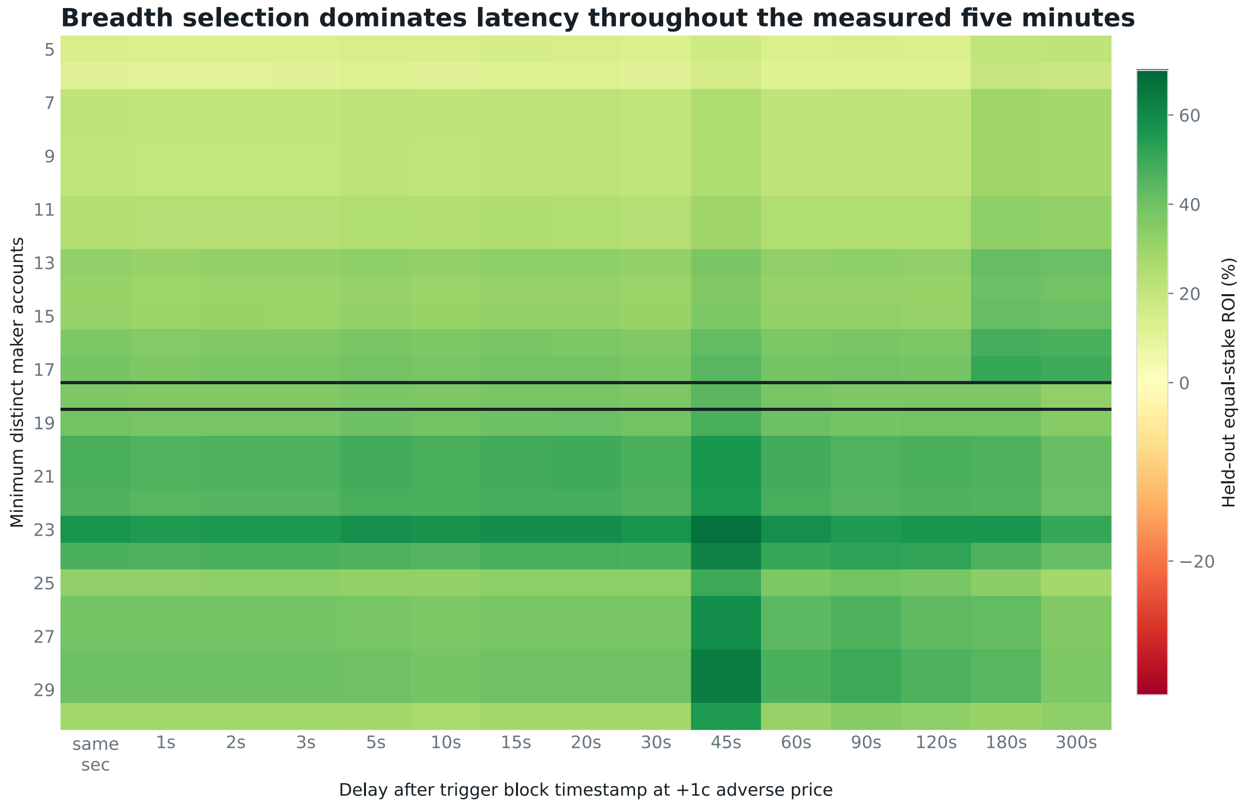

The horizontal color bands are the key result: changing the maker cutoff moves returns much more than moving from 1s to 60s. Same-second remains an optimistic bound.

Breadth by latency: horizontal bands dominate vertical changes. Choosing broad sweeps mattered much more than whether the historical proxy entered at one, five, or sixty seconds.

Speed was the clue, not the final answer

The earlier investigation found "conviction compression": at least 80% of observed taker buying arriving in the final minute. Across the full sample, rapid signals returned +24.4% while slower signals returned -24.5%. That was useful, but incomplete.

The rapid rule lost 25.9% in its middle validation block. Atomic breadth returned +5.1% in that stage and +44.0% in the final test. Speed pointed toward urgency; calldata revealed the stronger form of urgency.

> **Flow speed asks:** did the target accumulate quickly? **Atomic breadth asks:** did one executable decision consume offers from many different signed accounts? The second measure is closer to the act of demanding liquidity now.

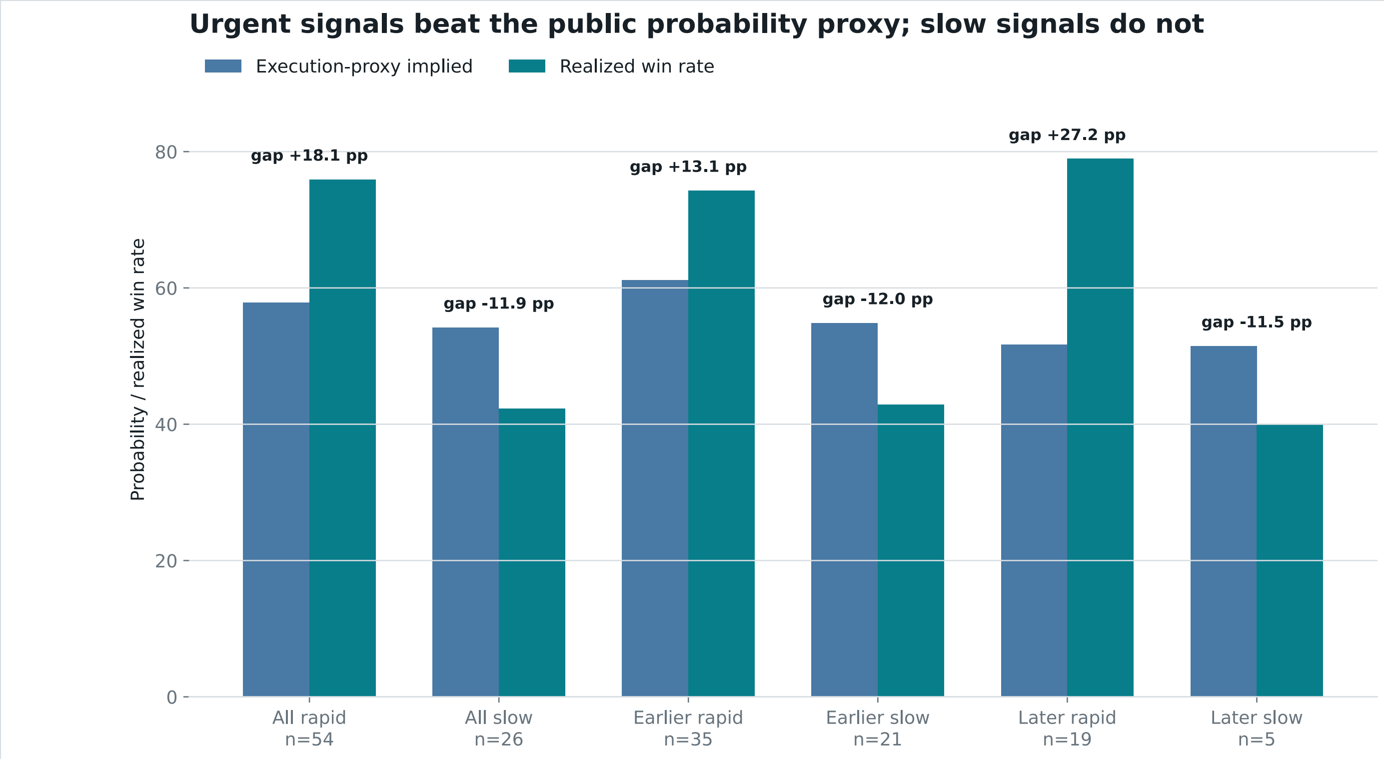

Implied probability uses the forced execution proxy before slippage and fees. Three no-print events retain trigger-price fallbacks.

Rapid buying was the first useful behavioral clue. It remains a confidence tag, not a hard requirement, because its own middle validation period lost money.

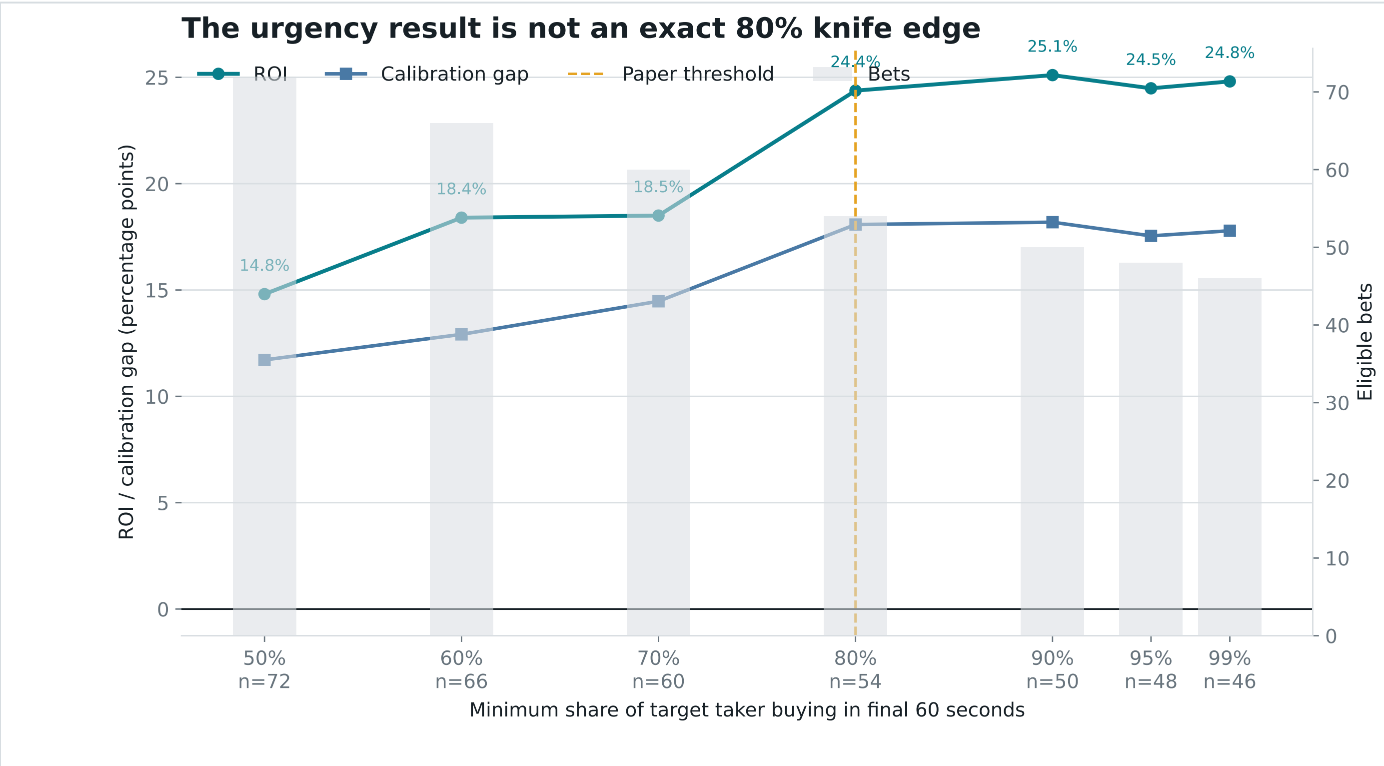

Threshold sweep is descriptive and correlated; it is not seven independent validations.

The urgency clue was not one lucky exact cutoff, but these overlapping thresholds are descriptive rather than independent confirmations.

What can actually run now: paper only

There are two auditable prototypes: the frozen atomic-breadth monitor and a live sports-to-book state probe. Both can emit paper decisions. Neither contains wallet keys, signing, approvals, or order submission.

Prototype A: atomic-breadth-18

1. Keep only the first signal for the real underlying event in core tennis, soccer, and esports.
2. Reject individual maps, single-game fragments, and short-horizon side markets.
3. Require the target to be at least 70% concentrated on one outcome and the trigger price to be 30 through 85 cents.
4. Decode the mined V2 `matchOrders` transaction and verify that the target is the BUY taker for the signaled token.
5. Count distinct addresses in `makerOrders[].maker`. Continue only when the count is at least **18**.
6. Record whether the sweep also has at most three price levels and median maker age at most 300 seconds. This is a shadow tag, not a production gate, until it has a larger prospective sample.
7. At the first observation at least one second after the block timestamp, snapshot every displayed ask level. Historical data cannot honestly distinguish 0.1 from 0.5 seconds.
8. Set the ordinary risk cap to the minimum of $100, 0.5% of bankroll, and 1% event exposure. Then reduce it again to at most **10% of displayed ask notional** through the price limit. Reject orders below $25.
9. Walk asks to calculate exact paper VWAP. The FOK limit is at most one cent above best ask, five cents above trigger, and 0.90 absolutely. If the complete size cannot fill, record a rejection rather than a partial fantasy fill.

10. Require the frozen model's predicted edge after fee-adjusted VWAP, keep one position per correlated event, cap portfolio exposure at 5%, and hold accepted paper fills to resolution.

The one-cent FOK buffer is a prospective paper convention, not a claim that historical depth existed. The replay's one-second plus one-cent price-only cell returned +35.6% in the held-out half, but only 38.1% had $100 of optimistic one-second all-print capacity. The original 60-second plus five-cent case remains the registered comparison. No live orders, private keys, approvals, martingale, inferred eventual whale size, discretionary exits, or threshold changes after a bad result are implemented.

Prototype B: live state and exact book depth

The recorder consumes Polymarket's public sports WebSocket, normalizes both observed payload shapes, extracts `gameId`, queries Gamma for the active moneyline, and dynamically subscribes the discovered token IDs on the market WebSocket. The paper engine waits one second after an external fair-value event, walks every eligible ask level, limits participation to 10% of displayed depth, charges the 3% fee model, requires at least five probability points after costs, and uses fill-or-kill semantics.

In a fresh 30-second local capture, 19 sports game IDs produced 7 active-market joins and 14 newly discovered outcome tokens. Every one of those tokens produced a market event. Gamma lookup latency was 51 ms median; median sports-message-to-first-book observation was 88 ms, with a range of 84 to 110 ms.

> **What changed:** the old historical tape cannot distinguish 0.1 seconds from 0.5 seconds, but the prospective collector can now measure both at millisecond resolution. **What did not change:** receiving a book in 88 ms is not knowing fair value, obtaining a fill, or earning profit. The capture validated plumbing, not a strategy.

The longer CS2 event study also changes the engineering priority. Optimizing this public-score path from one second to 100 ms does not recover the observed lead when the score itself arrives after the market. The next paper prototype needs an independently timestamped event source and must compare source-to-book latency, not merely message-to-order latency.

> **Why the repository will not “just arm the wallet”:** the closing-line test is negative, the compact-fresh held-out sample is seven, and post-sweep FOK depth is unknown. Connecting signing and submission now would convert unresolved research risk into unattended financial risk without adding evidence.

The public plumbing is sub-second; fair value remains the missing input

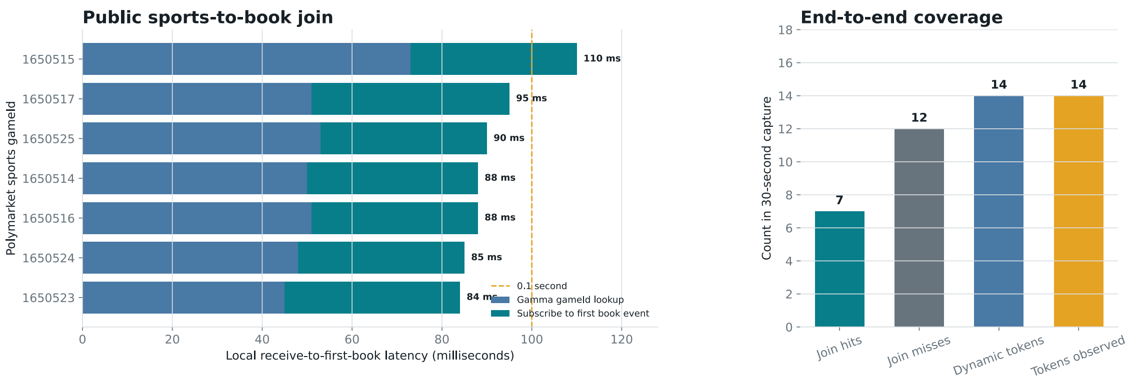

Paper-only local capture on August 27, 2026: 19 sports gameIds, 7 active-market joins, zero errors, 100% observation of newly subscribed tokens. Median sports-to-book latency was 88 ms. This measures data arrival, not order fill, signal quality, or parity with a licensed feed.

The public path joined seven active CS2 games to fourteen token books in 84–110 ms. This makes sub-second prospective measurement feasible; it does not close the missing fair-value and fill-quality gap.

Could this just be luck, sport mix, or bet size?

A promising pattern is easy to manufacture accidentally. The analysis attacked this one in several different ways.

CHRONOLOGY	**It stayed positive after selection.** The 21 held-out signals returned +27.3%, including +5.1% in validation and +44.0% in the final block.
COMPOSITION	**Similar markets still separated.** Relabeling broad sweeps only within discipline, three price bands, and fixed time period produced an effect this large with one-sided **p=0.0064** across 63 comparable bets.
TRADING DAYS	**One busy day did not create the result.** Resampling whole days put the broad-minus-narrow probability gap at +23.8 points, with a 95% interval from +6.2 points to +42.8 points.

THRESHOLD SEARCH

The declared cutoff search was simulated under market probabilities. After repeating selection and held-out scoring 20,000 times, a held-out effect this large occurred with one-sided **p=0.046**.

DOLLAR SIZE

Bigger notional did not explain breadth. A probability-offset model controlling for rapid flow, trigger notional, and time period estimated 4.94 times the outcome odds for a broad sweep, robust p=0.043. Trigger notional itself had odds ratio 0.94 and p=0.858.

> **Do not translate “18 addresses” into “18 independent humans.”** They are distinct signed maker accounts. One person or market-making system can control multiple addresses. The feature measures real contract-level execution breadth, not human headcount.

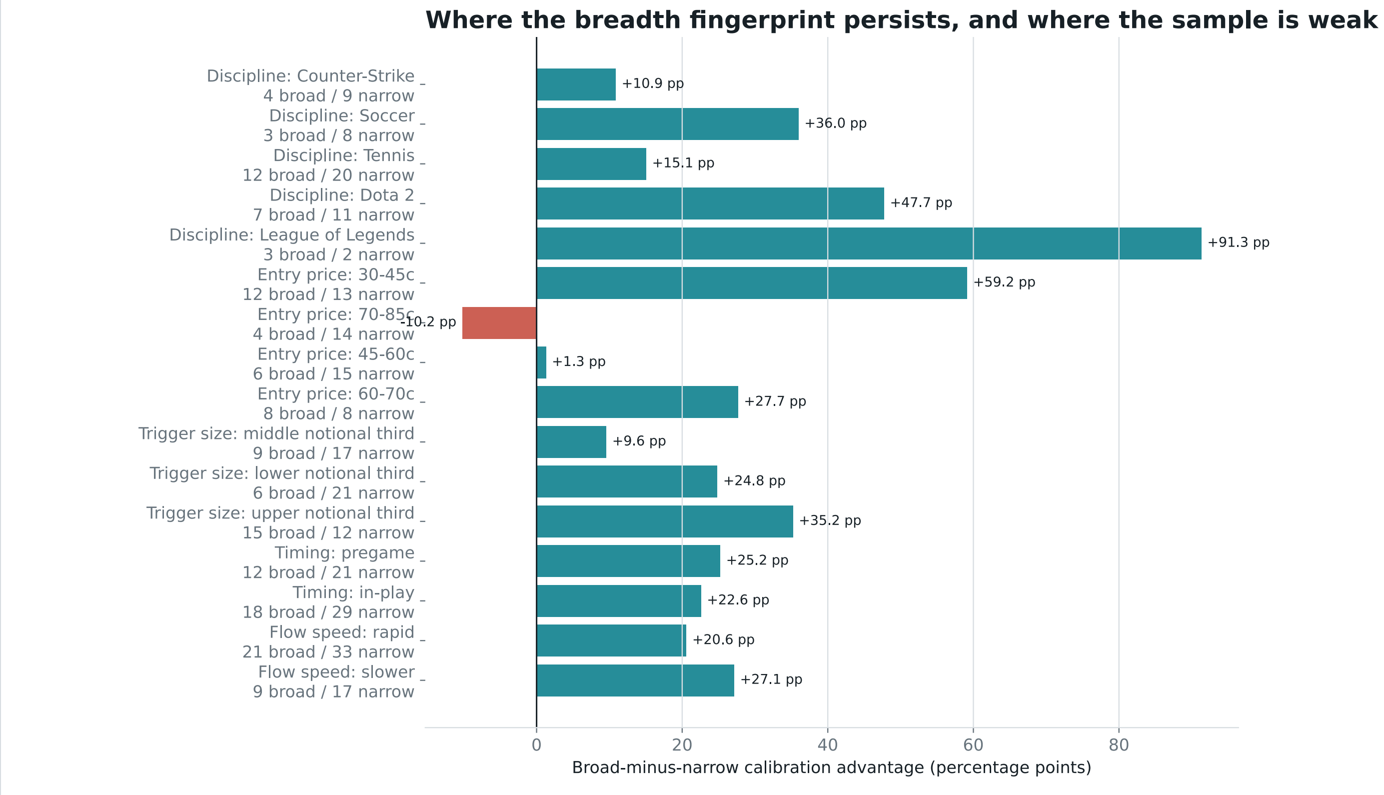

Descriptive subgroups with at least two broad and two narrow bets. These are not independent tests; small cells can swing sharply and should be treated as failure probes, not extra discoveries.

Most descriptive slices retain a positive breadth advantage. The clear weak spot is the 70-85 cent entry-price band, and several sport cells are tiny. This chart is a map of where to challenge the rule, not a menu for subgroup optimization.

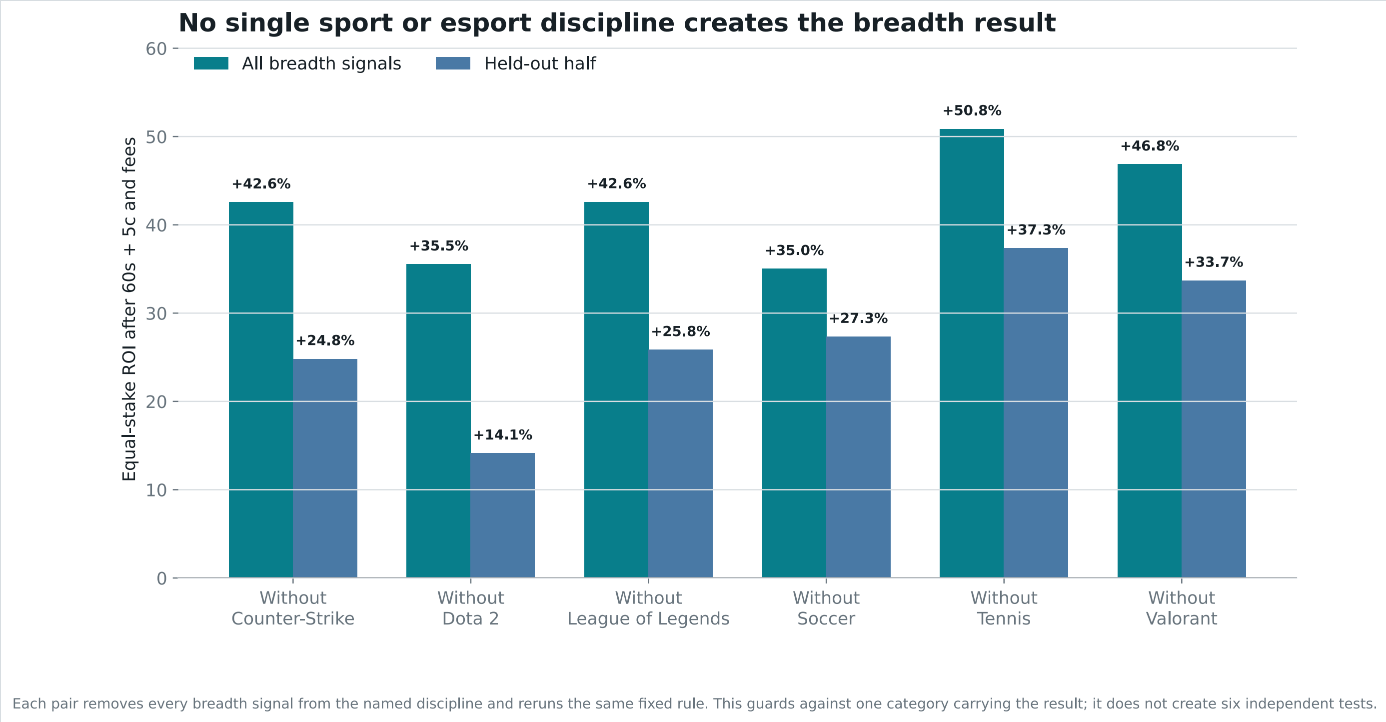

The weakest held-out leave-one-discipline-out result was still +14.1% after excluding Dota 2. No single sport or esport category creates the aggregate sign.

What the equity curve and drawdowns actually look like

Alpha is not the same as a smooth return stream. Losing contracts still lose the entire $100 stake, and a short sample can look stable by accident.

RISK VIEW	BLIND COPY	18+ MAKERS	HELD OUT
Bets	139	30	21
ROI	-6.2%	+41.9%	+27.3%
Maximum drawdown	$2,056	$172	$172
Longest losing streak	7	1	1
Worst rolling five bets	-100.0%	-14.4%	-14.4%

RISK VIEW	BLIND COPY	18+ MAKERS	HELD OUT
Profitable trading days	20 / 42	12 / 15	7 / 10
ROI after top five winners	-13.5%	+20.1%	-3.8%

The most important risk caveat: the full 30-trade breadth sample remains positive after removing its five best winners, but the 21-trade held-out half does not. That does not erase the discovery; it shows why 200 genuinely new paper signals are required before any live-money claim.

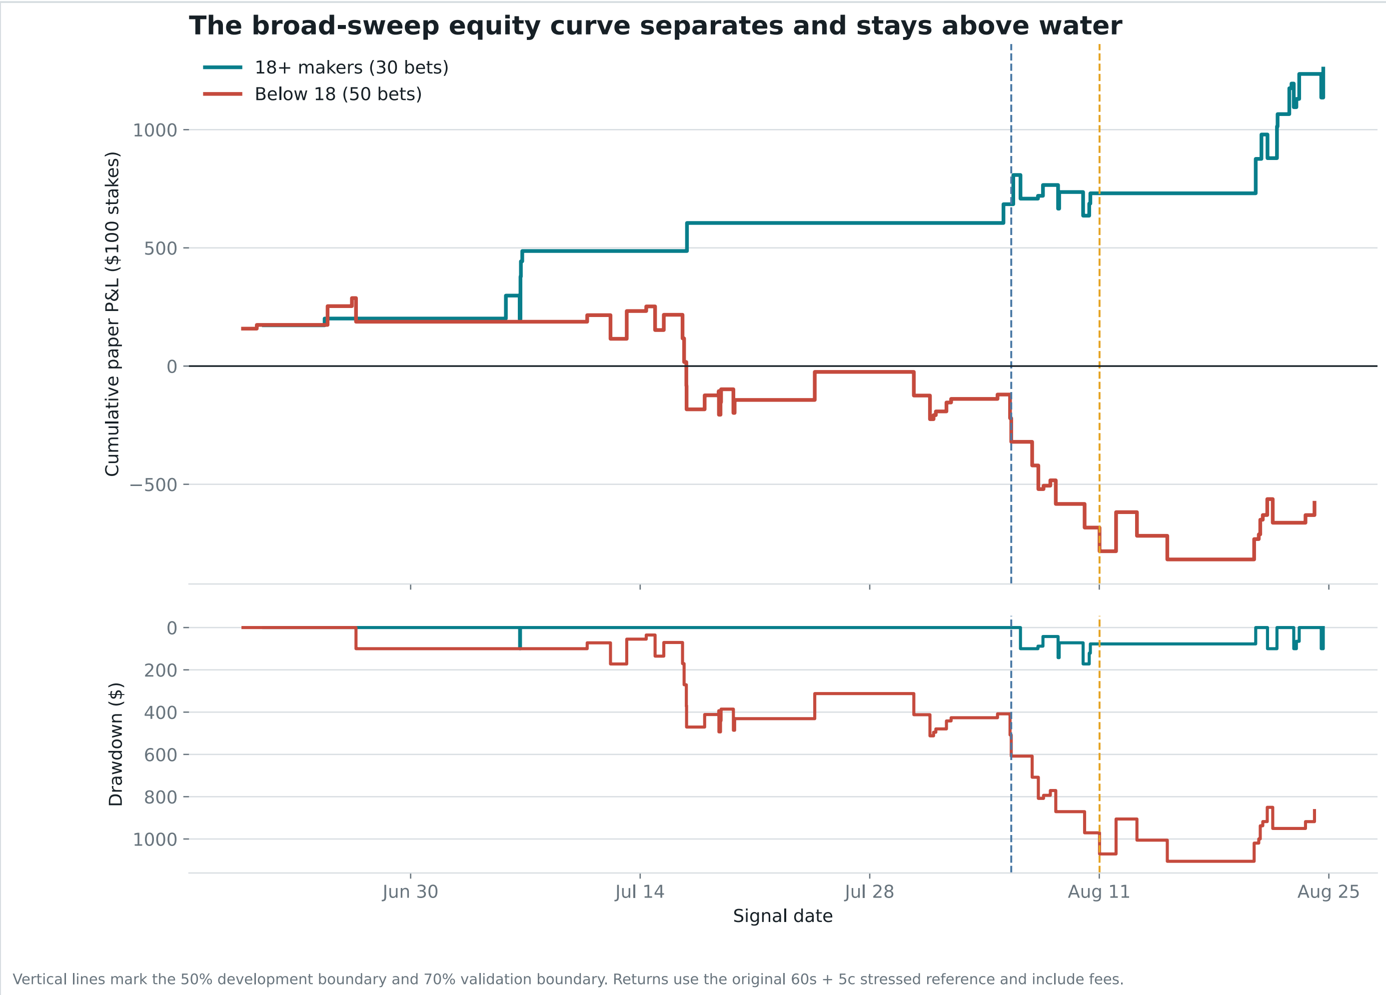

The broad curve separates early and stays positive through validation and final test. Narrow triggers deteriorate sharply. Vertical lines show the development and validation boundaries.

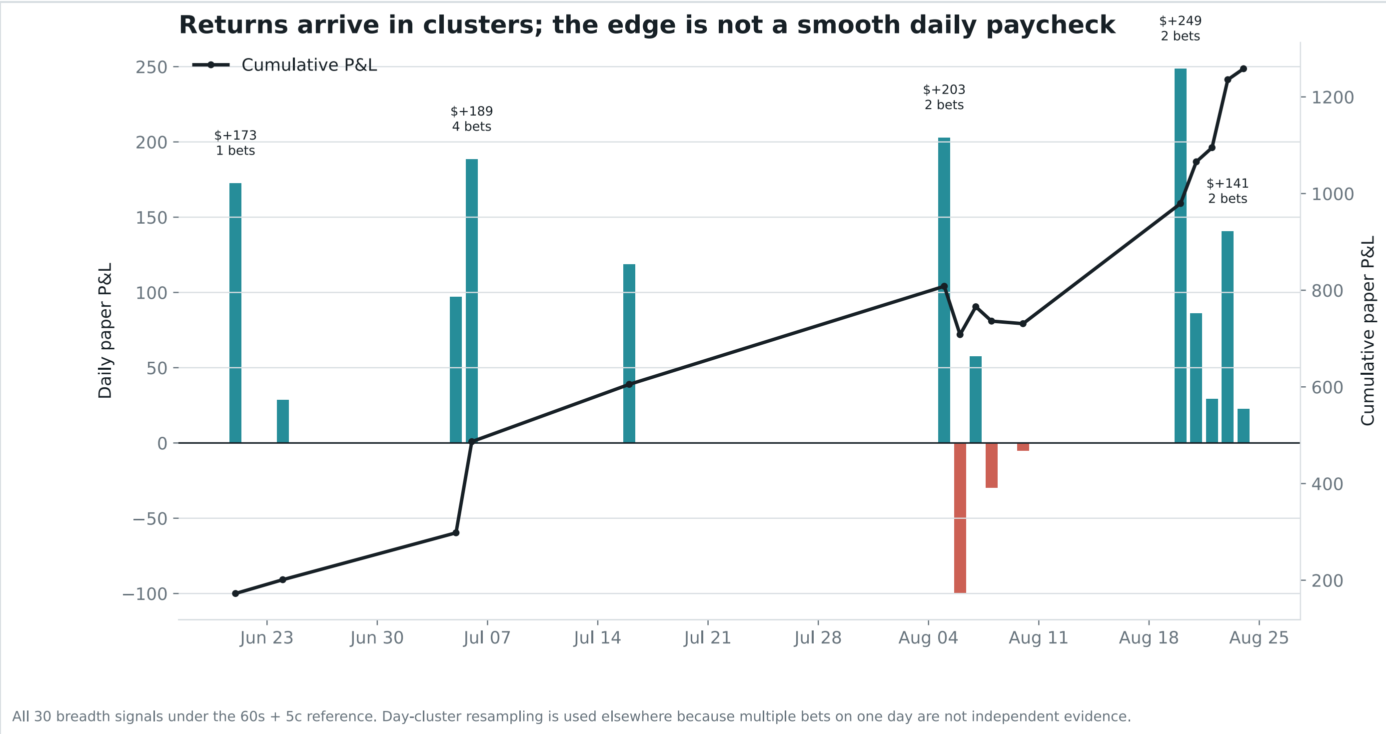

Profits arrive on clusters of signal days rather than evenly. That dependence is why the report resamples whole trading days instead of pretending every event is independent.

At the original 60-second mark, the breadth rule remains positive as the assumed price penalty rises. At ten cents, all broad signals return +29.3% and the held-out half returns +16.4%. The held-out line reaches break-even near 20 cents.

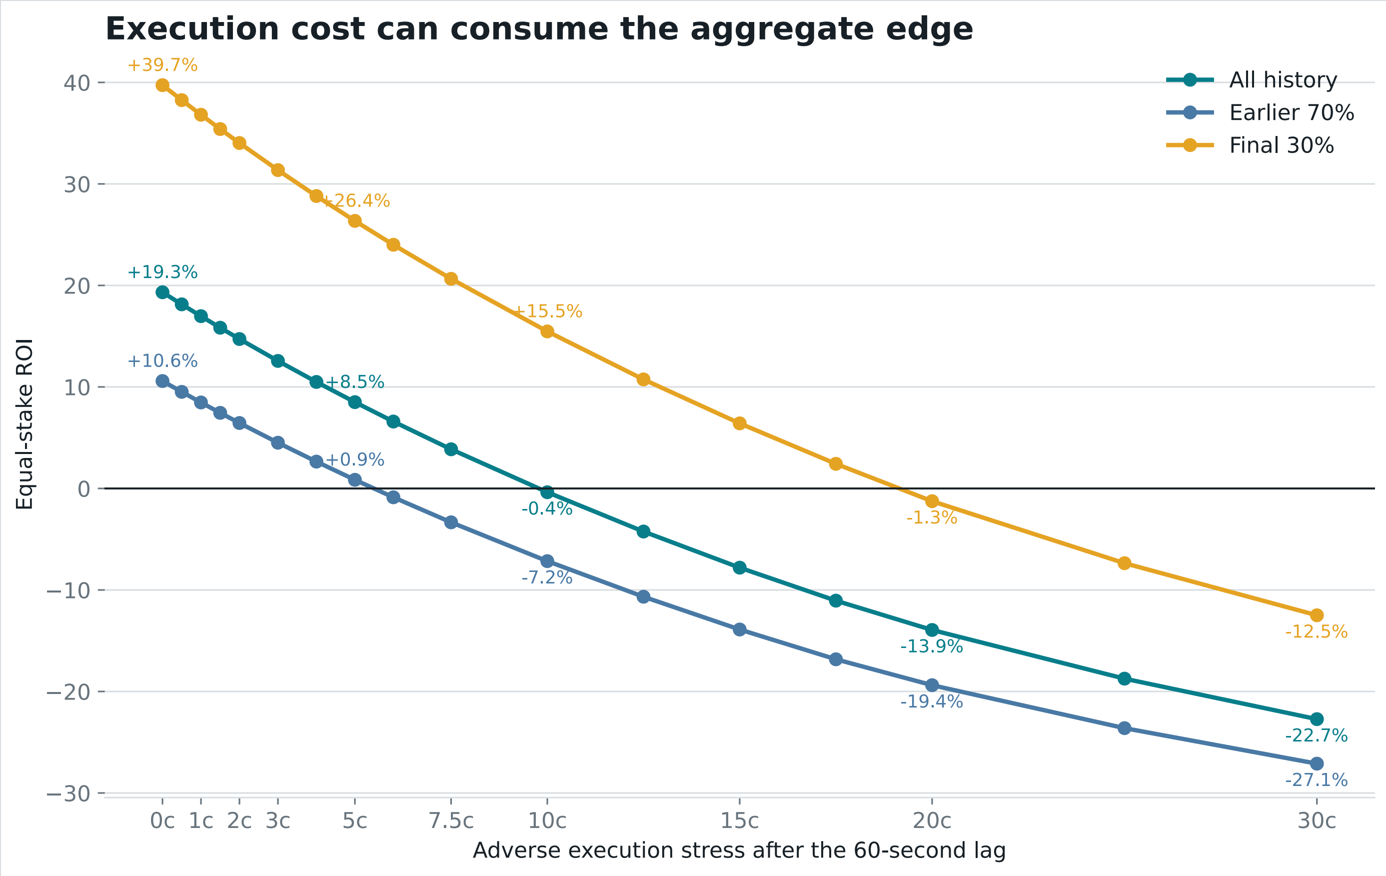

The wider eligible universe is much less forgiving. This is why a realistic copy test must report price impact rather than assuming the bot receives the whale's fill.

What the alpha reveals, and what remains hidden

CHAIN FACT

The wallet selectively takes broad liquidity. The triggering calldata directly names the maker orders consumed. Breadth is observable once the transaction is mined.

SAMPLE FACT

Those broad sweeps contain the forecasting value. They won above public-price expectations; narrower triggers did not. Dollar size does not account for the difference.

INTERPRETATION

The behavior looks like informed liquidity demand. The trader is willing to clear many standing offers immediately when waiting appears more dangerous than crossing the book.

SOURCE

The information source remains invisible. Public data cannot distinguish a superior sports model, faster public feeds, private information, coordinated research, or disciplined human judgment.

The discovery is not the trader's secret model. It is the best externally observable signature of when that hidden process is expressing unusually strong conviction.

Why the rest of the wallet still matters

- Maker orders were 92.8% of fills but only 27.9% of dollars. Fill count alone exaggerates routine activity.
- True multi-map series earned $7,569,468 at +31.2%. Single-game and map contracts lost $4,888,740.
- The first visible buy poorly predicted eventual position size; the chronological sizing model had out-of-sample R-squared -0.108.
- The whale can size, make markets, sell, and manage inventory. A follower can reproduce none of that by copying a public feed row.

The first frozen prospective window produced no eligible bet

After the historical cutoff, the refreshed wallet snapshot contained 875 trades across 19 markets. 7 crossed the raw $25,000 / 70% trigger. 2 remained in the frozen sport, format, and price universe. Neither reached 18 makers.

2 post-cutoff transactions did consume at least 18 makers, but both were MLB, a discipline excluded before the window. Dodgers–Braves consumed 40 makers with median maker age 5,597 seconds. The newer Twins–Athletics sweep consumed 19 makers at one price level with median maker age only 15.0 seconds. That second trade prospectively repeats the compact, fresh transaction geometry, but it is unresolved and outside the frozen universe, so it is not a scored strategy observation.

> **Zero eligible signals is not a zero-return backtest.** There was no frozen breadth bet to win or lose. The window provides exposure information only; assigning it 0% ROI or treating the excluded MLB sweep as validation would manufacture evidence.

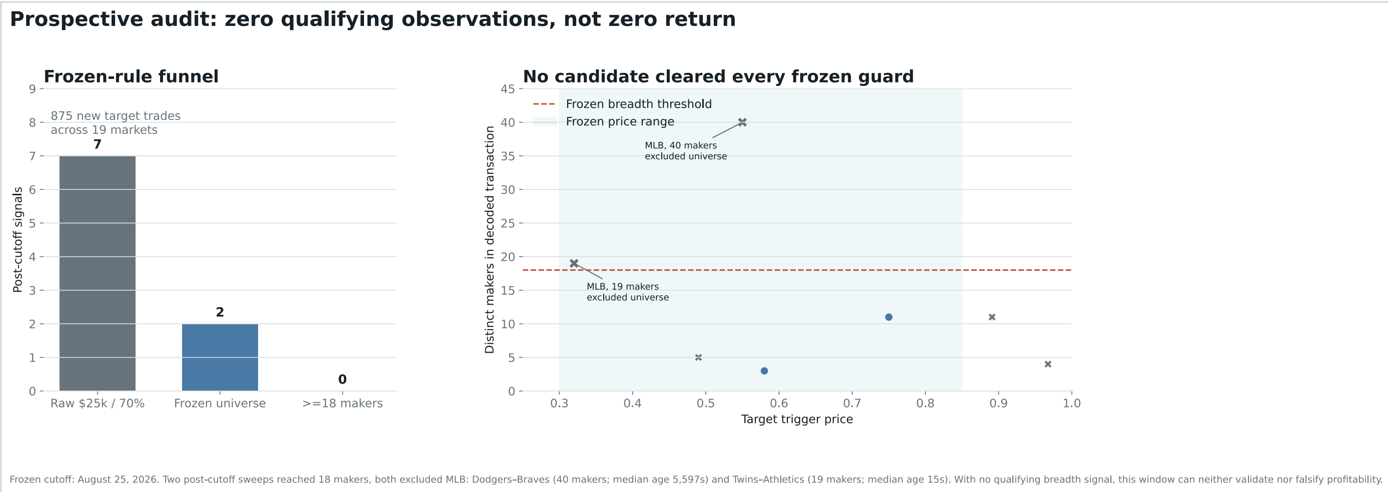

All 7 raw post-cutoff candidates are shown. Both points at or above 18 makers are outside the predeclared discipline universe; the 2 in-universe candidates stop below the breadth threshold.

The honest reasons for doubt

- **Thirty is small.** The full breadth result has 30 bets; only 21 occurred after threshold selection.
- **The held-out win-count test is suggestive, not decisive.** Its one-sided Poisson-binomial p-value is 0.086.
- **The threshold-search simulation is barely below 0.05.** Its p-value is 0.046, not overwhelming evidence.
- **The 1,444 atlas cells are not 1,444 confirmations.** They map sensitivity around one frozen rule. Treating the best cell as a new discovery would be parameter mining.
- **The 4,800 capacity cells are also not independent evidence.** They reuse the same events while varying stake, window, buffer, participation, and print proxy.
- **The wallet and feature family were chosen retrospectively.** The null simulation corrects the stated maker-count cutoff search, not every idea considered during the investigation.
- **The compact-fresh mechanism is second-stage research.** Its held-out block contains only seven bets, and its cluster interval includes a loss.
- **The history spans roughly two months.** A market regime, one operator, or one sports calendar can change.
- **The held-out result is winner-concentrated.** Removing its five most profitable winners changes held-out ROI to -3.8%.
- **Execution is still a proxy.** Public prints show activity, not guaranteed historical ask depth, queue position, partial fills, or API publication latency.
- **Closing-line value is negative.** The pregame market did not independently confirm the broad-sweep side before play.

- **The generic Dota mechanism failed independently.** A calibrated game-state model returned -6.87% in the later market-wide replay.
- **The prospective breadth audit has zero qualifying signals.** It adds no new return observation in either direction.

> **Correct conclusion:** this is a real, testable discovery in the available sample, not proof of future profit. It earns a frozen prospective paper test. It does not earn live capital yet.

Facts, interpretation, and speculation

CLAIM	CLASSIFICATION	WHY
Blindly copying every large signal lost money.	**Observed fact**	139 forced bets under declared execution assumptions.
All trigger transactions made the wallet the BUY taker.	**On-chain fact**	149 of 149 decoded V2 transactions.
Sweeps across 18 or more makers carried the edge.	**Observed in sample**	23 wins from 30, +41.9% ROI, with a positive held-out half.
Compact, fresh breadth is the closest observable mechanism.	**Exploratory lead**	6/7 held out, but selected post-hoc and the day-cluster interval includes losses.
Esports is the wallet's unique moat.	**Not established**	Esports carried 58.6% of cost, but frozen broad-sweep ROI was +40.1% versus +43.7% in traditional sports.
The target sometimes trades from live Dota state.	**Verified in one case**	The Falcons–Liquid sweep followed a target-benefiting tier-two tower by seven seconds with Liquid +13.4k gold.
A generic gold/XP/time value strategy	**Falsified in first independent window**	9 conservative fills returned -6.87% despite a calibrated state model.

CLAIM	CLASSIFICATION	WHY
reproduces his Dota alpha.		
The public esports-to-book path can be observed below one second.	**Measured plumbing fact**	Median local sports-message-to-book observation was 88 ms across 7 active joins; no order fill or alpha is implied.
The M80-NAVI buy anticipated a hidden round result.	**Not supported**	At the first fill, the aligned broadcast already showed M80 9-6 with a planted bomb and a 5-v-3. The result remained unresolved, but the favorable state was public.
The target's passive CS2 fill breadth can be copied mechanically.	**Falsified historically**	4/9 qualifying reverse-breadth clusters returned -73.46%; all tested 10-30 counterparty cutoffs were negative.
The target is consistently correct in late CS2 states.	**False**	The G2-M80 passive control lost $83,606; the unfinished Nemesis-1WIN map resolved 50-50 and cost $1,248.
Blind copy bots must be under one second.	**Not supported**	No sub-minute latency cliff appeared; roughly two cents of adverse price erased the blind-copy margin.
A $10,000 copy fills like a $100 copy.	**False**	Current +1c FOK coverage fell from 94.7% at $100 to 46.1% at $10,000 even in a favorable liquid-market sample.
Breadth is only a disguise for dollar size.	**Not supported**	Notional was null after breadth, urgency, and period controls: p=0.858.
Closing prices confirm the broad-sweep information.	**Not supported**	Median broad pregame CLV was -0.67c; only 4/12 were positive.
The trader has private information.	**Not established**	Behavior is consistent with informed trading, but public data cannot identify the source.

CLAIM	CLASSIFICATION	WHY
The first prospective audit validates atomic breadth.	**No observation**	Zero post-cutoff signals passed every frozen guard, so no return can be scored.
The algorithm is ready for live money.	**No**	State-only Dota lost independently, breadth has no prospective qualifying bet, and historical execution still relies on print proxies.

The honest verdict

The discovery is not "copy this whale." That strategy lost.

His edge, as far as public evidence can reveal it, has three layers: estimate fair probability from live match state and prior context, select only the events where that estimate is trustworthy, and execute aggressively or passively according to price and available liquidity. The M80-NAVI reconstruction directly exposes all three layers except the fair-value formula itself. The losing controls prove that state and execution shape alone are insufficient.

The frozen algorithm is **atomic-breadth-18**. Historical result: 23 wins from 30, +41.94% after the original stressed costs. Post-selection result: 15 wins from 21, +27.32%.

The strongest esports result is now a paired finding. Positively, the M80-NAVI timestamps show state-aware entry and deliberate passive deployment. Negatively, a pre-wallet Dota state model returned -6.87%, generic reverse breadth returned -73.46%, and the same-team G2 control lost. **The moat is most consistent with selective state valuation, not raw score, maker/taker count, team loyalty, or copy speed.**

For an indiscriminate bot, the historical break-even execution allowance at one second was only 1.53 cents. Size adds a second failure mode: the order can be completely rejected. The lesson is not "buy a faster server." It is "reject weak signals, cap depth participation, and control the paid price."

Decision: this evidence does not justify arming a live wallet. Continue the unchanged 18-maker taker rule until it accumulates a genuinely unseen sample, but reject passive reverse breadth as a strategy. Build the CS2 prototype as a paper market maker with independently timestamped round, economy, roster, map, and series features; record quote placement, queue assumptions, fills, cancels, adverse selection, and resolution rules. Capital waits for positive multi-day prospective P&L after exact fill constraints and after removing the largest winners.

Small glossary

Maker

A trader who leaves an order waiting in the book and provides liquidity.

Taker

A trader who accepts an available price immediately. Takers usually reveal more urgency and pay more.

Atomic sweep

One mined transaction in which a taker order matches several already-signed maker orders together.

Maker breadth

The number of distinct maker addresses consumed by that transaction. It counts accounts, not verified people.

Adverse price

How many extra cents a follower pays above the first public execution reference because of spread, consumed depth, or market reaction.

FOK

Fill or kill. The entire requested order trades immediately at or below its limit, or none of it trades.

VWAP

Volume-weighted average price across every ask level needed to fill an order. It is the true average paid when one order walks several prices.

Closing-line value

The final pregame probability minus the entry probability for the selected side. Positive means the market moved toward the bet before play.

Maker-order age

How long a signed maker order had been resting when the target consumed it. Fresh means recently posted, not necessarily well informed.

Implied probability

A 60-cent contract roughly represents a market-estimated 60% chance before trading costs.

Calibration gap

The actual win rate minus the win rate implied by prices. Positive is good only if it survives costs and honest testing.

ROI

Profit divided by total stake. A 10% ROI means $10 profit for each $100 staked.

p-value

A narrow measure of how often a result at least this large appears under a specified random null. It is not the probability that the strategy is true.

Confidence interval

A range showing statistical uncertainty under a particular resampling method. It does not include every source of bias.

Where the facts came from

This essay is a human-readable rendering of committed machine-readable evidence. The core calculations can be audited in:

- edge_analysis.json: blind-copy and atomic-breadth backtests, chronology, null simulation, controls, 1,444 parameter-atlas cells, risk diagnostics, and break-even frontiers.
- liquidity_capacity.json: timestamped current CLOB ask ladders and exact FOK walks across stake and price-buffer scenarios.
- closing_lines.json: final non-target pregame public prints used for the independent closing-line audit.
- edge_features.csv: one row per reconstructed signal, including decoded maker breadth and information available by signal time.
- trigger_transactions.json: decoded `matchOrders` calldata and derived fill anatomy for all 149 trigger transactions.
- deep_analysis.json: wallet-level execution, market format, profit, fee, and concentration facts.
- market_tape.json: 143,507 unrelated public prints used for execution proxies and market response.
- esports_state_analysis.json: deterministic Dota-to-OpenDota joins, replay state at each target signal, the 10,000-match pre-wallet model, and its chronological test.
- dota_independent_backtest.json: the target-independent August 25–26 Dota replay, every primary bet, and execution sensitivity.
- prospective_audit.json: every post-cutoff candidate and every frozen-rule rejection.
- live_probe_validation.json: event counts, dynamic game-to-market joins, measured local latency, and the SHA-256 digest of the compressed raw capture.
- esports_reaction_analysis.json: every qualifying CS2 score transition, local receive-time quote path, feed cadence, persistent-crossing timing, and the SHA-256 digest of the filtered raw trace.
- cs2_case_audit.json: wallet fill roles, exact public-tape counterparties, the M80 and control cases, and the complete reverse-breadth threshold audit.
- cs2_external_evidence.json: source URLs, page and video hashes, official map times, timestamp-aligned broadcast observations, and explicit evidence limits.

External context:

- Official Polymarket CTF Exchange V2 repository, including `matchOrders` execution and the signed order structure.
- Official Polymarket order lifecycle, distinguishing off-chain `MATCHED`, on-chain `MINED`, and final `CONFIRMED` trade states.
- Official Polymarket order-book endpoint and batch order-book endpoint, used for the current displayed-depth snapshot.
- Official FOK lifecycle documentation: a fill-or-kill order must execute completely or reject.
- Official Polymarket fee documentation, including the probability-dependent fee formula used in the paper calculations.
- Official public market WebSocket, documenting millisecond-stamped book and trade events available prospectively but absent from this historical second-resolution tape.
- Official Polymarket sports WebSocket, the live `gameId` and state source used by the paper recorder.
- OpenDota's open-source parser, plus the professional-match index and parsed match endpoint used to verify Dota state. Replay-derived availability is not represented as a low-latency feed.
- Official BLAST NAVI-M80 match record, the HLTV NAVI-M80 match page, and the timestamp-aligned broadcast VOD used for the winning state reconstruction.
- HLTV G2-M80 match page and its timestamp-aligned broadcast VOD, plus the HLTV Nemesis-1WIN record used for the losing controls.
- Dubach, The Anatomy of a Decentralized Prediction Market, supporting use of authoritative on-chain direction rather than potentially ambiguous public labels.
- Easley and O'Hara, Price, Trade Size, and Information in Securities Markets, classic context for why aggressive trade structure can contain information.
- Cheng, Yang, and Zou, Arbitrage Analysis in Polymarket NBA Markets, independent order-book evidence that executable opportunities can be sharply bounded by shallow depth. It is microstructure context, not validation of this wallet strategy.
- Bailey et al., The Probability of Backtest Overfitting.

Limit: this is public-wallet research, not proof of identity, distinct human counterparties, private information, causality, historical executable depth, or future profit. It is not financial advice.